



Hewlett Packard
Enterprise

Deployment Guide

HPE Express Containers with Docker Enterprise Edition on HPE SimpliVity

Contents

Executive Summary	5
Solution overview	5
New in this release	5
Solution configuration	6
High availability	7
Sizing considerations	7
Disaster Recovery	9
Security	9
Solution components	10
Hardware	10
Software	10
Application software	11
Preparing the environment	12
Verify prerequisites	12
Enable vSphere High Availability (HA)	12
Install vSphere Docker Volume Service driver on all ESXi hosts	13
Create the Ansible node	13
Create the Red Hat Linux template	14
Configuring the solution components	15
Ansible configuration	15
Editing the inventory	16
VMware configuration	18
HPE SimpliVity configuration	18
HPE SimpliVity backup configuration	20
Networking configuration	20
Environment configuration	21
Docker configuration	21
Orchestrator configuration	21
Kubernetes configuration	22
Protecting sensitive information	24
Inventory group variables	25
Overview of the playbooks	26
Core components	26
Optional components	26
Backup and restore playbooks	26
Convenience playbooks	27
Convenience scripts	27



Deploying the core components.....	27
Provisioning RHEL VMs.....	27
Provisioning load balancers for UCP and DTR.....	27
Installing Docker UCP and DTR on RHEL VMs.....	29
Deploying RHEL workers.....	29
HPE SimpliVity backup playbooks.....	30
Post deployment.....	30
Installing kubectf.....	30
Installing the client bundle.....	31
Installing Helm.....	31
Post-deploy validation.....	33
UCP metrics in Prometheus.....	39
Configuring storage.....	41
Deploying the NFS provisioner for Kubernetes.....	41
Manually testing the NFS provisioner.....	42
Validating the NFS provisioner using WordPress and MySQL.....	44
Deploying Windows workers.....	51
Create the Windows Template.....	51
Playbooks for adding Windows workers.....	52
Windows configuration.....	53
Deploying Windows workers behind a proxy.....	54
Windows operating system and Docker EE.....	54
Deploying Sysdig monitoring.....	55
Monitoring with Sysdig.....	55
Playbooks for installing Sysdig on RHEL.....	56
Sysdig configuration.....	56
Registering for Sysdig trial.....	57
Deploying Sysdig monitoring on Kubernetes.....	59
Deploying Sysdig monitoring on Docker Swarm.....	60
Deploying Splunk.....	61
Monitoring with Splunk.....	61
Playbooks for installing Splunk.....	63
Splunk configuration.....	63
Accessing Splunk UI.....	66
Redeploying Splunk demo.....	68
Deploying Prometheus and Grafana on Kubernetes.....	69
Monitoring Kubernetes with Prometheus and Grafana.....	69
Playbooks for installing Prometheus and Grafana on Kubernetes.....	69
Prometheus UI.....	71
Node Exporter.....	73
cAdvisor.....	74



Grafana UI	74
Deploying Prometheus and Grafana on Docker swarm	77
Monitoring with Prometheus and Grafana.....	77
Playbooks for installing Prometheus and Grafana.....	78
Prometheus and Grafana configuration	78
Accessing Grafana UI.....	78
Backup and restore	80
Backup and restore UCP and DTR.....	80
HPE SimpliVity backups.....	87
Solution lifecycle management.....	95
HPE SimpliVity environment	95
vSphere Docker Volume Service Plug-in.....	96
Red Hat Enterprise Linux operating system.....	96
Docker EE Environment.....	97
Monitoring Tools.....	97
Summary.....	97
Appendix A: Software Licenses.....	98
Appendix B: Using customer supplied certificates for UCP and DTR.....	98
Generating and testing certificates.....	98
Verify your certificates.....	101
Appendix C: Enabling SSL between the universal forwarders and the Splunk indexers using your certificates.....	101
Limitations	101
Prerequisites.....	102
Before you deploy	102
Hybrid environment Linux / Windows.....	103
Appendix D: How to check that certs were deployed correctly	104
Resources and additional links	106



Executive Summary

HPE Express Containers with Docker Enterprise Edition on HPE SimpliVity is a complete solution from Hewlett Packard Enterprise that includes all the hardware, software, professional services, and support you need to deploy a Containers-as-a-Service (CaaS) platform, allowing you to get up and running quickly and efficiently. The solution takes HPE SimpliVity infrastructure and combines it with Docker's enterprise-grade container platform and popular open source tools, along with deployment and advisory services from HPE Pointnext.

HPE Express Containers with Docker Enterprise Edition on HPE SimpliVity is ideal for customers migrating legacy applications to containers, transitioning to a container DevOps development model or needing a hybrid environment to support container and non-containerized applications on a common VM platform. This Reference Configuration provides a solution for IT operations, addressing the need for a production-ready environment that is easy to deploy and manage.

This release supports Kubernetes 1.11 via Docker Enterprise Edition (EE) 2.1, which is the only platform that manages and secures applications on Kubernetes in multi-Linux, multi-OS and multi-cloud customer environments. This document describes the best practices for deploying and operating HPE Enterprise Containers as a Service with Docker Enterprise Edition (EE). It shows how to automate the provisioning of the environment using a set of Ansible playbooks. It also outlines a set of manual steps to harden, secure and audit the overall status of the system.

Target Audience: This document is primarily aimed at technical individuals working in the operations side of the software pipeline, such as infrastructure architects, system administrators and infrastructure engineers, but anybody with an interest in automating the provisioning of virtual servers and containers may find this document useful.

Assumptions: The present document assumes a minimum understanding in concepts such as virtualization and containerization and also some knowledge around Linux®, Microsoft Windows® and VMware® technologies.

Solution overview

The HPE Express Containers with Docker Enterprise Edition on HPE SimpliVity solution consists of a set of Ansible playbooks that run on top of a VMware virtualization platform on HPE SimpliVity hardware. The solution allows you to configure a flexible OS environment (with both RHEL and Windows workers) providing built-in high availability (HA), container monitoring and security, and backup and restore functionality.



Figure 1. Solution overview

Figure 1 provides an overview of the steps used to deploy the solution. Deploying your hardware and HPE SimpliVity is specific to your environment and is not covered here. This document shows you how to:

- Prepare the VM templates
- Create the Ansible host
- Configure the Ansible parameters
- Run the Ansible playbooks

Once you are up and running, you should regularly back up the system using the scripts provided as part of this solution.

New in this release

Version 2.1 of the solution provides support for Kubernetes 1.11 via Docker EE 2.1. New features include:

- **Prometheus/Grafana on Kubernetes:** The playbooks now set up a full monitoring stack for the deployed Kubernetes infrastructure using Prometheus Operator. They install kube-state-metrics and node-exporter components, as well as supporting Kubelet and Apiserver metrics. Sample dashboards for Grafana are installed to help you monitor your Kubernetes infrastructure.



- **Docker UCP metrics for Kubernetes:** A separate, standalone Prometheus/Grafana deployment is provided to support visualization of UCP metrics. This will be integrated into the full stack deployment in a future release.
- **Sysdig for Kubernetes:** The Sysdig deployment has been updated to use Kubernetes 1.11 RBAC and config maps for sensitive data.
- **NFS Provisioner for Kubernetes:** The NFS Provisioner has been updated to use Kubernetes 1.11 RBAC.
- **WordPress and MySQL using NFS Provisioner:** Playbooks are provided to validate the NFS Provisioner, featuring a WordPress and MySQL deployment with persistent storage.
- **kubectrl:** A convenience playbook is provided to download and install kubectrl.
- **Client bundle:** A convenience playbook is available to download and configure the client bundle from UCP.
- **Helm charts:** Playbooks for downloading, installing and configuring Helm are provided, with the use of sample charts for validation purposes.

For more details on what is new in this release, see the release notes at <https://hewlettpackard.github.io/Docker-SimpliVity/rel-notes/new-features.html>.

Solution configuration

The Ansible playbooks are available to download at <https://github.com/HewlettPackard/Docker-SimpliVity>. By default, the playbooks are configured to set up a 3 node environment. This is the minimal starter configuration recommended by HPE and Docker for production.

HPE SimpliVity configuration

The operational environment is comprised of three HPE SimpliVity 380 Gen10 servers. HPE recommends dual socket HPE SimpliVity systems with at least 14 CPU cores per socket (28 total cores per system) for optimal performance and support during HA failover scenarios. Since the HPE SimpliVity technology relies on VMware virtualization, the servers are managed using vCenter.

Linux-only VM configuration

- 3 Docker Universal Control Plane (UCP) VM nodes for HA and cluster management
- 3 Docker Trusted Registry (DTR) VM nodes for HA of the container registry

The Docker UCP and DTR nodes are spread across 3 physical nodes, with one on each physical node. An odd number of manager nodes is recommended to avoid split-brain issues. It is possible to restrict the deployment to 1 UCP and 1 DTR, or to expand to more than 3, but the recommended minimum for an enterprise production deployment is 3 UCPs and 3 DTRs.

- 3 Docker Linux worker VM nodes for container workloads - Kubernetes or Docker swarm or a mix

The Docker worker nodes will be co-located with the UCP and DTR nodes in a 3 physical node deployment. Where more than 3 physical nodes are available, the worker nodes will typically be separated onto the extra nodes. It is possible to specify that more than one worker node is deployed per physical node but this decision will depend on the requirements of your applications.

- 1 Docker UCP load balancer VM to ensure access to UCP in the event of a node failure
- 1 Docker DTR load balancer VM to ensure access to DTR in the event of a node failure

By default, two load balancers are deployed to increase availability of UCP and DTR and these are placed on separate physical nodes. Load balancing for applications running on worker nodes can be achieved by using the playbooks to deploy additional load balancers, or by manually configuring the existing two to support your applications in addition to supporting UCP and DTR.

- 1 Logging server VM for central logging
- 1 NFS server VM for storage of Docker DTR images

With the addition of the NFS and logging VMs, a total of 13 VMs are created for the default Linux-only deployment. In addition to these VMs, the playbooks also set up the Docker persistent storage plug-in from VMware. The vSphere Docker volume plug-in facilitates the storage of data in a shared datastore that can be accessed from any machine in the cluster.

Hybrid VM configuration (Windows and Linux)

The hybrid deployment will typically add 3 Windows worker nodes to the above configuration, co-located with the Linux workers.

- 3 Docker swarm Windows worker VM nodes for container workloads (optional)



Note

Some of the application software supported by this configuration does not currently run on Windows, for example, the Sysdig Software Agent (see the section [Monitoring with Sysdig](#)).

High availability

Uptime is paramount for businesses implementing Docker containers in business critical environments. The HPE Express Containers with Docker Enterprise Edition on HPE SimpliVity solution offers various levels of high availability (HA) to support continuous availability. The Docker EE system components run on multiple manager nodes in the cluster. The management plane continues to operate even in the event of a manager node failure. Application containers can be protected through the use of **services** running on top of swarm. The swarm orchestrator works to maintain the number of containers declared as part of the service. The Ansible playbooks can be modified to fit your environment and your high availability (HA) needs.

Load Balancers

This solution also deploys load balancers in the system to help with container traffic management. There are two load balancer VMs – the UCP load balancer and DTR load balancer. The playbooks can be configured to deploy one or more worker load balancers depending on the requirements of your applications. A typical load balancer architecture for applications running on Docker EE is shown in Figure 2. The playbooks now support load balancers based on VRRP, using HAproxy and **keepalived**. The solution can be deployed using these loadbalancers, or external load balancers, or no load balancers or the legacy version of standalone load balancers. For more information on HAproxy, see <http://www.haproxy.com/solutions/high-availability/>.

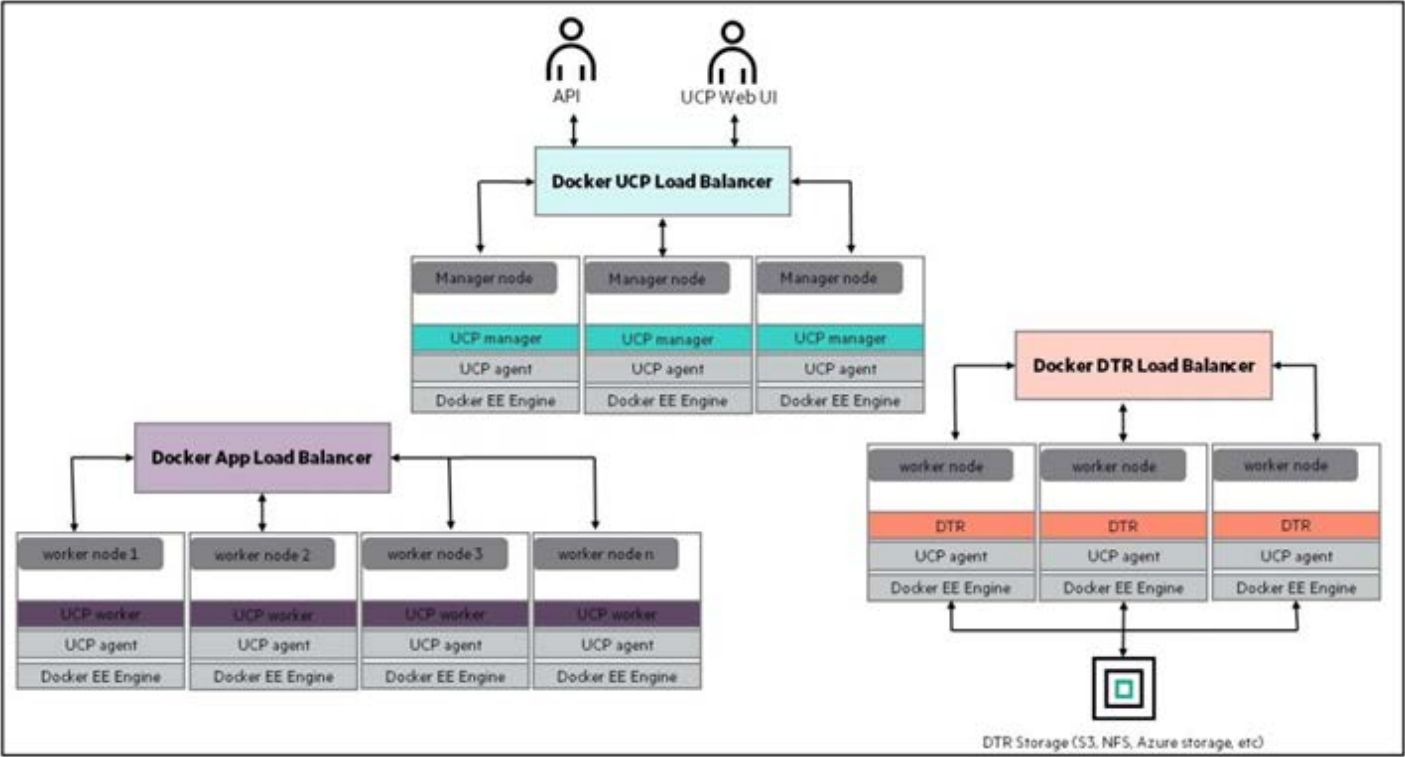


Figure 2. Load balancer architecture

Sizing considerations

A node is a machine in the cluster (virtual or physical) with Docker Engine running on it. There are two types of nodes: managers and workers. UCP will run on the manager nodes. Although DTR runs on a worker node, Docker does not recommend running other application containers on them. To decide what size the node should be in terms of CPU, RAM, and storage resources, consider the following:



1. All nodes should at least fulfil the minimal requirements, for UCP 3.0, 8GB of RAM and 4GB of storage. For production systems, 16GB of RAM is recommended for manager nodes. More detailed requirements are in the Docker EE UCP documentation at <https://docs.docker.com/ee/ucp/admin/install/system-requirements/>.
2. UCP controller nodes should be provided with more than the minimal requirements, but won't need much more if nothing else runs on them.
3. Ideally, worker node size will vary based on your workloads so it is impossible to define a universal standard size.
4. Other considerations like target density (average number of containers per node), whether one standard node type or several are preferred, and other operational considerations might also influence sizing.

If possible, node size should be determined by experimentation and testing actual workloads; and they should be refined iteratively. A good starting point is to select a standard or default machine type for all nodes in the environment. If your standard machine type provides more resources than the UCP controller nodes need, it makes sense to have a smaller node size for these. Whatever the starting choice, it is important to monitor resource usage and cost to improve the model.

For this solution, the following tables describe sizing configurations, assuming 3 Linux workers and 3 Windows workers. The vCPU allocations are described in Table 1.

Table 1. vCPU

vCPUs	node01	node02	node03
ucp1	4		
ucp2		4	
ucp3			4
dtr1	2		
dtr2		2	
dtr3			2
worker1	4		
worker2		4	
worker3			4
win-worker1	4		
win-worker2		4	
win-worker3			4
lb1	2		
lb2		2	
nfs			2
logger		2	
Total vCPU per node	16	18	16

Note

In the case of one ESX host failure, 2 nodes are enough to accommodate the amount of vCPU required.



The memory allocation for this solution (3 Linux workers and 3 Windows workers), is described in Table 2.

Table 2. Memory allocation

RAM (GB)	node01	node02	node03
ucp1	16		
ucp2		16	
ucp3			16
dtr1	16		
dtr2		16	
dtr3			16
worker1	64		
worker2		64	
worker3			64
win-worker1	64		
win-worker2		64	
win-worker3			64
lb1	4		
lb2		4	
nfs			4
logger		4	
Total RAM required (per node)	164	168	164
Available RAM	384	384	384

Note

In the case of one ESX host failure, the two surviving hosts can accommodate the amount of RAM required for all VMs.

Disaster Recovery

Recovery Time Objective (RTO) refers to the time that it takes to recover your data and applications while Recovery Point Objective (RPO) refers to the point in time you can recover to, in the event of a disaster. In essence, RPO tells you how often you will need to make new backups.

In order to protect your installation from disasters, you need to take regular backups and transfer the backups to a safe location. This solution provides a range of convenience scripts and Ansible playbooks to help automate the backup of UCP, DTR, your swarm and your Docker volumes. See the section [Backup and restore](#) for best practices, procedures and utilities for implementing disaster recovery.

Security

The Docker Reference architecture for Securing Docker EE and Security Best Practices is available at https://success.docker.com/article/Docker_Reference_Architecture-_Securing_Docker_EE_and_Security_Best_Practices

In addition to having all logs centralized in a single place and the image scanning feature enabled for the DTR nodes, there are other guidelines that should be followed in order to keep your Docker environment as secure as possible. The HPE Reference Configuration paper for securing Docker on HPE Hardware places a special emphasis on securing Docker in DevOps environments and covers best practices in terms of Docker security. The document can be found at <http://h20195.www2.hpe.com/V2/GetDocument.aspx?docname=a00020437enw>.

In addition, the Sysdig product also provides a strong level of container security and monitoring (see the section [Monitoring with Sysdig](#)).



Solution components

This section describes the various components that were utilized in this Reference Configuration.

Hardware

HPE SimpliVity is an enterprise-grade hyper-converged platform uniting best-in-class data services with the world's best-selling server.

About HPE SimpliVity

Rapid proliferation of applications and the increasing cost of maintaining legacy infrastructure causes significant IT challenges for many organizations. With HPE SimpliVity, you can streamline and enable IT operations at a fraction of the cost of traditional and public cloud solutions, by combining your IT infrastructure and advanced data services into a single, integrated solution. HPE SimpliVity is a powerful, simple, and efficient hyperconverged platform that joins best-in-class data services with the world's best-selling server and offers the industry's most complete guarantee.

More information about HPE SimpliVity can be found at: <https://www.hpe.com/us/en/integrated-systems/simplivity.html>

Software

The software components used in this Reference Configuration are listed in Table 3 and Table 4.

Table 3. Third-party software

Component	Version
Ansible	2.7
Docker EE	2.1 with Docker EE Engine 18.09 (tested with UCP 3.1.3 and DTR 2.6.2)
Red Hat Enterprise Linux	7.6
Microsoft Windows	Server 2016
VMware	ESXi 6.5.0 and vCenter 6.5.0

Table 4. HPE Software

Component	Version
HPE SimpliVity OmniStack	3.7.6

About Ansible

Ansible is an open-source automation engine that automates software provisioning, configuration management and application deployment.

As with most configuration management software, Ansible has two types of servers: the controlling machine and the nodes. A single controlling machine orchestrates the nodes by deploying modules to the Linux nodes over SSH. The modules are temporarily stored on the nodes and communicate with the controlling machine through a JSON protocol over the standard output. When Ansible is not managing nodes, it does not consume resources because no daemons or programs are executing for Ansible in the background. Ansible uses one or more inventory files to manage the configuration of the multiple nodes in the system.

When deploying Windows nodes in a hybrid deployment, the Ansible playbooks make use of the Python `pywinrm` module which carries out actions via the Windows remote manager.

More information about Ansible can be found at <http://docs.ansible.com>.

About Docker Enterprise Edition

Docker Enterprise Edition (EE) is the leading enterprise-ready container platform for IT that manages and secures diverse applications across disparate infrastructure, both on-premises and in the cloud. Docker EE provides integrated container management and security from development to production. Enterprise-ready capabilities like multi-architecture orchestration and secure software supply chain give IT teams the ability to manage and secure containers without breaking the developer experience.



Docker EE provides:

- Integrated management of all application resources from a single web admin UI.
- Frictionless deployment of applications and Compose files to production in a few clicks.
- Multi-tenant system with granular role-based access control (RBAC) and LDAP/AD integration.
- Self-healing application deployment with the ability to apply rolling application updates.
- End-to-end security model with secrets management, image signing and image security scanning.

More information about Docker Enterprise Edition can be found at <https://www.docker.com/enterprise-edition>.

Application software

A number of different logging and monitoring solutions are supported by this solution:

- Splunk
- Sysdig
- Prometheus and Grafana

The application software components used in this Reference Configuration are listed in Table 5.

Table 5. Application software

Component	Version
Splunk	7.1.2
Sysdig	latest
Prometheus	V2.3.2
Grafana	5.2.3

Monitoring with Splunk and Sysdig

The solution can be configured to use either Splunk or Sysdig or to enable both simultaneously. While there is some overlap in the functionality provided by these tools, they are ultimately complimentary in what they offer. Splunk aggregates logging and tracing for a wide variety of sources and provides a clean, high-level dashboard for all your enterprise systems. Sysdig, on the other hand, has been engineered from the ground up to focus on containerized environments and includes both monitoring and security features, with built-in understanding of the different workloads running on your cloud.

More information on configuring Splunk and running the relevant playbooks can be found in the section Deploying Splunk.

For more information on configuring Sysdig and running the relevant playbooks, see the section Deploying Sysdig monitoring.

Monitoring with Prometheus and Grafana

The solution can be configured to enable the use of Prometheus and Grafana for monitoring. In this setup, there is no need for native installs and all the required monitoring software runs in containers, deployed as either services or stacks.

The solution supports two separate monitoring stacks, with one running on Kubernetes and the other using Docker swarm.

For more information on running Prometheus and Grafana on Kubernetes, see section Monitoring Kubernetes with Prometheus and Grafana.

For more information on running Prometheus and Grafana on Docker swarm, see section Monitoring Docker swarm with Prometheus and Grafana



Preparing the environment

This section describes in detail how to prepare the environment that was outlined in the architecture section. The following high level steps are required:

- Verify prerequisites
- Enable vSphere High Availability (HA)
- Install vSphere Docker Volume Service driver on all ESXi hosts
- Create the Ansible node
- Create the Red Hat Linux Template and configure the `yum` repositories
- Create the Windows Template (optional)
- Finalize the template

Verify prerequisites

Before you start deployment, you must assemble the information required to assign values for each and every variable used by the playbooks. The variables are fully documented in the section [Configuring the solution components](#). A brief overview of the information required is presented in Table 6.

Table 6. Summary of information required

Component	Details
Virtual Infrastructure	The FQDN of your vCenter server and the name of the Datacenter. You will also need administrator credentials in order to create templates and spin up virtual machines.
HPE SimpliVity Cluster	The name of the SimpliVity cluster and the names of the members of this cluster as they appear in vCenter. You will also need to know the name of the SimpliVity datastore where you want to land the various virtual machines. You may have to create this datastore if you just installed your SimpliVity cluster. In addition, you will need the IP addresses of the OmniStack virtual machines. Finally you will need credentials with admin capabilities for using the OmniStack API. These credentials are typically the same as your vCenter admin credentials.
L3 Network requirements	You will need one IP address for each and every VM configured in the Ansible inventory (see the section Configuring the solution components). The recommended minimal deployment (Linux-only) configures 13 virtual machines so you would need to allocate 13 IP addresses to use this example inventory. If you have a hybrid environment with Windows workers, you will need to increase the allocation. Note that the Ansible playbooks do not support DHCP so you need static IP addresses. All the IPs should be in the same subnet. You will also have to specify the size of the subnet (for example /22 or /24) and the L3 gateway for this subnet.
DNS	You will need to know the IP addresses of your DNS server. In addition, all the VMs you configure in the inventory must have their names registered in DNS prior to deployment. In addition, you will need to know the domain name to use for configuring the virtual machines (such as <code>example.com</code>)
NTP Services	You need time services configured in your environment. The deployed solution uses certificates that are time-sensitive. You will need to specify the IP addresses of your time servers (NTP).
RHEL Subscription	A RHEL subscription is required to pull extra packages that are not on the DVD.
Docker Prerequisites	You will need a URL for the official Docker EE software download and a license file. Refer to the Docker documentation to learn more about this URL and the licensing requirements at: https://docs.docker.com/engine/installation/linux/docker-ee/rhel/ in the section entitled "Docker EE repository URL"
Proxy	The playbooks pull the Docker packages from the Internet. If your environment accesses the Internet through a proxy, you will need the details of the proxy including the fully qualified domain name and the port number.

Enable vSphere High Availability (HA)

You must enable vSphere High Availability (HA) to support virtual machine failover during a HA event such as a host failure. Sufficient CPU and memory resources must be reserved across the system so that all VMs on the affected host(s) can fail over to remaining available hosts in the system. You configure an Admission Control Policy (ACP) to specify the percentage CPU and memory to reserve on all the hosts in the cluster to support HA functionality.



Note

You should not use the default Admission Control Policy. Instead, you should calculate the memory and CPU requirements that are specific to your environment.

Install vSphere Docker Volume Service driver on all ESXi hosts

vSphere Docker Volume Service technology enables stateful containers to access the storage volumes. Setting this up is a one-off manual step. In order to be able to use Docker volumes using the vSphere driver, you must first install the latest release of the vSphere Docker Volume Service (vDVS) driver, which is available as a vSphere Installation Bundle (VIB). To perform this operation, log in to each of the ESXi hosts and then download and install the latest release of vDVS driver.

```
# esxcli software vib install -v /tmp/vmware-esx-vmrkops-<version>.vib --no-sig-check
```

More information on how to download and install the driver can be found at <http://vmware.github.io/vsphere-storage-for-docker/documentation/install.html>. The version of the driver tested in this configuration is 0.21.2.

Create the Ansible node

The Ansible node will act as the driver to automate the provisioning of the environment and it is essential that it is properly installed.

1. Create a Virtual Machine and install your preferred OS (in this example, and for the sake of simplicity, RHEL7 will be used). The rest of the instructions assume that, if you use a different OS, you understand the possible differences in syntax for the provided commands. If you use RHEL 7, select **Infrastructure Server** as the base environment and the **Guests Agents** add-on during the installation.
2. Log in to the `root` account and create an SSH key pair. Do not protect the key with a passphrase (unless you want to use `ssh-agent`).

```
# ssh-keygen
```

3. Configure the following yum repositories, `rhel-7-server-rpms` and `rhel-7-server-extras-rpms` as explained in the section [Configure the yum repositories](#). The "extras" repo can be enabled as follows:

```
# subscription-manager repos --enable=rhel-7-server-extras-rpms
```

4. Configure the EPEL repository. For more information, see: <http://fedoraproject.org/wiki/EPEL>. Note that `yum-config-manager` comes with the Infrastructure Server base environment. If you did not select this environment, you will have to install the `yum-utils` package.

```
# rpm -ivh https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm
# yum-config-manager --enable rhel-7-server-extras-rpms
```

5. Install Ansible 2.7 or higher.

```
# yum install ansible
```

6. Install the following packages which are a mandatory requirement for the playbooks to function as expected. (Update `pip` if requested).

```
# yum install python-pymomi python-netaddr python2-jmespath python-pip gcc python-devel openssl-devel git
# pip install --upgrade pip
# pip install cryptography
# pip install pysphere
# pip install --ignore-installed "pywinrm>0.2.2"
```

Configure the yum repositories

The Red Hat packages required during the deployment of the solution come from two repositories: `rhel-7-server-rpms` and `rhel-7-server-extras-rpms`. The first repository is on the Red Hat DVD but the second is not. There are two options, with both options requiring a Red Hat Network account. Logon to your VM template using SSH with the credentials you configured for the root account and then implement one of the two options below:

Option 1: Use Red Hat subscription manager to register your system. This is the easiest way and will automatically give you access to the official Red Hat repositories.

1. Use the `subscription-manager register` command as follows.



```
# subscription-manager register --auto-attach
```

2. If you are behind a proxy, you must configure this before running the above command to register.

```
# subscription-manager config --server.proxy_hostname=<proxy IP> --server.proxy_port=<proxy port>
```

3. Verify that you don't have the issue described here: <https://access.redhat.com/solutions/3317671> by entering the following command.

```
# yum repolist
```

4. If you have the issue, fix it with the following command

```
# subscription-manager repos --disable=rhel-7-server-rt-beta-rpms
```

The playbooks will later automatically enable the **extras** repository on the VMs that need it.

Option 2: Use an internal repository. Instead of pulling the packages from Red Hat, you can create copies of the required repositories on a dedicated node. You can then configure the package manager to pull the packages from the dedicated node. Your `/etc/yum.repos.d/redhat.repo` could look as follows.

```
[RHEL7-Server]
name=Red Hat Enterprise Linux $releasever - $basearch
baseurl=http://yourserver.example.com/rhel-7-server-rpms/
enabled=1
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
```

```
[RHEL7-Server-extras]
name=Red Hat Enterprise Linux Extra pkg $releasever - $basearch
baseurl=http://yourserver.example.com/rhel-7-server-extras-rpms/
enabled=1
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
```

To see how you can create a local mirror of the Red Hat repositories and how to share them, check the Red Hat documentation at <https://access.redhat.com/solutions/23016>, <https://access.redhat.com/solutions/265523> and at <https://access.redhat.com/solutions/7227>.

Create the Red Hat Linux template

To create the Red Hat Linux VM template that you will use as the base for all your nodes, you first create a Virtual Machine with the OS installed and then convert the Virtual Machine to a VM Template. The VM Template is created as lean as possible, with any additional software installs and/or system configuration performed subsequently using Ansible.

As the creation of the template is a one-off task, this procedure has not been automated. The steps required to manually create a VM template are outlined below.

Log in to vCenter and create a new Virtual Machine with the following characteristics:

- Guest OS Family: Linux, Guest OS Version: Red Hat Enterprise Linux (64-bit)
- Hard Disk size: 50GB, (Thin provisioning)
- A single network controller connected to the network or VLAN of your choice. All VMs will connect to this same network.
- Optionally you can remove the floppy drive

Install Red Hat Enterprise 7:

1. Select a language which is supported by Docker



2. For the software selection, choose **Infrastructure Server** as the base environment and add the **Guest Agents** from the lists of add-ons available for this environment. The Infrastructure Server environment is selected here versus the Minimal Install because Customization of Linux guest operating systems requires that Perl is installed in the Linux guest operating system.
3. Configure the network settings so that you can later access the VM using SSH. Specify an IP address for the network interface, a default gateway, DNS settings and possibly any HTTP/HTTPS proxies that apply in your environment.
4. Specify a password for the root account and optionally created an admin user.
5. Wait for the installation to finish and for the VM to reboot.

Finalize the template

Log in to the `root` account on the Ansible box and copy the SSH public key to the VM Template. This will allow your Ansible node to SSH to all the Virtual Machines created from the VM Template without the need for a password.

```
ssh-copy-id root@<IP of your VM_Template>
```

Perform the following steps on the VM Template to finalize its creation:

1. Clean up the template by running the following commands from the **Virtual Machine Console**:

```
# rm /etc/ssh/ssh_host_*
# nmcli con del ens192
# logrotate -f /etc/logrotate.conf
# rm /var/log/*-201?*
# history -c
```

2. Shutdown the VM

```
# shutdown -h now
```

3. Turn the VM into a template by right-clicking on your VM and selecting **Template -> Convert to Template**. This will create a new template visible under VM Templates in Folders, ready for future use.

Note

In both the Ansible node and the VM Template, you might need to configure the network so one node can reach the other. Instructions for this step have been omitted since it is a basic step and could vary depending on the user's environment.

Configuring the solution components

Once you have prepared your environment, you need to download the solution software and edit the configuration variables to match your setup.

Ansible configuration

1. On the Ansible node, retrieve the latest version of the playbooks using Git.

```
# git clone https://github.com/HewlettPackard/Docker-SimpliVity.git
```

2. Change to the directory that you just cloned:

```
# cd ~/Docker-SimpliVity
```

Note

All subsequent file names are relative to the `Docker-SimpliVity` directory. For example `vm_hosts` is located in `~/Docker-SimpliVity/group_vars/vars` corresponds to `~/Docker-SimpliVity/group_vars/vars`.



You now need to prepare the configuration to match your own environment, prior to deploying Docker EE and the rest of the nodes. To do so, you will need to modify a number of files including:

- `site.yml`, the main entry point for the playbooks.
- `vm_hosts`, the inventory file.

You also need to create and populate a number of files:

- `group_vars/vars`, the group variables file.
- `group_vars/vault`, containing sensitive information that needs to be protected.
- `group_vars/backups`, containing backup-related variables.

For the latter group, a set of sample files has been provided to help you get started:

- `group_vars/vars.sample`, a sample group variables file.
- `group_vars/vault.sample`, a sample vault file.
- `group_vars/backups.sample`, a sample backup configuration file.

The file `group_vars/win_worker.yml` supports advanced configuration of Windows remote management and in general should not require modification.

You should work from the `root` account for the configuration steps and also later on when you run the playbooks.

Editing the inventory

The inventory is the file named `vm_hosts` in the `~/Docker-Simplivity` directory. You need to edit this file to describe the configuration you want to deploy.

The nodes inside the inventory are organized in groups. The groups are defined by brackets and the group names are static so they must not be changed. Other fields (hostnames, specifications, IP addresses...) are edited to match your setup. The groups are as follows:

- `[ucp_main]`: A group containing one single node which will be the main UCP node and swarm leader. Do not add more than one node under this group.
- `[ucp]`: A group containing all the UCP nodes, including the main UCP node. Typically you should have either 3 or 5 nodes under this group.
- `[dtr_main]`: A group containing one single node which will be the first DTR node to be installed. Do not add more than one node under this group.
- `[dtr]`: A group containing all the DTR nodes, including the main DTR node. Typically you should have either 3 or 5 nodes under this group.
- `[worker]`: A group containing all the Linux worker nodes.
- `[win_worker]`: A group containing all the Windows worker nodes.
- `[nfs]`: A group containing one single node which will be the NFS node. Do not add more than one node under this group.
- `[logger]`: A group containing one single node which will be the logger node. Do not add more than one node under this group.
- `[local]`: A group containing the local Ansible host. It contains an entry that should not be modified.

If you are deploying the new active-active load balancers, using floating IPs managed by **keepalived**:

- `[loadbalancer]`: A group containing the UCP, DTR and any worker load balancers you are deploying.

If you are using the legacy, standalone load balancers:

- `[ucp_lb]`: A group containing one single node which will be the load balancer for the UCP nodes. Do not add more than one node under this group.



- **[dtr_lb]**: A group containing one single node which will be the load balancer for the DTR nodes. Do not add more than one node under this group.
- **[worker_lb]**: A group containing one single node which will be the load balancer for the worker nodes. Do not add more than one node under this group.
- **[lbs]**: A group containing all the load balancers. This group will have 3 nodes, also defined individually in the three groups above.

There are also a few special groups:

- **[docker:children]**: A group of groups including all the nodes where Docker will be installed.
- **[vms:children]**: A group of groups including all the Virtual Machines involved, with the exception of the local host.

Finally, you will find some variables defined for each group:

- **[vms:vars]**: A set of variables defined for all VMs. Currently only the size of the boot disk is defined here.
- **[ucp:vars]**: A set of variables defined for all nodes in the **[ucp]** group.
- **[dtr:vars]**: A set of variables defined for all nodes in the **[dtr]** group.
- **[worker:vars]**: A set of variables defined for all nodes in the **[worker]** group.
- **[win_worker:vars]**: A set of variables defined for all nodes in the **[win_worker]** group.
- **[loadbalancer:vars]**: A set of variables defined for all nodes in the **[loadbalancer]** group.
- **[lbs:vars]**: A set of variables defined for all nodes in the **[lbs]** group.
- **[nfs:vars]**: A set of variables defined for all nodes in the **[nfs]** group.
- **[logger:vars]**: A set of variables defined for all nodes in the **[logger]** group.

If you wish to configure your nodes with different specifications to the ones defined by the group, it is possible to declare the same variables at the node level, overriding the group value. For instance, you could have one of your Linux workers with higher specifications by setting:

```
[worker]
worker01 ip_addr='10.0.0.10/16' esxi_host='esxi1.domain.local'
worker02 ip_addr='10.0.0.11/16' esxi_host='esxi1.domain.local'
worker03 ip_addr='10.0.0.12/16' esxi_host='esxi1.domain.local' cpus='16' ram='32768'

[worker:vars]
cpus='4' ram='16384' disk2_size='200'
```

In the example above, the **worker03** node would have 4 times more CPU and double the RAM compared to the rest of the worker nodes.

The different variables you can use are described in Table 7 below. They are all mandatory unless otherwise specified.

Table 7. Variables

Variable	Scope	Description
ip_addr	Node	IP address in CIDR format to be given to a node
esxi_host	Node	ESXi host where the node will be deployed. If the cluster is configured with DRS, this option will be overridden
cpus	Node/Group	Number of CPUs to assign to a VM or a group of VMs
ram	Node/Group	Amount of RAM in MB to assign to a VM or a group of VMs
disk2_usage	Node/Group	Size of the second disk in GB to attach to a VM or a group of VMs. This variable is only mandatory on Docker nodes (UCP, DTR, worker) and NFS node. It is not required for the logger node or the load balancers.
node_policy	Node/Group	HPE SimpliVity backup policy to assign to a VM or a group of VMs. The name has to match one of the backup policies defined in the group_vars/vars file described in the section HPE SimpliVity backup configuration.



VMware configuration

All VMware-related variables are mandatory and are described in Table 8.

Table 8. VMware variables

Variable	File	Description
vcenter_hostname	group_vars/vars	IP or hostname of the vCenter appliance
vcenter_username	group_vars/vars	Username to log in to the vCenter appliance. It might include a domain, for example, 'administrator@vsphere.local'.
vcenter_password	group_vars/vault	The password corresponding to the <code>vcenter_username</code> user above.
vcenter_validate_certs	group_vars/vars	'no'
datacenter	group_vars/vars	Name of the datacenter where the environment will be provisioned
vm_username	group_vars/vars	Username to log into the VMs. It needs to match the one from the VM Template, so unless you have created a user, you must use 'root'.
vm_password	group_vars/vault	The password for the <code>vm_username</code> user above.
vm_template	group_vars/vars	Name of the RHEL VM Template to be use. Note that this is the name from a vCenter perspective, not the hostname.
folder_name	group_vars/vars	vCenter folder to deploy the VMs. If you do not wish to deploy in a particular folder, the value should be /. Note: If you want to deploy in a specific folder, you need to create this folder in the inventory of the selected datacenter before starting the deployment.
datastores	group_vars/vars	List of datastores to be used, in list format, i.e. [<code>Datastore1</code> ; <code>Datastore2</code> ...]. The datastores must exist before you run the playbooks. Note that each datastore should be mounted on each of the ESXi hosts. Please note that from a HPE SimpliVity perspective, it is a best practice to use only one Datastore. Using more than one will not provide any advantages in terms of reliability and will add additional complexity.
disk2	group_vars/vars	UNIX® name of the second disk for the Docker VMs. Typically <code>/dev/sdb</code>
disk2_part	group_vars/vars	UNIX name of the partition of the second disk for the Docker VMs. Typically <code>/dev/sdb1</code>
vsphere_plugin_version	group_vars/vars	Version of the vSphere plugin for Docker. The default is 0.21.2 which is the latest version at the time of writing this document. The version of the plugin should match the version of the vSphere Installation Bundle (VIB) that you installed on the ESXi servers.
vm_portgroup	group_vars/vars	Used by the playbook <code>create_vms.yml</code> , this variable is used to specify the portgroup connected to the network that connects all the VMs. There is currently only one network. It is recommended that the template which is used as the base for all deployed VMs specifies a network adapter but it is not required. If a network adapter is specified, you should not attach this adapter to a standard switch if the portgroup designated by <code>vm_portgroup</code> is connected to a distributed vSwitch. In addition, you should make sure that the adapter specifies <code>Connect At Power On</code> .

HPE SimpliVity configuration

Variables related to your HPE SimpliVity deployment are mandatory and are described in Table 9.

Table 9. SimpliVity variables

Variable	File	Description
simplivity_username	group_vars/vars	Username to log in to the SimpliVity Omnistack appliances. It might include a domain, for example, <code>administrator@vsphere.local</code> . Note: The corresponding password is stored in the variable named <code>simplivity_password</code> .
simplivity_password	group_vars/vault	The password for the <code>simplivity_username</code> user.
omnistack_ovc	group_vars/vars	List of Omnistack hosts to be used, in list format, i.e. [<code>omni1.local</code> ; <code>omni2.local</code> ...]. If your OmniStack virtual machines do not have their names registered in DNS, you can use their IP addresses.



VM placement and number of HPE SimpliVity servers in the cluster

The placement of the various VMs deployed by the playbooks depends on whether DRS is enabled or not:

1. If DRS is not enabled, the placement of the VMs is specified in the Ansible inventory file `vm_hosts`
2. If DRS is enabled, the placement of the VMs is outside the control of the playbooks

The playbooks have only been tested with three nodes in the ESX cluster, but the following sections provide guidance on how to use more than three nodes.

Using more than three nodes when DRS is not enabled

The default `vm_hosts` file in the solution GitHub repository corresponds to a deployment on a 3-node HPE SimpliVity cluster. For each Ansible host in the inventory, you use the `esxi_hosts` variable to specify on which ESX hosts the VM should be placed. The following code extract shows 3 UCP VMs distributed across the three members of the cluster. This is the recommended placement as you do not want one node to host two UCP VMs since failure of that node would result in the cluster losing quorum.

```
[ucp]
hpe-ucp01 ip_addr='10.10.174.112/22' esxi_host='simply01.am2.cloudra.local'
hpe-ucp02 ip_addr='10.10.174.113/22' esxi_host='simply02.am2.cloudra.local'
hpe-ucp03 ip_addr='10.10.174.114/22' esxi_host='simply03.am2.cloudra.local'
```

In the above example, the first UCP VM will be placed on the ESX host named `simply01.am2.cloudra.local`. Note that the value for `esxi_host` is the name of the ESX host in the vCenter inventory.

The default `vm_hosts` inventory configures three Docker worker nodes and distributes them across the three ESX hosts:

```
[worker]
hpe-worker01 ip_addr='10.10.174.122/22' esxi_host='simply01.am2.cloudra.local'
hpe-worker02 ip_addr='10.10.174.123/22' esxi_host='simply02.am2.cloudra.local'
hpe-worker03 ip_addr='10.10.174.124/22' esxi_host='simply03.am2.cloudra.local'
```

If you have more than three ESX hosts in your cluster, you can add an additional worker node as follows:

```
[worker]
hpe-worker01 ip_addr='10.10.174.122/22' esxi_host='simply01.am2.cloudra.local'
hpe-worker02 ip_addr='10.10.174.123/22' esxi_host='simply02.am2.cloudra.local'
hpe-worker03 ip_addr='10.10.174.124/22' esxi_host='simply03.am2.cloudra.local'
hpe-worker04 ip_addr='10.10.174.150/22' esxi_host='simply04.am2.cloudra.local'
```

You can also distribute the infrastructure VMs across four nodes rather than across the default nodes. For example, the default placement for the NFS server VM is as follows:

```
[nfs]
hpe-nfs ip_addr='10.10.174.121/22' esxi_host='simply03.am2.cloudra.local'
```

Instead, you can change the placement NFS server VM, leveraging a fourth ESX node:

```
[nfs]
hpe-nfs ip_addr='10.10.174.121/22' esxi_host='simply04.am2.cloudra.local'
```

When you specify the placement of the VM, you should ensure that you follow these placement guidelines:

- Do not place two UCP VMs on the same node. If the node fails, the UCP cluster will lose quorum and the service will go down.
- Do not place two DTR replicas (VMs) on the same node. Once again, the cluster will lose quorum if that node fails.



Note

The OmniStack software maintains two replicas on two different hosts for each VM. As a result, when a VM is scheduled on an ESX server that does not have local access to one of the replicas, the VM will report the warning “SimpliVity VM Data Access Not Optimized”. You can safely ignore this warning.

Using more than three nodes when DRS is enabled

When DRS is enabled, it controls the placement of the VMs and as a result, the placement you have specified in the `vm_hosts` inventory is ignored. Instead, you use DRS rules to make sure that the UCP and DTR VMs are distributed across three nodes for the reasons explained earlier.

Warning

If you do not specify DRS rules to determine the placement, DRS will automatically move the VMs that report the “SimpliVity VM Data Access Not Optimized” warning to a node with a replica of the VM which may break the earlier placement guideline.

HPE SimpliVity backup configuration

Variables related to HPE SimpliVity backups are described in Table 10.

Table 10. HPE SimpliVity backup variables

Variable	Description
<code>backup_policies</code>	<p>List of dictionaries containing the different backup policies to be used along with the scheduling information. Any number of backup policies can be created and they need to match the <code>node_policy</code> variables defined in the inventory. Times are indicated in minutes. All month calculations use a 30-day month. All year calculations use a 365-day year. The format is as follows:</p> <pre>backup_policies: - name: daily' days: 'All' start_time: '11:30' frequency: '1440' retention: '43200' - name: 'hourly' days: 'All' start_time: '00:00' frequency: '60' retention: '2880'</pre>
<code>dummy_vm_prefix</code>	In order to be able to back up the Docker volumes, a number of “dummy” VMs need to spin up. This variable will set a recognizable prefix for them.
<code>docker_volumes_policy</code>	Backup policy to use for the Docker Volumes.

Networking configuration

All network-related variables are mandatory and are described in Table 11.

Table 11. Network variables

Variable	File	Description
<code>nic_name</code>	<code>group_vars/vars</code>	Name of the device, for RHEL this is typically <code>ens192</code> and it is recommended to leave it as is.
<code>gateway</code>	<code>group_vars/vars</code>	IP address of the gateway to be used
<code>dns</code>	<code>group_vars/vars</code>	List of DNS servers to be used, in list format, i.e. <code>[10.10.173.1;10.10.173.2;...]</code>



Variable	File	Description
domain_name	group_vars/vars	Domain name for your Virtual Machines
ntp_server	group_vars/vars	List of NTP servers to be used, in list format, i.e. ['1.2.3.4','0.us.pool.net.org'...]

Environment configuration

All Environment-related variables are described in Table 12 below.

Table 12. Environment variables

Variable	File	Description
env	group_vars/vars	<p>Dictionary containing all environment variables. It contains three entries described below. Please leave the proxy related settings empty if not required:</p> <p>http_proxy: HTTP proxy URL, such as 'http://15.184.4.2:8080'. This variable defines the HTTP proxy URL if your environment is behind a proxy.</p> <p>https_proxy: HTTPS proxy URL, such as 'http://15.184.4.2:8080'. This variable defines the HTTPS proxy URL if your environment is behind a proxy.</p> <p>no_proxy: List of hostnames or IPs that don't require proxy, such as 'localhost,127.0.0.1,.cloudra.local,10.10.174.'</p>

Docker configuration

All Docker-related variables are mandatory and are described in Table 13.

Table 13. Docker variables

Variable	File	Description
docker_ee_url	group_vars/vault	Note: This is a private link to your Docker EE subscription. The value for docker_ee_url is the URL documented at the following address: https://docs.docker.com/engine/installation/linux/docker-ee/rhel/ .
docker_ee_reponame	group_vars/vars	For Docker EE 2.1, this variable must be set to the value stable-18.09
rhel_version	group_vars/vars	For the Docker installation, this sets the version of your RHEL OS, such as 7.6 . The playbooks were tested with RHEL 7.6.
dtr_version	group_vars/vars	Version of the Docker DTR you wish to install. You can use a numeric version or latest for the most recent one. The playbooks were tested with 2.6.2.
ucp_version	group_vars/vars	Version of the Docker UCP you wish to install. You can use a numeric version or latest for the most recent one. The playbooks were tested with UCP 3.1.3.
images_folder	group_vars/vars	Directory in the NFS server that will be mounted in the DTR nodes and that will host your Docker images.
license_file	group_vars/vars	Full path to your Docker EE license file on your Ansible host. The license file is available from the Docker Store
ucp_username	group_vars/vars	Username of the administrator user for UCP and DTR, typically admin .
ucp_password	group_vars/vault	The password for the ucp_username account.
docker_storage_driver	group_vars/vars	Storage driver for Docker nodes. Accepted values are overlay2 (the default) and devicemapper .

To see how to use customer-supplied certificates with UCP and DTR, see Appendix B.

Orchestrator configuration

The variable **orchestrator** in the **[worker]** group is used to specify if a worker node should be assigned to the Kubernetes orchestrator (**orchestrator: 'kubernetes'**) or to the swarm orchestrator (**orchestrator: 'swarm'**). In general, you should only change the orchestrator for worker nodes.



Note

Docker supports a third type, **mixed**, that enables workloads to be scheduled by both Kubernetes and Docker swarm on the same node. Mixing orchestrator types on the same node is not recommended for production deployments because of the likelihood of resource contention. As a result, these playbooks do not support the **mixed** type.

The following example shows how to set Kubernetes as the default orchestrator for worker nodes, and how to override the default to use Docker swarm on one specific node instead.

```
## WORKER
[worker]
hpe-worker01 ip_addr='10.60.59.21/16' esxi_host='esxi-hpe-1.cloudra.local'
hpe-worker02 ip_addr='10.60.59.22/16' esxi_host='esxi-hpe-2.cloudra.local'
hpe-worker03 ip_addr='10.60.59.23/16' esxi_host='esxi-hpe-3.cloudra.local' orchestrator=swarm

[worker:vars]
cpus='4'
ram='65536'
disk2_size='500'
orchestrator=kubernetes
```

Note

The playbooks do not change Docker's default orchestrator type which is **swarm**. Instead, the inventory is used to configure worker nodes for Kubernetes workloads or swarm workloads as explained above. If you want to change the default orchestrator type, use the method explained in the Docker documentation at <https://docs.docker.com/ee/ucp/admin/configure/set-orchestrator-type/#set-the-default-orchestrator-type-for-new-nodes>.

It is possible to manually change the orchestrator type for a node. When you do this, existing workloads are evicted and they are not migrated automatically to the new orchestrator. If you want the workloads to be scheduled by the new orchestrator, you must migrate them manually. More information is available in the Docker documentation at <https://docs.docker.com/ee/ucp/admin/configure/set-orchestrator-type/#what-happens-when-you-change-a-nodes-orchestrator>.

Kubernetes configuration

The current playbooks support the deployment of UCP 3.1.* which deploys Kubernetes version 1.11.*. This version of the playbooks will not work with a version of UCP that is lower than 3. If you wish to deploy using UCP 2.*, you will need to download a previous release of the playbooks, which is available on the GitHub site.

The preceding section [Orchestrator configuration](#) explains how to assign a worker node to the Kubernetes orchestrator. This section covers specific Kubernetes configuration, including how to set the pod CIDR and how to configure Kubernetes Persistent Volumes.

Pod CIDR

The variable `k8s_pod_cidr` is specified in `group_vars/vars` and configures a custom range of IP addresses to be used by pods. The specific range that you use should be dedicated to the cluster.

The default value is `192.168.0.0/16`. To set an alternative value, use the variable as shown in the example:

```
k8s_pod_cidr: 192.168.128.0/17
```

Kubernetes Persistent Volume configuration

Variables related to the configuration of Kubernetes Persistent Volumes are shown in



Table 14.



Table 14. Kubernetes Persistent Volume variables

Variable	File	Description
nfs_provisioner_namespace	group_vars/vars	The Kubernetes namespace, for example, <code>nfsstorage</code>
nfs_provisioner_role	group_vars/vars	Name of the role to create, for example, <code>nfs-provisioner-runner</code> .
nfs_provisioner_serviceaccount	group_vars/vars	The Kubernetes service account name to use for RBAC purposes, for example, <code>nfs-provisioner</code>
nfs_provisioner_name	group_vars/vars	Name of the provisioner, for example, <code>hpe.com/nfs</code>
nfs_provisioner_storage_class_name	group_vars/vars	Name of the storage class to create, for example, <code>nfs</code>
nfs_provisioner_server_ip	group_vars/vars	IP address (or FQDN) of your external NFS server, for example, <code>hpe2-nfs.cloudra.local</code>
nfs_provisioner_server_share	group_vars/vars	Name of the NFS share where all the persistent volume data will be stored, for example, <code>/k8s</code>

Related playbooks

The playbook `playbooks/k8s-nfs-provisioner.yml` is used to enable a dynamic NFS provisioner which can be used to automatically create and allocate Kubernetes persistent volumes. The backend storage is provided by an NFS backend. This playbook is run from the Ansible box after downloading a UCP client bundle for the `admin` account and sourcing the downloaded `env.sh` file. For more information on using this playbook, see the section Deploying the NFS provisioner for Kubernetes.

Protecting sensitive information

A vault file is used to protect any sensitive variables that should not appear in clear text in your `group_vars/vars` file. The vault file will be encrypted and will require a password to be entered before it can be read or updated.

A sample vault file is provided named `group_vars/vault.sample` that you can use as a model for your vault file. To create a vault, you create a new file called `group_vars/vault` and add entries similar to:

```
---
docker_ee_url: 'your_url_here'
vcenter_password: 'xxxx'
vm_password: 'xxxx'
ucp_password: 'zzzz'
win_password: 'yourpass'
sysdig_access_key: 'enter-sysdig-access-key'
rhn_orgid: "YourOrgId"
rhn_key: "YourActivationKey"
#password for the splunk universal forwarder. Must meet password complexity requirement
splunk_uf_password: 'YourPa$$word12'
```

`rhn_orgid` and `rhn_key` are the credentials needed to subscribe the virtual machines with Red Hat Customer Portal. For more information regarding activation keys, see the following URL: <https://access.redhat.com/articles/1378093>

To encrypt the vault you need to run the following command:

```
# ansible-vault encrypt group_vars/vault
```

You will be prompted for a password that will decrypt the vault when required. You can update the values in your vault by running:

```
# ansible-vault edit group_vars/vault
```

In order for Ansible to be able to read the vault, you need to specify a file where the password is stored, for instance, in a file called `.vault_pass`. Once the file is created, take the following precautions to avoid illegitimate access to this file:

- Change the permissions so only `root` can read it using `# chmod 600 .vault_pass`
- Add the file to your `.gitignore` file if you are using a Git repository to manage your playbooks.



Inventory group variables

Additional configuration files for each group in the inventory are available, including `group_vars/vms.yml`, `group_vars/ucp.yml`, `group_vars/dtr.yml`, `group_vars/worker.yml` and `group_vars/nfs.yml`.

These group files facilitate more sophisticated settings, such as additional drives and additional network interfaces. For example, here is the `group_vars/nfs.yml` file.

```
networks:
  - name: '{{ vm_portgroup }}'
    ip: "{{ ip_addr | ipaddr('address') }}"
    netmask: "{{ ip_addr | ipaddr('netmask') }}"
    gateway: "{{ gateway }}"

disks_specs:
  - size_gb: '{{ disk1_size }}'
    type: thin
    datastore: "{{ datastores | random }}"
  - size_gb: '{{ disk2_size }}'
    type: thin
    datastore: "{{ datastores | random }}"
  - size_gb: 10
    type: thin
    datastore: "{{ datastores | random }}"
```

In this example, the size of the first two drives is specified using the values of the variables `disk1_size` and `disk2_size` that are declared in the `group_vars/vars` file. This maintains compatibility with `vm_hosts` inventories from the previous release of the playbooks. However, it is possible to provide explicit values, depending on your requirements, for the individual UCP, DTR, worker or NFS VMs. For example, you may want to increase the size of the second disk for the NFS VM as this is used to store the DTR images, so the default value of 500GB may not be sufficient to meet your needs.

In this release, support has been added for configuring a third drive that can be used to hold Kubernetes persistent volume data. The default size (10GB) is set low as the use of the NFS VM for storing persistent volume data is only considered suitable for demo purposes and should not be used in a production environment.

In the following example, the `group_vars/nfs.yml` has been modified to configure the NFS VM with a 50GB boot disk, a 500GB drive for DTR images and an 800GB drive for Kubernetes persistent volumes data.

```
networks:
  - name: '{{ vm_portgroup }}'
    ip: "{{ ip_addr | ipaddr('address') }}"
    netmask: "{{ ip_addr | ipaddr('netmask') }}"
    gateway: "{{ gateway }}"

disks_specs:
  - size_gb: 50
    type: thin
    datastore: "{{ datastores | random }}"
  - size_gb: 500
    type: thin
    datastore: "{{ datastores | random }}"
  - size_gb: 800
    type: thin
    datastore: "{{ datastores | random }}"
```



Note

The number of drives and the purpose of each drive is determined by the role of the VM and the specific playbooks that use the information. The first disk is always used as the boot disk, irrespective of VM role, while the purpose of the second or third disk is specific to the role.

Overview of the playbooks

The Ansible playbooks are available to download at <https://github.com/HewlettPackard/Docker-SimpliVity>. Once you have cloned the repository, change directory to `/root/Docker-SimpliVity`.

You can use the playbook `site.yml` as the day 0 playbook to deploy the solution. It is simply a wrapper around a number of required and optional playbooks that allow you to configure the deployment to your needs.

To start a deployment, use the following command:

```
# ansible-playbook -i vm_hosts site.yml --vault-password-file .vault_pass
```

The playbooks should run for approximately 35-40 minutes for the default deployment with 3 UCP, 3 DTR and 3 Linux worker nodes (depending on your server specifications and the size of your environment).

Core components

The playbooks for deploying the core components are described in the following sections:

- Provisioning RHEL VMs
- Provisioning load balancers for UCP and DTR
- Installing Docker UCP and DTR on RHEL VMs
- Deploying RHEL workers

Optional components

The playbooks for deploying optional components are described in the following sections:

- Playbooks for adding Windows workers
- Playbooks for installing Sysdig on RHEL
- Playbooks for installing Splunk
- Playbooks for installing Prometheus and Grafana

Backup and restore playbooks

Best practices and procedures are described in the section [Backup and restore](#). The following playbooks are used to perform backups:

- `playbooks/backup_swarm.yml` is used to back up the swarm data
- `playbooks/backup_ucp.yml` is used to back up UCP
- `playbooks/backup_dtr_meta.yml` is used to back up DTR metadata
- `playbooks/backup_dtr_images.yml` is used to back up DTR images

The following playbooks are used to restore the system:

- `playbooks/restore_dtr_images.yml` is used to restore DTR images
- `playbooks/restore_dtr_metadata.yml` is used to restore DTR metadata
- `playbooks/restore_ucp.yml` is used to restore UCP



Convenience playbooks

- `playbooks/clean_all.yml` powers off and deletes all VMs in your inventory.
- `playbooks/distribute_keys.yml` distributes public keys between all nodes, to allow each node to password-less log in to every other node. As this is not essential and can be regarded as a security risk (a worker node probably should not be able to log in to a UCP node, for instance), this playbook is commented out in `site.yml` by default.

Convenience scripts

- `backup.sh` can be used to take a backup of the swarm, UCP, DTR metadata and the DTR images in one go.
- `restore_dtr.sh` can be used to restore DTR metadata and DTR images.
- `scale_worker.sh` can be used to scale the worker nodes.

Deploying the core components

At this point, the system is ready to be deployed. Make sure you are logged on as `root` in your ansible box and that your current directory is `/root/Docker-SimpliVity`

Note

As well as configuring your `vars` and `vault` files, you must also provide a `backups` configuration file in the `group_vars` folder when running `site.yml` or `hybrid.yml`. An example file is provided in the repository named `backups.sample`. Rename it to `backups` before running the playbooks. Details on how to configure this file are available in the section [Backup and restore](#).

Provisioning RHEL VMs

The following playbooks are used to provision RHEL VMs:

- `playbooks/create_vms.yml` will create all the necessary virtual machines for the environment from the VM Template defined in the `vm_template` variable. All Linux VMs are now created in one go, regardless of the number of drives they have. This playbook also has the potential to configure additional network adapters.
- `playbooks/config_networking.yml` will configure the network settings in all the virtual machines.
- `playbooks/resize_syspart.yml` resizes the logical volume that holds the `/` partition of the Linux VMs to use all the space available on the drive.
- `playbooks/config_subscription.yml` registers and subscribes all virtual machines to the Red Hat Customer Portal. This is only needed if you pull packages from Red Hat. This playbook is commented out by default but you should uncomment it to make sure each VM registers with the Red Hat portal. It is commented out so that you can test the deployment first without having to unregister all the VMs from the Red Hat Customer Portal between each test. If you are using an internal repository, as described in the section [Create the Red Hat Linux template](#), you can keep this playbook commented out.
- `playbooks/config_ntp.yml` configures the **chrony** client package in all virtual machines in order to have a synchronized clock across the environment. It will use the list of servers specified in the `ntp_servers` variable in the file `group_vars/vars`.

Provisioning load balancers for UCP and DTR

The playbook `playbooks/loadbalancer.yml` is used to deploy load balancers in an **active-active** configuration to provide highly-available access to UCP and DTR.

At least two nodes are specified in the `[loadbalancer]` group in the inventory, along with group variables defining CPU and RAM requirements. These nodes run `keepalived` and `HAProxy`.

```
[loadbalancer]
hpe-lb1 ip_addr='10.10.174.248/22' esxi_host='simply04.am2.cloudra.local' ucp=true
hpe-lb2 ip_addr='10.10.174.249/22' esxi_host='simply05.am2.cloudra.local' dtr=true
```



```
[loadbalancer:vars]
cpus='2'
ram='4096'
node_policy='hpe-bronze'
```

The virtual IP for UCP will be handled by **hpe-lb1** by default, which will split the traffic across the three UCP VMs **hpe-ucp01**, **hpe-ucp02** and **hpe-ucp03**. In the case of a failure of **hpe-lb1**, the virtual IP for UCP will automatically move to the second load balancer node **hpe-lb2** which will again distribute the traffic to the UCP VMs.

Similarly, the virtual IP for DTR will be handled by default by the load balancer **hpe-lb2**, splitting the traffic across the three DTR VMs **hpe-dtr01**, **hpe-dtr02** and **hpe-dtr03**. In the case of a failure of **hpe-lb2**, the virtual IP for DTR will automatically move to the first load balancer node **hpe-lb1** which will again distribute the traffic to the DTR VMs.

To configure the virtual IPs for UCP and DTR, you need to add a **loadbalancers** dictionary to your **group_vars/vars** file as shown in the excerpt below:

```
loadbalancers:
  ucp:
    public_interface: 'ens192'
    public_vip: '10.60.59.251'
    public_fqdn: hpe-ucpvip.cloudra.local
    virtual_router_id: 54
  dtr:
    public_interface: 'ens192'
    public_vip: '10.60.59.252'
    public_fqdn: hpe-dtrvip.cloudra.local
    virtual_router_id: 55
```

Warning

If you re-run **playbooks/loadbalancer.yml** after a configuration change, you may need to subsequently run **playbooks/reconfigure_dtr.yml** as the latter playbook configures the virtual IP address for accessing the UCP Single-Sign-On (SSO) page. If there is no virtual IP or FQDN defined for UCP in the variables file, the playbook will choose the address of the first UCP node in the **[ucp]** group. This scenario introduces a single point of failure and should be avoided.

Note

By default, the playbook supports ports **443** and **6443** for UCP and port **433** for DTR. If you deploy Prometheus and Grafana on Docker Swarm, the Grafana port **3000** will be handled as well.

Note

The playbook **playbooks/loadbalancer.yml** can be used to create one or more load balancers for applications running on your worker nodes. However, it is impossible for the playbooks to know what ports to support, so manual configuration of HAProxy and **keepalived** may be required. By default, the playbooks support ports **80** and **443** for worker nodes

Legacy stand-alone load balancers

The playbook **playbooks/install_haproxy.yml** is used to deploy three separate load balancers, for the UCP, DTR and worker nodes. It is recommended that you use the HAProxy and **keepalived** solution documented above instead of this option.

Deploying without load balancers

If you do not want to deploy load balancers when running **site.yml**, you should comment out any declarations in the inventory and variables files. This includes any legacy stand-alone load balancers.



Deploying with your own load balancers

If you are using external load balancers for UCP and DTR, you can configure UCP and DTR to use these external load balancers by specifying FQDNs in the loadbalancers dictionary in `group_vars/vars`:

```
loadbalancers:
  ucp:
    public_fqdn: external-ucpvip.am2.cloudra.local
  dtr:
    public_fqdn: external-dtrvip.am2.cloudra.local
```

Installing Docker UCP and DTR on RHEL VMs

The following playbooks are used to install Docker UCP and DTR on RHEL VMs.

- `playbooks/config_storage_driver.yml` prepares drives for local Docker volumes and container images. It also configures Docker with either the overlay2 storage driver (the default) or the devicemapper storage driver, depending on the value of the `docker_storage_driver` variable in `group_vars/vars`. This playbook was previously called `playbooks/config-docker-lvs.yml` in earlier releases of the solution.
- `playbooks/install_docker.yml` installs Docker along with all of its dependencies.
- `playbooks/install_rsyslog.yml` installs and configures **rsyslog** in the logger node and in all Docker nodes. The logger node will be configured to receive all syslogs on port 514 and the Docker nodes will be configured to send all logs (including container logs) to the logger node.
- `playbooks/docker_post_config.yml` performs a variety of tasks to complete the installation of the Docker environment, including configuration of the HTTP/HTTPS proxies, if any, and installation of the VMware vSphere Storage for Docker volume plugin.
- `playbooks/install_nfs_server.yml` installs and configures an NFS server on the NFS node.

This playbook has been updated to configure a third drive which is used to hold the data of the persistent volumes created with the NFS provisioner. The default size for this drive is purposefully kept small because using the NFS VM to store persistent volumes is not recommended for production use. However, this can be useful for demo purposes.

- `playbooks/install_nfs_clients.yml` installs the required packages on the DTR nodes to be able to mount an NFS share.
- `playbooks/create_main_ucp.yml` installs and configures the first Docker UCP instance on the target node defined by the group `ucp_main` in the `vm_hosts` inventory.
- `playbooks/scale_ucp.yml` installs and configures additional instances of UCP on the target nodes defined by the group `ucp` in the `vm_hosts` inventory, except for the node defined in the group `ucp_main`.
- `playbooks/create_main_dtr.yml` installs and configures the first Docker DTR instance on the target node defined by the group `dtr_main` in the `vm_hosts` inventory.
- `playbooks/config_scheduler.yml` configures the scheduler to prevent regular users (i.e. non-admin users) scheduling containers on the Docker nodes running instances of UCP and DTR.
- `playbooks/scale_dtr.yml` installs and configures additional instances (or replicas) of DTR on the target nodes defined by the group `dtr` in the `vm_hosts` inventory, with the exception of the node defined in the group `dtr_main`.
- `playbooks/reconfigure_dtr.yml` is used to reconfigure DTR with the FQDN of the UCP Load Balancer and also enables image scanning.

Deploying RHEL workers

By default, `site.yml` will automatically deploy any RHEL (and / or Windows) worker nodes that are declared in the inventory.

If you subsequently want additional RHEL worker nodes, add them to the inventory as appropriate and then run the playbooks for [Provisioning RHEL VMs](#), followed by the specific playbooks for RHEL worker nodes outlined below:



- `playbooks/scale_workers.yml` installs and configures additional Linux workers on the target nodes defined by the group `worker` in the `vm_hosts` inventory.

A utility script `scale_worker.sh` is provided to assist you in adding worker nodes after the initial deployment.

HPE SimpliVity backup playbooks

Two playbooks are provided to support the backup of Docker volumes using HPE SimpliVity functionality. The playbooks are run by default when using `site.yml` to deploy the solution.

Configure dummy VMs to backup Docker volumes

The playbook `playbooks/config_dummy_vms_for_docker_volumes_backup.yml` ensures that you can back up Docker volumes that have been created using the vSphere plugin (vDVS) in SimpliVity. There is not a straight-forward way to do this, so you need to use a workaround. Since all Docker volumes are going to be stored in the `dockvols` folder in the datastore(s), you need to create a 'dummy' VM per datastore. The `vmx`, `vmfsd` and `vmkd` files from this VM will have to be inside the `dockvols` folder, so when these VMs are backed up, the volumes are backed up as well. Obviously these VMs don't need to take any resources and you can keep them powered off.

Configure SimpliVity backups

The playbook `playbooks/config_simplivity_backups.yml` configures the defined backup policies in the group variables file in HPE SimpliVity and will include all Docker nodes plus the 'dummy' VMs created before, so the existing Docker volumes are also taken into account. The playbook will mainly use the SimpliVity REST API to perform these tasks. A reference to the REST API can be found at: https://api.simplivity.com/rest-api_getting-started_overview/rest-api_getting-started_overview_rest-api-overview.html.

Post deployment

The playbooks are intended to be used to deploy a new environment. You should only use them for Day 0 deployment purposes.

The Ansible log is stored in the folder `/root/Docker-SimpliVity`. If the deployment fails, you may find useful hints in this log. To see how to check if your certs have been deployed correctly, see Appendix D.

Installing kubectl

A convenience playbook is provided to make it easy to install `kubectl`. This playbook uses variables in `group_vars/vars` to determine which version to download. The default version specified by the variable `kubectl_version` in the sample variables file is `1.11.5`.

To run the playbook:

```
# cd ~/Docker-SimpliVity
# ansible-playbook -i vm_hosts playbooks/install_kubectl.yml --vault-password-file .vault_pass
```

Test the installation by running the `kubectl version` command:

```
# kubectl version
```

```
Client Version: version.Info{Major:"1", Minor:"11", GitVersion:"v1.11.5",
GitCommit:"753b2dbc622f5cc417845f0ff8a77f539a4213ea", GitTreeState:"clean", BuildDate:"2018-11-
26T14:41:50Z", GoVersion:"go1.10.3", Compiler:"gc", Platform:"linux/amd64"}
```

The connection to the server localhost:8080 was refused - did you specify the right host or port?

The client version is reported correctly. However, `kubectl` cannot connect to the server until you set up a client bundle - this is described in the section titled Installing the client bundle.

Manually installing kubectl

You can find the version number for the current stable version of `kubectl` at <https://kubernetes.io/docs/tasks/tools/install-kubectl/>. At the time of writing, the stable version is `1.13`.

The following is an example of manually downloading and installing a specific version of `kubectl`.

```
# version=v1.10.4
# wget -O kubectl https://storage.googleapis.com/kubernetes-
```



```
release/release/${version}/bin/linux/amd64/kubectl
# chmod +x ./kubectl
# sudo mv ./kubectl /usr/local/bin/kubectl

# kubectl version
Client Version: version.Info{Major:"1", Minor:"10", GitVersion:"v1.10.4",
GitCommit:"5ca598b4ba5abb89bb773071ce452e33fb66339d", GitTreeState:"clean", BuildDate:"2018-06-
06T08:13:03Z", GoVersion:"go1.9.3", Compiler:"gc", Platform:"linux/amd64"}

Server Version: version.Info{Major:"1", Minor:"8+", GitVersion:"v1.8.11-docker-8d637ae",
GitCommit:"8d637aedef46b9c21dde723e29c645b9f27106fa5", GitTreeState:"clean", BuildDate:"2018-04-
26T16:51:21Z", GoVersion:"go1.8.3", Compiler:"gc", Platform:"linux/amd64"}
```

More details on installing `kubectl` are available at <https://kubernetes.io/docs/tasks/tools/install-kubectl/>.

Installing the client bundle

A convenience playbook is provided to install and apply the client bundle. To run the playbook:

```
# cd ~/Docker-SimpliVity
# ansible-playbook -i vm_hosts playbooks/install_client_bundle.yml --vault-password-file .vault_pass
```

The client bundle is downloaded to `~/certs.<<ucp_instance>>.<<ucp_username>>` where `ucp_instance` will be specific to the cluster you are running against, for example, `hpe2-ucp01` and the `ucp_username` is typically `admin`.

The playbook downloads the client bundle, but does not configure it for use. Change to the download folder and execute `eval "${<env.sh}"`

```
# cd ~/certs.hpe2-ucp01.admin
# eval "${<env.sh}"
```

Test the configuration by again running the `kubectl version` command. It should now report the server version as well as the client version:

```
# kubectl version

Client Version: version.Info{Major:"1", Minor:"11", GitVersion:"v1.11.5",
GitCommit:"753b2dbc622f5cc417845f0ff8a77f539a4213ea", GitTreeState:"clean", BuildDate:"2018-11-
26T14:41:50Z", GoVersion:"go1.10.3", Compiler:"gc", Platform:"linux/amd64"}

Server Version: version.Info{Major:"1", Minor:"11+", GitVersion:"v1.11.5-docker-1",
GitCommit:"d512ba512d0de40cd80258f480ff66bf71f2d8a4", GitTreeState:"clean", BuildDate:"2018-12-
03T19:55:14Z", GoVersion:"go1.10.3", Compiler:"gc", Platform:"linux/amd64"}
```

More information on the client bundle is available at <https://docs.docker.com/ee/ucp/user-access/cli/#download-client-certificates-by-using-the-rest-api>.

Installing Helm

Prerequisites

- Install the `kubectl` binary on your Ansible box
- Install the UCP Client bundle for the admin user
- Confirm that you can connect to the cluster by running a test command, for example, `kubectl get nodes`

Running the playbook

To run the playbook:

```
# cd ~/Docker-SimpliVity
# ansible-playbook -i vm_hosts playbooks/install_helm.yml --vault-password-file .vault_pass
```



Install sample charts

A number of sample charts are delivered with the solution, for the purposes of demonstration.

Alpine

A simple chart is provided in the `~/Docker-SimpliVity/test/files/helm/alpine` directory to run a single pod of Alpine Linux.

The `templates/` directory contains a very simple pod resource with a couple of parameters. The `values.yaml` file contains the default values for the `alpine-pod.yaml` template.

```
# cd ~/Docker-SimpliVity
# helm install test/files/helm/alpine
```

The output shows that a single pod was deployed.

```
NAME:      old-mole
LAST DEPLOYED: Fri Feb  8 17:27:35 2019
NAMESPACE: default
STATUS:    DEPLOYED
```

RESOURCES:

```
==> v1/Pod
```

NAME	READY	STATUS	RESTARTS	AGE
old-mole-alpine	1/1	Running	0	0s

Nginx

An example chart is provided in the `~/Docker-SimpliVity/test/files/helm/nginx` directory to install a simple nginx server according to the following pattern:

- A ConfigMap is used to store the files the server will serve. (`templates/configmap.yaml`)
- A Deployment is used to create a Replica Set of nginx pods. (`templates/deployment.yaml`)
- A Service is used to create a gateway to the pods running in the replica set (`templates/service.yaml`)

The `values.yaml` exposes a few of the configuration options in the charts.

```
NAME:      worn-olm
LAST DEPLOYED: Fri Feb  8 16:23:21 2019
NAMESPACE: default
STATUS:    DEPLOYED
```

RESOURCES:

```
==> v1/Deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
worn-olm-nginx	1	1	1	1	14s

```
==> v1/Pod[related]
```

NAME	READY	STATUS	RESTARTS	AGE
worn-olm-nginx-7d648f7dfb-gg2jk	1/1	Running	0	14s
worn-olm-nginx-vhwc7	0/1	Completed	0	14s

```
==> v1/ConfigMap
```

NAME	DATA	AGE
worn-olm-nginx	2	14s

```
==> v1/Service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
worn-olm-nginx	NodePort	10.96.30.222	<none>	80:34567/TCP	14s



Helm also allows you to easily delete installed releases. List the installed releases to find the name of the release you wish to delete.

```
# helm list
NAME                REVISION    UPDATED                               STATUS    CHART          APP
VERSION            NAMESPACE
worn-olm            1           Fri Feb  8 16:23:21 2019        DEPLOYED  nginx-0.1.0
default
```

Use the `helm delete` command to remove the named release.

```
# helm delete worn-olm
release "worn-olm" deleted
```

Post-deploy validation

Many sample Kubernetes applications are available at <https://kubernetes.io/docs/tutorials/>. This section details how to deploy the stateless `guestbook` application with Redis as documented at <https://kubernetes.io/docs/tutorials/stateless-application/guestbook/>.

When deploying applications, you must be aware that Kubernetes version 1.11 shipped with Docker 2.1. If you are testing examples that are designed to work with a newer (or older) version of Kubernetes, you may have to make changes in some places to the configuration files.

Prerequisites

- Install the `kubect1` binary on your Ansible box
- Install the UCP Client bundle for the admin user
- Confirm that you can connect to the cluster by running a test command, for example, `kubect1 get nodes`

Kubernetes guestbook example with Redis

The playbook for the Kubernetes example `guestbook` is based on the example taken from the GitHub repo at <https://github.com/kubernetes/examples>.

```
# cd ~/Docker-Simplivity
# ansible-playbook -i vm-hosts test/playbooks/k8s-guestbook.yml --vault-password-file .vault_pass
```

You can run the playbook directly, but it can be informative to walk through the individual files to see what is going on under the covers.

Quickstart

```
# cd ~/Docker-Simplivity/test/files/k8s-examples/guestbook
# kubect1 apply -f redis-master-deployment.yaml
# kubect1 apply -f redis-master-service.yaml
# kubect1 apply -f redis-slave-deployment.yaml
# kubect1 apply -f redis-slave-service.yaml
# kubect1 apply -f frontend-deployment.yaml
# kubect1 apply -f frontend-service.yaml
# kubect1 get svc frontend
```

Details

Change to the directory containing the `guestbook` YAML files.

```
# cd ~/Docker-Simplivity/test/files/k8s-examples/guestbook
```

The manifest file `redis-master-deployment.yaml`, included below, specifies a deployment controller that runs a single replica Redis master pod.

```
# cat redis-master-deployment.yaml
```

```
apiVersion: apps/v1 # for k8s versions before 1.9.0 use apps/v1beta2 and before 1.8.0 use extensions/v1beta1
kind: Deployment
metadata:
```



```

name: redis-master
spec:
  selector:
    matchLabels:
      app: redis
      role: master
      tier: backend
  replicas: 1
  template:
    metadata:
      labels:
        app: redis
        role: master
        tier: backend
    spec:
      containers:
        - name: master
          image: k8s.gcr.io/redis:e2e # or just image: redis
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          ports:
            - containerPort: 6379

```

Apply the Redis master deployment from the redis-master-deployment.yaml file:


```
# kubectl apply -f redis-master-deployment.yaml
```

Query the list of Pods to verify that the Redis master pod is running.

```
# kubectl get pods | grep redis
redis-master-57657796fc-psvhc      1/1      Running    0      32s
```

Use the `kubectl logs` command to view the logs from the Redis master pod:

```
# kubectl logs -f redis-master-57657796fc-psvhc
```



```
Redis 2.8.19 (00000000/0) 64 bit
```

```
Running in stand alone mode
```

```
Port: 6379
```

```
PID: 1
```

```
http://redis.io
```



```
[1] 07 Feb 15:04:32.189 # Server started, Redis version 2.8.19
[1] 07 Feb 15:04:32.189 # WARNING you have Transparent Huge Pages (THP) support enabled in your
kernel. This will create latency and memory usage issues with Redis. To fix this issue run the command
'echo never > /sys/kernel/mm/transparent_hugepage/enabled' as root, and add it to your /etc/rc.local
in order to retain the setting after a reboot. Redis must be restarted after THP is disabled.
[1] 07 Feb 15:04:32.189 # WARNING: The TCP backlog setting of 511 cannot be enforced because
/proc/sys/net/core/somaxconn is set to the lower value of 128.
[1] 07 Feb 15:04:32.190 * The server is now ready to accept connections on port 6379
```

The guestbook application needs to communicate with the Redis master to write its data. You need to apply a service to proxy the traffic to the Redis master pod. A service defines a policy to access the pods.

```
# cat redis-master-service.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: redis-master
  labels:
    app: redis
    role: master
    tier: backend
spec:
  ports:
    - port: 6379
      targetPort: 6379
  selector:
    app: redis
    role: master
    tier: backend
```

Apply the Redis master service from the `redis-master-service.yaml` file:

```
# kubectl apply -f redis-master-service.yaml
service "redis-master" created
```

Query the list of services to verify that the Redis master service is running.

```
# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
redis-master	ClusterIP	10.96.240.18	<none>	6379/TCP	1m

Although the Redis master is a single pod, you can make it highly available to meet traffic demands by adding replica Redis slaves.

```
# cat redis-slave-deployment.yaml
```

```
apiVersion: apps/v1 # for k8s versions before 1.9.0 use apps/v1beta2 and before 1.8.0 use
extensions/v1beta1
kind: Deployment
metadata:
  name: redis-slave
spec:
  selector:
    matchLabels:
      app: redis
      role: slave
      tier: backend
```



```

replicas: 2
template:
  metadata:
    labels:
      app: redis
      role: slave
      tier: backend
  spec:
    containers:
      - name: slave
        image: gcr.io/google-samples/gb-redisslave:v1
        resources:
          requests:
            cpu: 100m
            memory: 100Mi
        env:
          - name: GET_HOSTS_FROM
            value: dns
            # If your cluster config does not include a dns service, then to
            # instead access an environment variable to find the master
            # service's host, comment out the 'value: dns' line above, and
            # uncomment the line below:
            # value: env
        ports:
          - containerPort: 6379

```

Create the Redis slaves from the `redis-slave-deployment.yaml` file.

```

# kubectl apply -f redis-slave-deployment.yaml
deployment.apps "redis-slave" created

```

Query the list of Pods to verify that the Redis slave pods are running.

```

# kubectl get pods | grep redis
redis-master-57657796fc-psvhc      1/1      Running    0           7m
redis-slave-5cb5956459-bqq1g      1/1      Running    0           19s
redis-slave-5cb5956459-gql5x      1/1      Running    0           19s

```

The guestbook application needs to communicate to Redis slaves to read data. To make the Redis slaves discoverable, you need to set up a service that provides transparent load balancing to the set of pods.

```

# cat redis-slave-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: redis-slave
  labels:
    app: redis
    role: slave
    tier: backend
spec:
  ports:
    - port: 6379
  selector:

```



```

app: redis
role: slave
tier: backend

```

Deploy the Redis slave service from the `redis-slave-service.yaml` file.

```

# kubectl apply -f redis-slave-service.yaml
service "redis-slave" created

```

Query the list of services to verify that the Redis slave service is running.

```

# kubectl get services | grep redis
redis-master   ClusterIP   10.96.240.18   <none>         6379/TCP      4m
redis-slave    ClusterIP   10.96.200.85   <none>         6379/TCP      22s

```

The guestbook application has a web frontend written in PHP serving the HTTP requests. It is configured to connect to the `redis-master` service for write requests and the `redis-slave` service for read requests.

```

# cat frontend-deployment.yaml
apiVersion: apps/v1 # for k8s versions before 1.9.0 use apps/v1beta2 and before 1.8.0 use extensions/v1beta1
kind: Deployment
metadata:
  name: frontend
spec:
  selector:
    matchLabels:
      app: guestbook
      tier: frontend
  replicas: 3
  template:
    metadata:
      labels:
        app: guestbook
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google-samples/gb-frontend:v4
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          env:
            - name: GET_HOSTS_FROM
              value: dns
            # If your cluster config does not include a dns service, then to
            # instead access environment variables to find service host
            # info, comment out the 'value: dns' line above, and uncomment the
            # line below:
            # value: env
          ports:
            - containerPort: 80

```

Create the frontend deployment using the `frontend-deployment.yaml` file.



```
# kubectl apply -f frontend-deployment.yaml
deployment.apps "frontend" created
```

Query the list of pods to verify that the three frontend replicas are running.

```
# kubectl get pods -l app=guestbook -l tier=frontend
NAME                                READY    STATUS    RESTARTS   AGE
frontend-7f5cd767dc-28j6b          1/1      Running   0           23s
frontend-7f5cd767dc-mqcbv          1/1      Running   0           23s
frontend-7f5cd767dc-v6lwc          1/1      Running   0           23s
```

If you want guests to be able to access your guestbook, you must configure the frontend service to be externally visible, so a client can request the service from outside the container cluster.

```
# cat frontend-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # comment or delete the following line if you want to use a LoadBalancer
  type: NodePort
  # if your cluster supports it, uncomment the following to automatically create
  # an external load-balanced IP for the frontend service.
  # type: LoadBalancer
  ports:
    - port: 80
  selector:
    app: guestbook
    tier: frontend
```

Deploy the frontend service using the `frontend-service.yaml` file.

```
# kubectl apply -f frontend-service.yaml
service "frontend" created
```

Query the list of services to verify that the frontend service is running.

```
# kubectl get services | grep frontend
frontend    NodePort    10.96.16.200    <none>    80:33444/TCP    25s
```

Access the UI using the identified port on any node in your cluster, for example, <http://hpe2-ucp01.am2.cloudra.local:33444/> as shown in Figure 3.






Figure 3. Guestbook UI

Teardown

A playbook is provided to remove the deployed `guestbook` artifacts.

```
# cd ~/Docker-Simplivity
# ansible-playbook -i vm_hosts test/playbooks/k8s-guestbook-teardown.yml --vault-password-file
.vault_pass
```

UCP metrics in Prometheus

Docker EE 2.1 uses a built-in deployment of Prometheus to power the performance graphs in the web UI for UCP. The metrics that UCP generates can be routed to a separate Prometheus, if required. A convenience playbook has been provided to configure a minimal Prometheus and Grafana deployment that can help visualize all of the metrics that UCP generates.

For more information on UCP cluster metrics, see the article at <https://docs.docker.com/ee/ucp/admin/configure/collect-cluster-metrics/>.

Prerequisites

- Install the `kubectl` binary on your Ansible box
- Install the UCP Client bundle for the admin user
- Confirm that you can connect to the cluster by running a test command, for example, `kubectl get nodes`

Deploy Prometheus and Grafana

The playbook `playbooks/ucp-metrics-prometheus.yml` deploys pods for Prometheus and Grafana and configures them to use the client bundle to access the UCP metrics. To run the playbook:

```
# cd ~/Docker-Simplivity
# ansible-playbook -i vm_hosts playbooks/ucp-metrics-prometheus.yml --vault-password-file .vault_pass
```

Prometheus UI

The playbook exposes a port to access the user interface for Prometheus - to find the port, get the details of the `prometheus` service:

```
# kubectl get svc Prometheus
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
prometheus	NodePort	10.96.216.220	<none>	9090:34713/TCP	6d

The Prometheus UI can be accessed on any node in your cluster, using the port returned by `kubectl get svc`. In this instance, it is accessed at `http://hpe2-ucp01.am2.cloudra.local:34713` as shown in Figure 4.






Figure 4. UCP metrics in Prometheus

Using Grafana to vizualize UCP metrics

The playbook also exposes a port to access the Grafana UI - to find the port, get the details of the grafana service:

```
# kubectl get svc Grafana
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
grafana   NodePort  10.96.177.108   <none>           3000:33118/TCP   6d
```

The Grafana UI can be accessed on any node in your cluster, using the port returned by `kubectl get svc`. In this instance, it is accessed at `http://hpe2-ucp01.am2.cloudra.local:33118`. The example UCP Dashboard shown in Figure 5 is taken from <https://grafana.com/dashboards/9309>.





Figure 5. UCP Dashboard in Grafana

Configuring storage

Deploying the NFS provisioner for Kubernetes

NFS can be provisioned using the NFS VM for proof of concept or demo systems.

Prerequisites

- Configure the variables described in the section Kubernetes Persistent Volume configuration
- Install the kubectl binary on your Ansible box
- Install the UCP Client bundle for the admin user
- Confirm that you can connect to the cluster by running a test command, for example, `kubectl get nodes`

Using NFS VM for post-deployment verification

In this example, it is assumed that the relevant variables are configured as shown in Table 15.

Table 15. NFS provisioner configuration values

Variable	Value
nfs_provisioner_namespace	nfsstorage
nfs_provisioner_role	nfs-provisioner-runner
nfs_provisioner_serviceaccount	nfs-provisioner
nfs_provisioner_name	hpe.com/nfs
nfs_provisioner_storage_class_name	nfs
nfs_provisioner_server_ip	hpe-nfs.cloudra.local
nfs_provisioner_server_share	/k8s

In this instance, the server IP is set to the NFS VM that has been deployed.



Running the playbook

Once the prerequisites are satisfied, run the appropriate playbook on your Ansible node.

```
# cd Docker-SimpliVity
# ansible-playbook -i vm-hosts playbooks/nfs-provisioner.yml --vault-password-file .vault_pass
```

For validation, the playbook creates a test claim and a pod, the pod writes content to a file, the pod is deleted and then the playbook checks that the contents of the file have been persisted.

```
kubectl -n {{ nfs_provisioner_namespace }} apply -f /tmp/nfs-provisioner-test-claim.yml
kubectl -n {{ nfs_provisioner_namespace }} apply -f /tmp/nfs-provisioner-test-pod.yml
```

```
sleep 5 # need sleep here to allow pod/container to start up and write file
```

```
ssh {{ nfs_provisioner_server_ip }} ls -R {{ nfs_provisioner_server_share }}
echo '*** delete test-pod ***'
kubectl -n {{ nfs_provisioner_namespace }} delete -f /tmp/nfs-provisioner-test-pod.yml
echo '*** cat bar.txt ***'
ssh {{ nfs_provisioner_server_ip }} "cd {{ nfs_provisioner_server_share }}/{{ nfs_provisioner_namespace }}*; cat bar.txt"
echo '*** delete test-claim ***'
kubectl -n {{ nfs_provisioner_namespace }} delete -f /tmp/nfs-provisioner-test-claim.yml
```

The output of the playbook shows the various steps taking place:

```
"pod/test-pod created",
"/k8s:",
"nfsstorage-test-claim-pvc-e6a09191-3b41-11e9-a830-0242ac11000b",
"",
"/k8s/nfsstorage-test-claim-pvc-e6a09191-3b41-11e9-a830-0242ac11000b:",
"bar.txt",
"*** delete test-pod ***",
"pod \"test-pod\" deleted",
"*** cat bar.txt ***",
"hello",
"*** delete test-claim ***",
"persistentvolumeclaim \"test-claim\" deleted"
```

Running the command `kubectl get sc` will show the storage class named `nfs`:

```
# kubectl get sc
NAME          PROVISIONER  AGE
nfs           hpe.com/nfs  5m
```

The following section shows how to manually perform a similar validation test to the one done by the playbook.

Manually testing the NFS provisioner

Create a temporary file `/tmp/pvc.yml` for a persistent volume claim (PVC) named `dynnfs-testpvc` with a storage class of `nfs`

```
# cat /tmp/pvc.yml <<EOF
---
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: dynnfs-testpvc
```



```

  annotations:
    volume.beta.kubernetes.io/storage-class: "nfs"
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 100Mi
EOF

```

Create the PVC resource by running `kubectl apply` on this file.

```

# kubectl apply -f /tmp/pvc.yml
persistentvolumeclaim "dynnfs-testpvc" created

```

Verify that the corresponding persistent volume (PV) was created at the same time.

```

# kubectl get pv
NAME                                CAPACITY  STATUS  CLAIM                                STORAGECLASS  AGE
pvc-e685a9d2-8a6f-11e8-...         100Mi     Bound   default/dynnfs-testpvc             nfs            4s

```

Define a pod that will mount the persistent volume by using the persistent volume claim. The persistent volume is mounted under `/tmp/foo`.

```

# cat /tmp/pod.yml <<EOF
apiVersion: apps/v1beta2
kind: Deployment
metadata:
  name: dynnfs-testpod
spec:
  selector:
    matchLabels:
      app: dynnfs-testpod
  replicas: 1
  template:
    metadata:
      labels:
        app: dynnfs-testpod
    spec:
      volumes:
        - name: pod-data
          persistentVolumeClaim:
            claimName: dynnfs-testpvc
      containers:
        - name: dynnfs-testpod
          command:
            - sh
            - -c
            - while true; do sleep 1; done
          image: radial/busyboxplus:curl
          volumeMounts:
            - mountPath: /tmp/foo
              name: pod-data
EOF

```

Create the pod resource by running `kubectl apply` on the file.



```
# kubectl apply -f /tmp/pod.yml
deployment.apps "dynnfs-testpod" created
```

Retrieve the pod ID and then execute a command in the pod to create a test file on the persistent volume. The file is named `/tmp/foo/bar.txt` and contains the string `hello`.

```
# pod=$(kubectl get pod | awk '/dynnfs-testpod-/ {print $1}')
# kubectl exec -it $pod -- sh -c "echo hello >/tmp/foo/bar.txt"
```

In this example, where the NFS VM is being used as the storage back-end, you can examine the content of the folder containing the persistent volumes. Given the values specified above, where the NFS VM is named `hpe-nfs` and the `nfs_provisioner_server_share` is `k8s`, you can connect to the VM and explore the folder as follows.

```
# ssh hpe-nfs ls -R /k8s
/k8s:
default-dynnfs-testpvc-pvc-e685a9d2-8a6f-11e8-9025-0242ac110010

/k8s/default-dynnfs-testpvc-pvc-e685a9d2-8a6f-11e8-9025-0242ac110010:
bar.txt
```

Examine the contents of the file to ensure that the string `hello` has been persisted in the file `bar.txt`.

```
# ssh hpe-nfs cat /k8s/default-dynnfs-testpvc-pvc-e685a9d2-8a6f-11e8-9025-0242ac110010/bar.txt
hello
```

Validating the NFS provisioner using WordPress and MySQL

A sample playbook has been provided to show how to use the NFS provioner for perstent storage for a WordPress and MySQL deployment.

Prerequisites

- Install the kubectl binary on your Ansible box
- Install the UCP Client bundle for the admin user
- Confirm that you can connect to the cluster by running a test command, for example, `kubectl get nodes`

Deploy the NFS provisioner as outlined in the preceeding section. The article assumes that the NFS configuration is the same as used in that section, as shown in Table 16:

Table 16. NFS provisioner configuration values

Variable	Value
nfs_provisioner_namespace	nfsstorage
nfs_provisioner_role	nfs-provisioner-runner
nfs_provisioner_serviceaccount	nfs-provisioner
nfs_provisioner_name	hpe.com/nfs
nfs_provisioner_storage_class_name	nfs
nfs_provisioner_server_ip	hpe2-nfs.cloudra.local
nfs_provisioner_server_share	/k8s

Running the playbook

The playbook `test/playbooks/wordpress-mysql-nfs.yml` creates Persistent Volume Claims for both Wordpress and MySQL, deploys both applications and makes the WordPress UI available via a NodePort.



```
# cd ~/Docker-SimpliVity
# ansible-playbook -i vm_hosts ./test/playbooks/wordpress-mysql-nfs.yml --vault-password-file
.vault_pass
```

The output shows the components created along with the NodePort for the `wordpress` service.

```
ok: [localhost] => {
  "ps.stdout_lines": [
    "Cluster \"ucp_hpe2-ucp01.am2.cloudra.local:6443_admin\" set.",
    "User \"ucp_hpe2-ucp01.am2.cloudra.local:6443_admin\" set.",
    "Context \"ucp_hpe2-ucp01.am2.cloudra.local:6443_admin\" modified.",
    "namespace/wordpress-mysql created",
    "secret/mysql-pass created",
    "persistentvolumeclaim/mysql-pv-claim created",
    "persistentvolumeclaim/wp-pv-claim created",
    "deployment.apps/wordpress-mysql created",
    "deployment.apps/wordpress created",
    "service/wordpress-mysql created",
    "service/wordpress created",
    "NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE",
    "wordpress           NodePort      10.96.216.103   <none>           80:33790/TCP     0s",
    "wordpress-mysql     ClusterIP     None            <none>           3306/TCP         0s"
  ]
}
```

Browse to the specified port on any node in your cluster.

```
http://hpe2-ucp01.am2.cloudra.local:33790
```

Configuring WordPress

You need to configure the language and password before WordPress is ready to use, as shown in Figure 6.

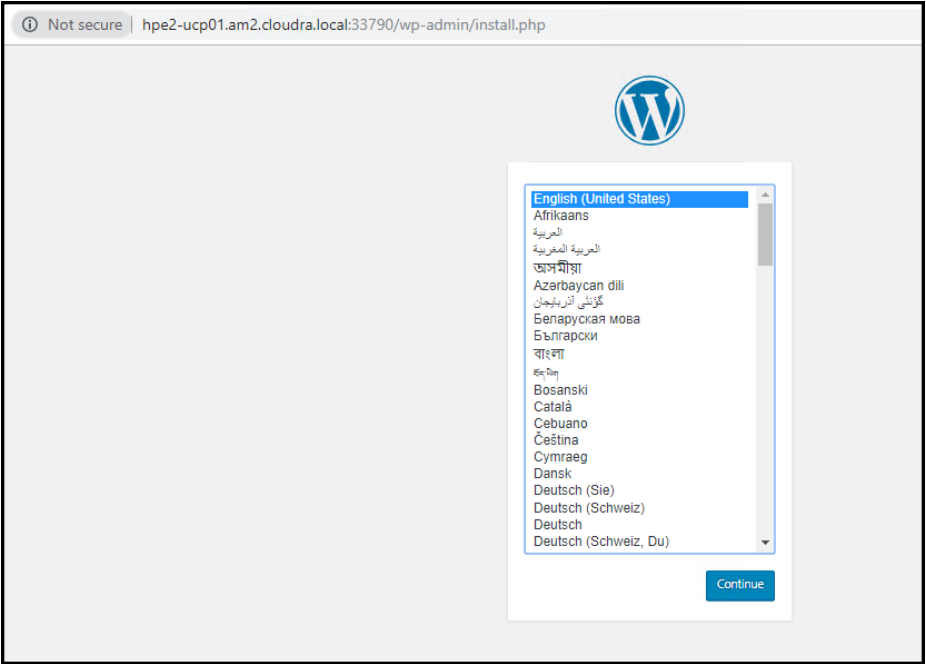


Figure 6. Configure WordPress language



Add a username, password and other configuration details, as shown in Figure 7.

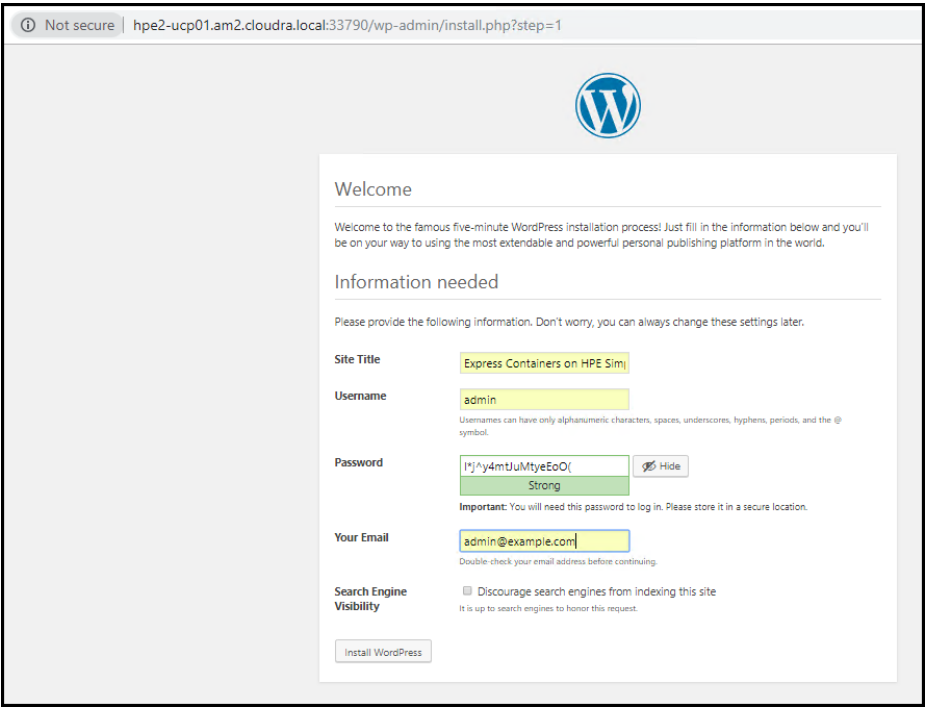


Figure 7. Configure WordPress password

Log in to WordPress, as shown in Figure 8, with the user name and password you have just set up.




Figure 8. WordPress login

The welcome page is displayed, as shown in Figure 9.



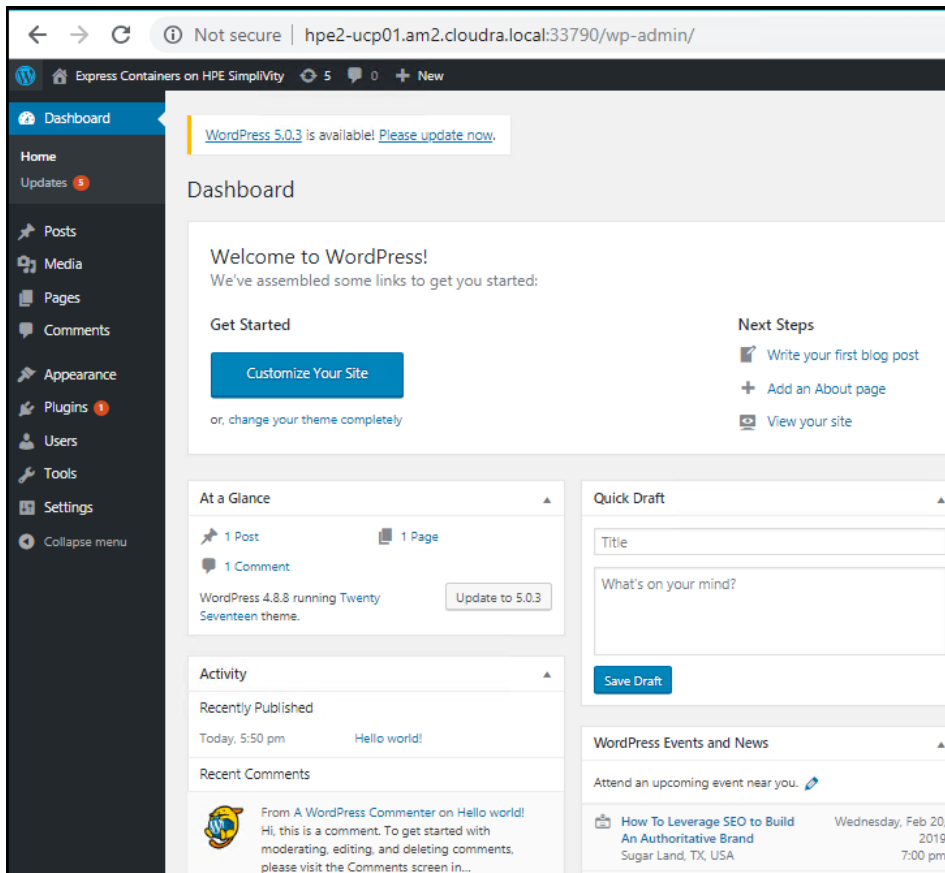


Figure 9. WordPress welcome

Create your first post

Click on **Write your first blog post** and start creating some content. Add a blog title and then click **Add Media** to upload an image to the Media Library and then **Insert** into post. In this example, as shown in Figure 10, the image is a file named **380 with OmniStack.jpg**.



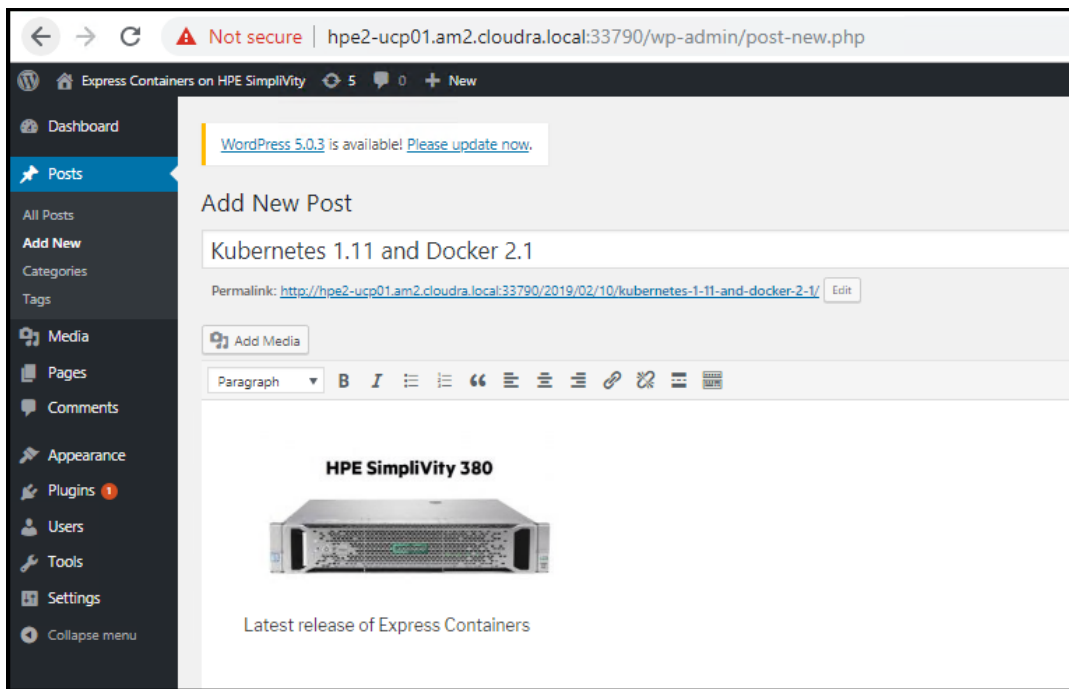


Figure 10. Create your first WordPress blog post

Click **Publish** and then **View post** to see your new blog post, as shown in Figure 11.




Figure 11. View your first post

Test persistence for WordPress

Find your WordPress Persistent Volume Claim (PVC).




```
# kubectl -n wordpress-mysql get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS	STORAGECLASS
mysql-pv-claim	Bound	pvc-d48880e3-2d58-11e9-adb2-0242ac110003	1Gi	RWO	nfs
wp-pv-claim	Bound	pvc-d4bc101f-2d58-11e9-adb2-0242ac110003	20Gi	RWO	nfs

Connect to the NFS VM and browse the /k8s folder to find the volume for the WordPress claim `wp-pv-claim`.

```
# ssh hpe2-nfs ls /k8s
wordpress-mysql-mysql-pv-claim-pvc-d48880e3-2d58-11e9-adb2-0242ac110003
wordpress-mysql-wp-pv-claim-pvc-d4bc101f-2d58-11e9-adb2-0242ac110003
```

Locate the `wp-content` folder.

```
# ssh hpe2-nfs ls /k8s/wordpress-mysql-wp-pv-claim-pvc-d4bc101f-2d58-11e9-adb2-0242ac110003
index.php
license.txt
readme.html
wp-activate.php
wp-admin
wp-blog-header.php
wp-comments-post.php
wp-config.php
wp-config-sample.php
wp-content
wp-cron.php
wp-includes
wp-links-opml.php
wp-load.php
wp-login.php
wp-mail.php
wp-settings.php
wp-signup.php
wp-trackback.php
xmlrpc.php
```

Now find the image used in the blog post.

```
# ssh hpe2-nfs ls /k8s/wordpress-mysql-wp-pv-claim-pvc-d4bc101f-2d58-11e9-adb2-0242ac110003/wp-
content/uploads/2019/02
380-with-OmniStack-100x100.jpg
380-with-OmniStack-150x150.jpg
380-with-OmniStack-300x150.jpg
380-with-OmniStack-768x384.jpg
380-with-OmniStack.jpg
```

Note that WordPress has created a number of variations of the original image, for different screen sizes. Shutdown wordpress (leave MySQL running for now).

```
# kubectl -n wordpress-mysql delete -f /tmp/wordpress-mysql-nfs/wordpress-deployment.yml
deployment.apps "wordpress" deleted
```

Refresh the page in the browser to confirm that WordPress is indeed inaccessible.






Figure 12. Cannot connect to WordPress

Now redeploy Wordpress

```
# kubectl -n wordpress-mysql apply -f /tmp/wordpress-mysql-nfs/wordpress-deployment.yml
deployment.apps/wordpress created
```

Refresh the page in the browser to confirm that WordPress is now accessible and that the image in the blog post has survived the shutdown, as shown in Figure 13.

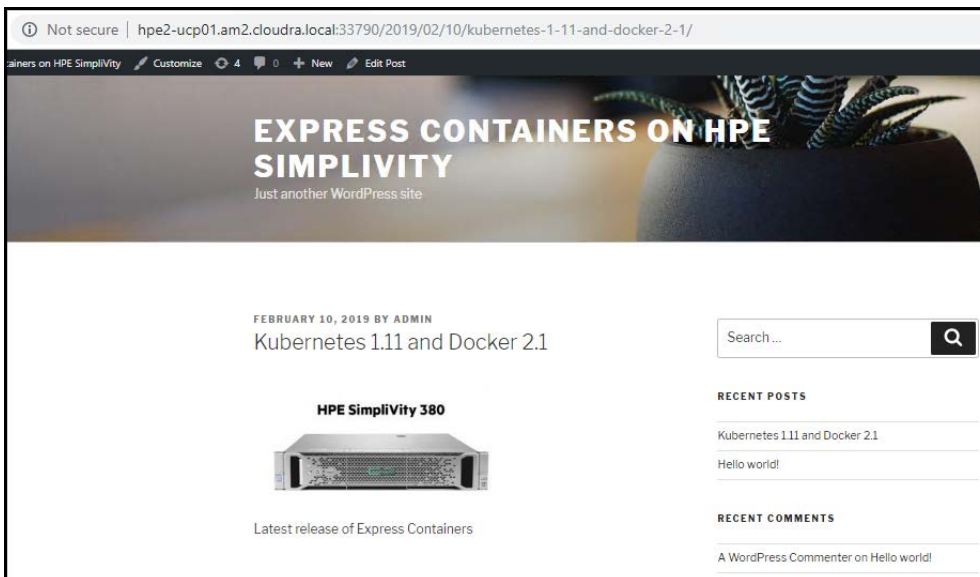


Figure 13. View restored post

Test persistence in MySQL

A similar procedure can be performed for MySQL. While assets such as images, CSS files, etc are stored in the WordPress volume, information about users, posts, comments, tags, etc are stored in the MySQL database. It is possible to browse the tables in the database and identify the rows related to the blog post you created.

Shut down MySQL as follows:

```
# kubectl -n wordpress-mysql delete -f /tmp/wordpress-mysql-nfs/mysql-deployment.yml
deployment.apps "wordpress-mysql" deleted
```



Refresh the page for your blog post, and you will see that WordPress can no longer connect to the database, as shown in Figure 14.



Figure 14. Cannot connect to MySQL

Restore the MySQL deployment:

```
# kubectl -n wordpress-mysql apply -f /tmp/wordpress-mysql-nfs/mysql-deployment.yml
deployment.apps/wordpress-mysql created
```

Refresh the page in the browser, as shown in Figure 15, to confirm that WordPress can now access the database and that the blog post has survived the database shutdown.




Figure 15. Check after MySQL restored

Deploying Windows workers

The `site.yml` playbook will automatically deploy any Windows workers declared in the inventory. The playbooks should run for approximately 70-80 minutes with 3 Windows workers added to the default deployment (depending on your server specifications and the size of your environment). The increase in running time is primarily due to the need to update Windows after creating the VMs.

This section describes the functionality and configuration of the Windows-specific playbooks. It also details how to create the initial Windows template and how to manage deploying Windows worker nodes behind a proxy.

Create the Windows Template

To create the Windows VM Template that you will use as the base for all your Windows worker nodes, you will first create a Virtual Machine with the OS installed and then convert the Virtual Machine to a VM Template. The VM Template is created as lean as possible, with any additional software installs and/or system configuration performed subsequently using Ansible.

As the creation of the template is a one-off task, this procedure has not been automated. The steps to create a VM template manually are outlined below.



Log in to vCenter and create a new Virtual Machine with the following characteristics:

- Guest OS Family: Windows, Guest OS Version: Microsoft Windows Server 2016 (64-bit)
- Hard Disk size: 100GB (Thin provisioning), 1 vCPU and 4 GB of RAM. Both vCPU and memory can be altered later after you deploy from this template.
- A single network controller connected to the network or VLAN of your choice. All VMs will connect to this same network.
- Change the network type to **VMXNET3**, and attach the Windows Server 2016 ISO image from a datastore ensuring you connect the CD/DVD drive on boot.
- Click on the **VM Options** tab, and in the **Boot Options** section, select **Force BIOS setup[*]** to ensure that the machine enters the BIOS setup screen on next boot of this VM. This will allow you to adjust the boot order, placing the virtual CDROM in front of your hard drive.
- Optionally you can remove the floppy drive.

Install Windows Server 2016:

1. Power on the selected VM and then Open Console. Once connected to the console, you will be placed in the BIOS setup screen.
2. Select the Boot tab, click on CD-ROM Drive and move up the CDROM drive above the hard drive. This allows your Windows Server 2016 ISO image to be loaded first on boot. F10 Save and exit is next step.
3. Enter your choices for Language, Time/Currency Format, Keyboard and then Install Now.
4. Select the OS you want to install, and then select Custom: Install Windows Only.
5. Select drive 0, the 100 GB drive you specified earlier, as the location for installing windows.
6. Add a password for the Administrator user.
7. Install VMware Tools and reboot.
8. Once the VM has re-booted, add a temporary network IP address.
9. Use the **sconfig** utility from (MS-DOS) command line to install Windows updates and enable remote desktop.
10. Perform any other customizations you require at this point.
11. Prior to converting the VM to Template, run **Sysprep: C:\Windows\System32\Sysprep\Sysprep.exe**
12. Ensure 'System Out-of-Box Experience (OOBE)' is selected.
13. Select the 'Generalize' option.
14. Select 'Shutdown' from the Shutdown Options.
15. Shutdown VM, and untick **Connect CD/DVD** so that the Windows Server 2016 ISO is no longer mounted. Boot the Windows VM one final time and enter regional settings applicable to your location and keyboard mapping, then enter a password and Shutdown VM.

Note

The `vmware_guest` module used by the playbooks will generate a new SID.

Turn the VM into a template by right-clicking on your VM and selecting **Template -> Convert to Template**. This will create a new template visible under VM Templates in Folders, ready for future use.

Playbooks for adding Windows workers

- `playbooks/create_vms.yml` will create all the necessary Windows 2016 VMs for the environment based on the Windows VM Template defined in the `win_vm_template` variable. Windows workers nodes are defined in the group `win_worker` in the `vm_hosts` inventory.
- `playbooks/install_docker.yml` installs Docker along with all its dependencies on your Windows VMs



- `playbooks/scale_workers.yml` installs and configures additional Windows workers on the target nodes defined by the group `win_worker` in the `vm_hosts` inventory.
- `playbooks/splunk_uf_win.yml` installs and configures the Splunk Universal Forwarder on each Windows machine in the inventory.

Windows configuration

Window-related variables are shown in Table 17.

Table 17. Windows variables

Variable	File	Description
<code>win_vm_template</code>	<code>group_vars/vars</code>	Name of the Windows 2016 VM Template to use. Note that this is the name from a vCenter perspective, not the hostname.
<code>win_username</code>	<code>group_vars/vars</code>	Windows user name. The default is <code>Administrator</code> .
<code>win_password</code>	<code>group_vars/vault</code>	The password for the Windows account.
<code>windows_vdvs_ps</code>	<code>group_vars/vars</code>	Variable used to download the PowerShell script that is used to install vDVS for Windows. For example, <code>https://raw.githubusercontent.com/vmware/vsphere-storage-for-docker/master/install-vdvs.ps1</code>
<code>windows_vdvs_path</code>	<code>group_vars/vars</code>	Variable used to download vSphere Docker Volume Service software. This variable is combined with <code>windows_vdvs_version</code> (below) to generate a URL of the form <code><windows_vdvs_path>_<windows_vdvs_version>.zip</code> to download the software. For example, to download version 0.21, set <code>windows_vdvs_path</code> equal to <code>https://vmware.bintaray.com/vDVS/vsphere-storage-for-docker_windows</code> and <code>windows_vdvs_version</code> equal to <code>0.21</code>
<code>windows_vdvs_version</code>	<code>group_vars/vars</code>	Combined with <code>windows_vdvs_path</code> , this variable is used to generate the URL for downloading the software.
<code>windows_vdvs_directory</code>	<code>group_vars/vars</code>	Variable used to determine where vDVS software will be unzipped and installed from. The default is <code>C:\Users\Administrator\Downloads</code>
<code>windows_winrm_script</code>	<code>group_vars/vars</code>	Variable used to determine where the <code>winrm</code> Powershell script will be downloaded from. See the section Deploying Windows workers behind a proxy for more information.
<code>docker_ee_version_windows</code>	<code>group_vars/vars</code>	It is important that the version of the Docker engine running on your Windows worker nodes is the same as that running on RHEL in the rest of your cluster. You should use this variable to explicitly match up the versions. For Docker 2.0, the recommended value is <code>'17.06'</code> . If you do not explicitly set this value, you may end up with an incompatible newer version running on your Windows workers.
<code>windows_update</code>	<code>group_vars/vars</code>	Variable used to determine if Windows updates are automatically downloaded when installing Docker on Windows worker nodes (in the <code>playbooks/install_docker.yml</code>). Defaults to <code>true</code> . See the section Deploying Windows workers behind a proxy for more information.

group_vars/win_worker.yml

There is a separate file in the `group_vars` directory named `win_worker.yml` for advanced, Windows-specific configuration. These variables are used in the following playbooks:

- `playbooks/create_windows_vms.yml`
- `playbooks/install_docker_window.yml`
- `playbooks/scale_workers_windows.yml`



In general, it should not be necessary to modify this file, but the variables are documented in Table 18 for the sake of completeness.

Table 18. Advanced windows variables

Variable	File	Description
ansible_user	group_vars/win_worker.yml	Defaults to the Windows user account <code>win_username</code> as specified in <code>group_vars/vars</code>
ansible_password	group_vars/win_worker.yml	Defaults to the Windows user password <code>win_password</code> as specified in <code>group_vars/vault</code>
ansible_port	group_vars/win_worker.yml	5986
ansible_connection	group_vars/win_worker.yml	winrm
ansible_winrm_server_cert_validation	group_vars/win_worker.yml	Defaults to <code>ignore</code>
ansible_winrm_operation_timeout_sec	group_vars/win_worker.yml	Defaults to 250
ansible_winrm_read_timeout_sec	group_vars/win_worker.yml	Defaults to 300
windows_timezone	group_vars/win_worker.yml	Defaults to 15

Deploying Windows workers behind a proxy

If you are deploying behind a proxy, you will need to make a local copy of the remoting script available.

Configuring the winrm remoting script

The playbooks for deploying Windows workers rely on a Powershell script for remote access from the Ansible machine. The script `ConfigureRemotingForAnsible.ps1` is available online on GitHub and you can configure the playbooks to use the script directly as follows:

```
windows_winrm_script: '
https://github.com/ansible/ansible/blob/devel/examples/scripts/ConfigureRemotingForAnsible.ps1'
```

However, if the deployment is running behind a proxy, you may need to make this script available locally.

1. Download the script:

```
wget https://github.com/ansible/ansible/blob/devel/examples/scripts/ConfigureRemotingForAnsible.ps1
```

2. Deploy a local HTTP server, enabling port 80, for example:

```
yum install httpd
systemctl enable httpd
systemctl start httpd
firewall-cmd --permanent --add-port 80/tcp --zone=public
firewall-cmd --reload
```

3. Copy the downloaded script to the web server:

```
cp ConfigureRemotingForAnsible.ps1 /var/www/html
```

4. Configure the variable to point at the local web server, for example,

```
windows_winrm_script: 'http://10.10.174.230/ConfigureRemotingForAnsible.ps1'
```

Windows operating system and Docker EE

Docker Enterprise Edition for Windows Server (Docker EE) enables native Docker containers on Windows Server. This solution has been tested with Windows worker nodes running Windows Server 2016 and with Docker EE 18.09. More recent versions of Windows Server may work but have not been tested.



Note

Docker Universal Control Plane is not currently supported on Windows Server 1709 due to image incompatibility issues. For more information, see the Docker documentation [Install Docker Enterprise Edition for Windows Server](#).

This solution recommends that you only run Windows Server 2016 on your Windows worker nodes and that you install any required updates to your Windows nodes in a timely manner.

For information on how to update Docker EE on Windows Server 2016, see the Docker documentation [Update Docker EE](#).

Deploying Sysdig monitoring

By default, the playbooks for deploying Sysdig are commented out in `site.yml` and must be explicitly enabled in that file if you want it included in the initial deployment. Alternatively, you can run the specific playbooks detailed in this section in a stand-alone manner, subsequent to the initial deployment.

Note

By default, you must have outgoing port 6666 open in your firewall, to allow data to flow to `collector.sysdigcloud.com`. You can configure the agent to use a different port by using the variable `sysdig_collector_port` in `group_vars/vars`.

If you are using a proxy, it must be configured to be "fully-transparent". Non-transparent proxies will not allow the agent to connect.

Monitoring with Sysdig

Sysdig's approach to Docker monitoring uses transparent instrumentation to see inside containers from the outside, with no need for agents in each container. Metrics from Docker containers, and from your applications running inside them, are aggregated in real-time across each service to provide meaningful monitoring dashboards and alerts for your application. Figure 16 provides an overview of the Sysdig architecture.

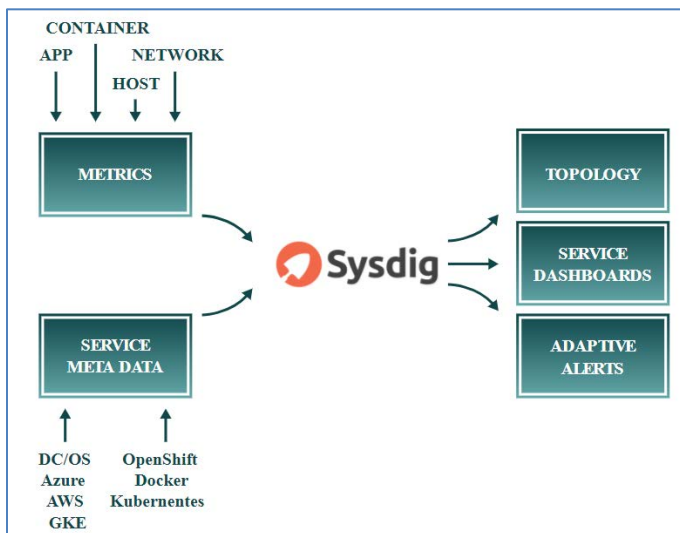


Figure 16. Sysdig architecture

Sysdig Monitor allows you to analyze response times, application performance metrics, container and server utilization metrics, and network metrics. You can build dashboards across applications, micro-services, container and networks, and explore metadata from Docker, Kubernetes, Mesos and AWS. For more information, see the [Sysdig Container Monitoring](#) video overview and the [Sysdig Monitor 101](#) training course.



Sysdig Secure provides security at the orchestrator as well as the container level. You create service-aware policies that allow you to take actions (like killing a container) or send alerts (to Slack, Splunk, etc) whenever a policy violation occurs. All commands are audited to help you identify anomalous actions, along with taking snapshots of all activities pre-and-post a policy violation. For more information, see the [Sysdig Secure](#) video overview and the [Sysdig Secure 101](#) training course.

The implementation in this solution uses the Software as a Service (SaaS) version of Sysdig. The playbooks deploy Sysdig Agent software on each UCP, DTR and Linux worker node, as well as the NFS, logger and load balancer VMs and captured data is relayed back to your Sysdig SaaS Cloud portal. The deployment provides access to a 90 day try-and-buy, fully featured version of the Sysdig software.

Note

The Sysdig functionality is not turned on by default in this solution - see the section on [Sysdig configuration](#) for more information on how to enable Sysdig. For more information on how to access the 90 day try-and-buy version, see the GitHub repository at <https://github.com/HewlettPackard/Docker-SimpliVity>.

Playbooks for installing Sysdig on RHEL

The following playbooks are used when deploying Sysdig:

- `playbooks/sysdig-k8s-rbac.yml` is used to configure Sysdig for Kubernetes.
- `playbooks/install-sysdig.yml` is used to configure Sysdig for Docker swarm. It opens the required port in the firewall, and installs the latest version of the Sysdig agent image on the nodes. By default, this playbook is commented out in `site.yml`, so if you want to use the solution to automatically configure Sysdig for Docker swarm, you must uncomment this line.

Sysdig configuration

Separate playbooks are used to install Sysdig for Docker swarm and Sysdig for Kubernetes.

Sysdig configuration for Docker swarm

The playbook `playbooks/install-sysdig.yml` is used to automate the configuration of the SaaS setup for Docker swarm. By default, this playbook is commented out in `site.yml` and must be explicitly enabled. The variables used to configure Sysdig for Docker swarm are detailed in Table 19.

Table 19. Sysdig variables for Docker swarm

Variable	File	Description
<code>sysdig_access_key</code>	<code>group_vars/vault</code>	After the activation of your account on the Sysdig portal, you will be provided with your access key. This is used by the playbooks to install the agent on each UCP, DTR and Linux worker node, as well as the NFS, logger and load balancer VMs.
<code>sysdig_agent</code>	<code>group_vars/vars</code>	Specifies the URL to the Sysdig Linux native install agent, for example, <code>https://s3.amazonaws.com/download.draios.com/stable/install-agent</code>
<code>sysdig_tags</code>	<code>group_vars/vars</code>	Tagging your hosts is highly recommended. Tags allow you to sort the nodes of your infrastructure into custom groups in Sysdig Monitor. Specify location, role, and owner in the format: <code>'location:City,role:Express Containers,owner:Customer Name'</code>

Sysdig configuration for Kubernetes

The playbook `playbooks/sysdig-k8s-rbac.yml` is used to automate the configuration of the SaaS setup for Kubernetes. The variables used to configure Sysdig for Kubernetes are detailed in Table 20.

Table 20. Sysdig variables for Kubernetes

Variable	File	Description
<code>sysdig_access_key</code>	<code>group_vars/vault</code>	After the activation of your account on the Sysdig portal, you will be provided with your access key. This is used by the playbooks to install the agent on each UCP, DTR and Linux Kubernetes worker nodes.



sysdig_collector	group_vars/vars	The URL for the Sysdig SaaS, by default, 'collector.sysdigcloud.com'
sysdig_collector_port	group_vars/vars	The port used by the agent, by default, '6666'
sysdig_tags	group_vars/vars	Tagging your hosts is highly recommended. Tags allow you to sort the nodes of your infrastructure into custom groups in Sysdig Monitor. Specify location, role, and owner in the format: 'location:City,role:Express Containers,owner:Customer Name'
k8s_cluster	group_vars/vars	<div>This should match the cluster name displayed when you source the environment setup script, for example:</div> <div># source env.sh Cluster "ucp_hpe-ucp.cloudra.local:6443_admin" set. User "ucp_hpe-ucp.cloudra.local:6443_admin" set.</div> <div>For more information, see the section on installing the UCP client bundle in the section Deploying Sysdig monitoring on Kubernetes,</div>

Registering for Sysdig trial

Hewlett Packard Enterprise has teamed up with Sysdig to offer a fully featured 90-day trial version of Sysdig Monitor and Secure as part of the HPE Express Containers with Docker Enterprise Edition on HPE SimpliVity solution. For more details on how to sign up, see the GitHub repository at <https://github.com/HewlettPackard/Docker-SimpliVity>.

After registering for the trial, you will be presented with options for setting up your environment, as shown in Figure 17.

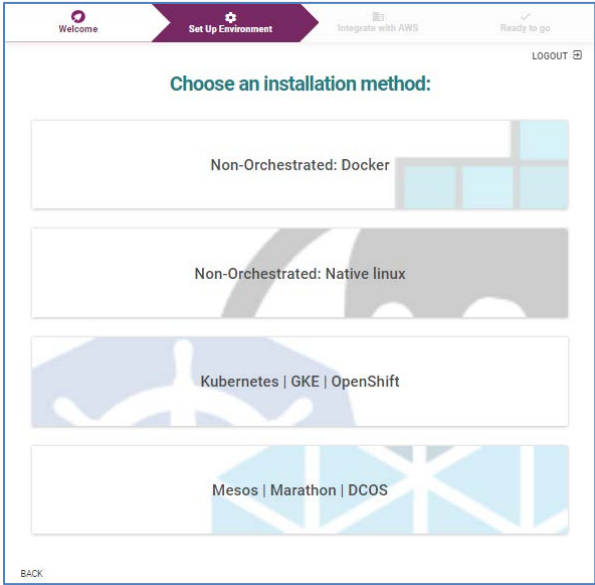


Figure 17. Sysdig Monitor set up environment

Sysdig Monitoring for Kubernetes

If you are deploying Sysdig monitoring on Kubernetes, select the **Kubernetes | GKE | OpenShift** option. You will be presented with an access code, as shown in Figure 18.



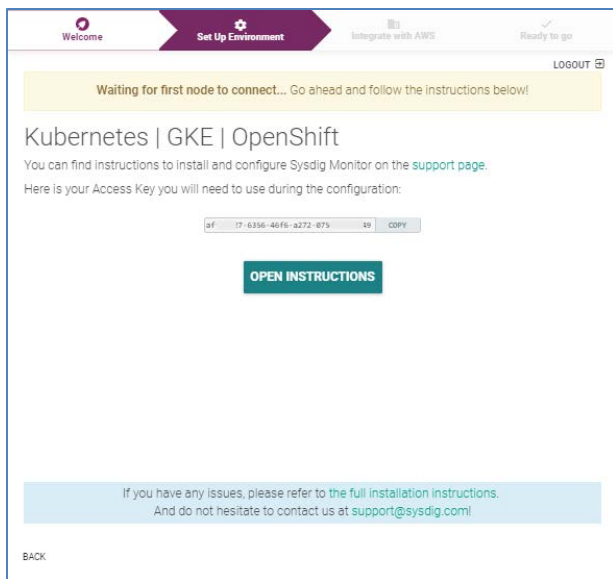


Figure 18. Sysdig Monitor access code for Kubernetes

Use the `sysdig_access_key` field in your `group_vars/vault`, as described in the section Sysdig configuration for Kubernetes. Once you deploy your environment and your Kubernetes nodes connect to the Sysdig SaaS platform, Sysdig will automatically display information regarding your setup, as shown in Figure 19.

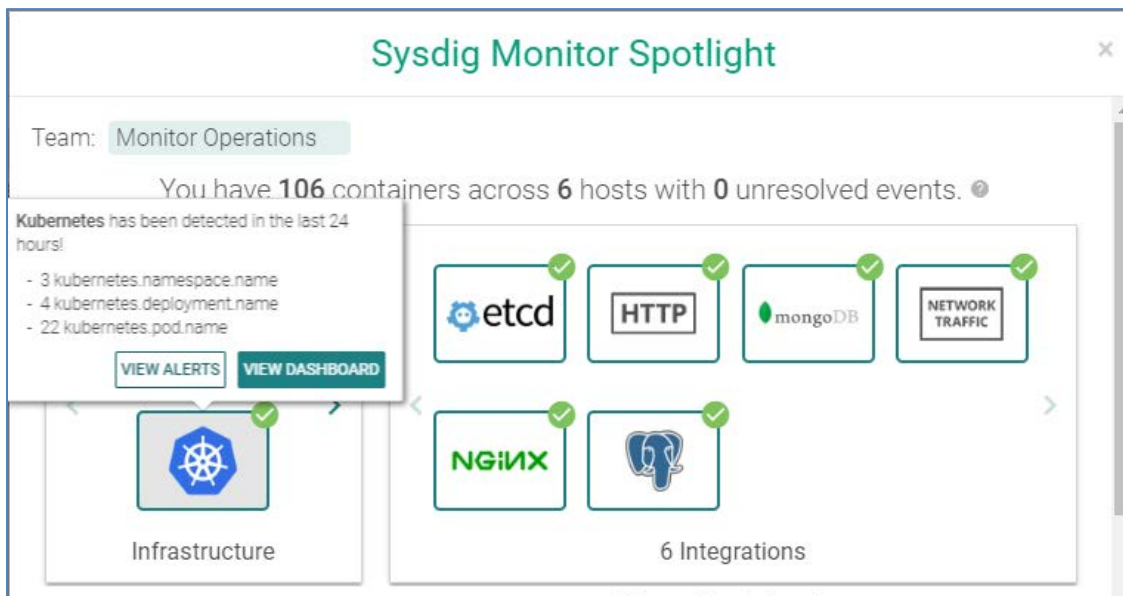


Figure 19. Sysdig Monitor Spotlight for Kubernetes

Select **View Dashboard** for an entry point to accessing all your monitoring data. Alternatively, you can browse to <https://app.sysdigcloud.com> at any time to access your dashboards.



Sysdig Monitor for Docker swarm

If you are deploying Sysdig monitoring on Docker swarm, select the **Non-Orchestrated: Native Linux** option. You will be presented with a screen containing details for the URL to download the Sysdig agent, along with your access code embedded in the command, as shown in Figure 20.

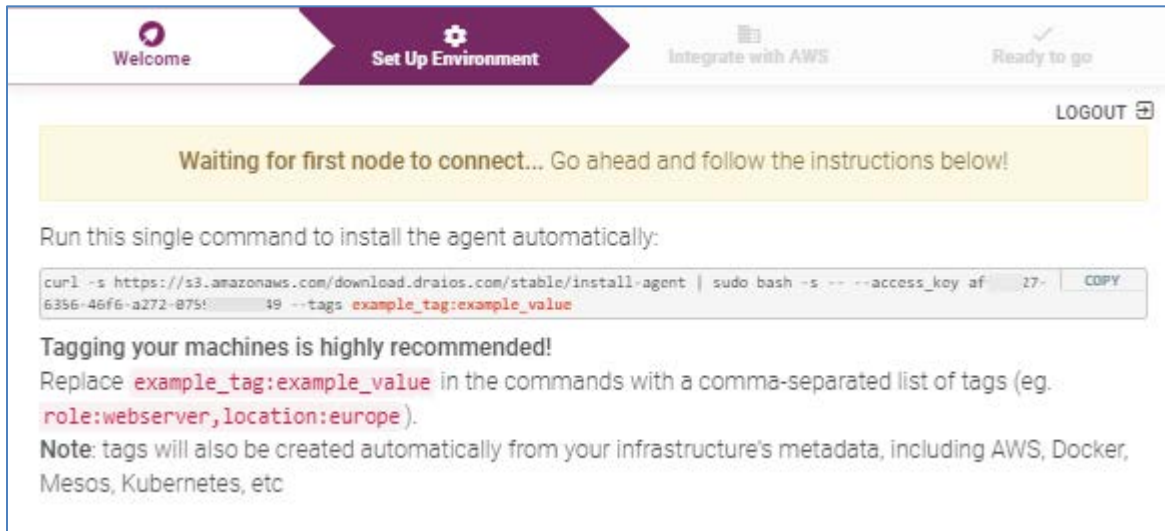


Figure 20. Sysdig Monitor download location and access code for Docker

The download URL is used in the `sysdig_agent` field in `group_vars/vars`, while the access code is stored in the `sysdig_access_key` field in your `group_vars/vault`, as described in the section Sysdig configuration for Docker swarm.

Once you deploy your environment and your Docker swarm nodes connect to the Sysdig SaaS platform, Sysdig will automatically display information regarding your setup. Alternatively, you can browse to <https://app.sysdigcloud.com> at any time to access your dashboards.

Deploying Sysdig monitoring on Kubernetes

The latest version of Sysdig supports monitoring of Kubernetes logs and metrics.

Prerequisites

- Install the kubectl binary on your Ansible box.
- Install the UCP Client bundle for the admin user.
- Confirm that you can connect to the cluster by running a test command, for example, `kubectl get nodes`
- Ensure that you have configured the required variables, as described in the section Sysdig configuration for Kubernetes

For example, you add the relevant variables in the `group_vars/vars` file.

```
sysdig_collector: 'collector.sysdigcloud.com'
sysdig_collector_port: '6666'
sysdig_tags: 'location:Enter city,role:Enter role,owner:Customer name'
k8s_cluster: 'ucp-hpe2-ucp.cloudra.local'
```

You should add the access key to the encrypted `group_vars/vault` using the command `ansible-vault edit group_vars/vault`.

```
sysdig_access_key: '10****97-9160-****-9061-84bfd0f****0'
```

Running the playbook

The playbook `playbooks/k8s-install-sysdig.yml` is used to automate the configuration of the SaaS setup for Docker swarm.



```
# cd Docker-SimpliVity
# ansible-playbook -i vm_hosts playbooks/sysdig-k8s-rbac.yml --vault-password-file .vault_pass
```

Using the Sysdig software as a solution (SaaS) website <https://app.sysdigcloud.com>, you are able to view, analyze and inspect various different dashboards. Initially, you will just see the monitoring information for the infrastructure itself. Deploy a sample application, as detailed in the section_Kubernetes guestbook example with Redis, and use the Sysdig solution to analyze the different facets of the deployed application.

Deploying Sysdig monitoring on Docker Swarm

The playbook `playbooks/install_sysdig.yml` is used to automate the configuration of the SaaS setup for Docker swarm. By default, this playbook is commented out in `site.yml` and must be explicitly enabled. An access key variable must be set in the `group_vars/vault` file as detailed in Table 19.

```
# cd Docker-SimpliVity
# ansible-playbook -i vm_hosts playbooks/install_sysdig.yml --vault-password-file .vault_pass
```

Using the Sysdig software as a solution (SaaS) website <https://app.sysdigcloud.com>, you are able to view, analyze and inspect various different dashboards.



Deploying Splunk

This section provides an overview of Splunk, outlines how to configure and run the relevant playbooks and shows how to access the UI to see the resultant Docker and Kubernetes dashboards.

Monitoring with Splunk

Splunk Enterprise allows you to collect and index any data from any source, and to monitor systems and infrastructure in real time to preempt issues before they happen. It allows you to analyze your data to understand trends, patterns of activity and behavior, giving you valuable intelligence across your entire organization. The solution architecture for Splunk is shown in Figure 21.

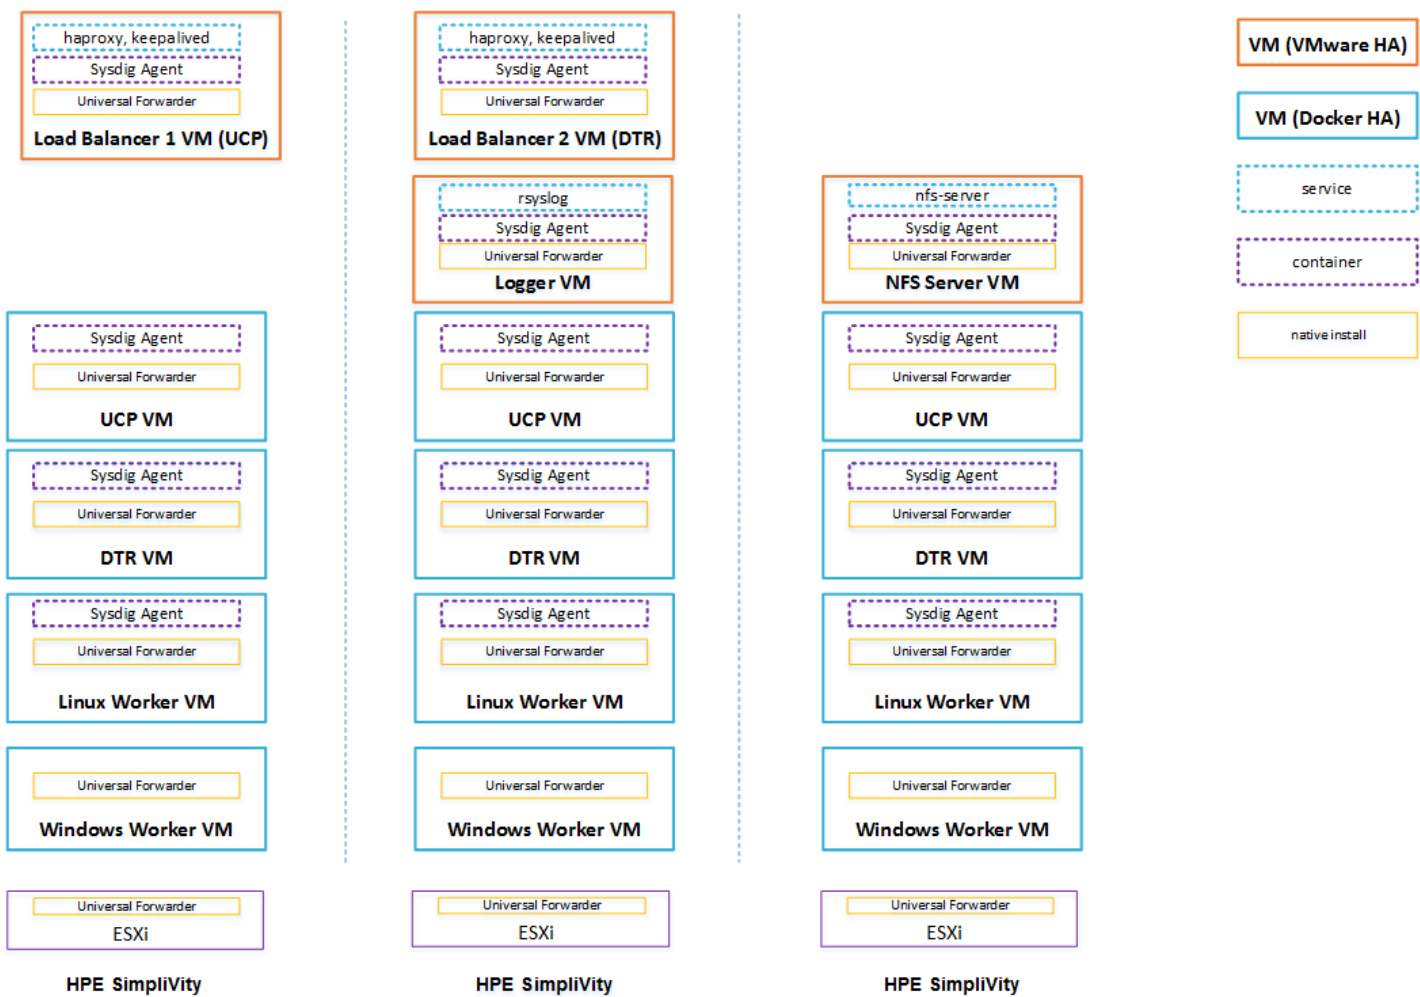


Figure 21. Solution architecture: Hybrid Linux and Windows workers with Splunk and Sysdig

This solution allows you to integrate your CaaS deployment with an existing Splunk Enterprise installation or to deploy a stand-alone Splunk Enterprise demo environment as a Docker stack in your cloud. In both instances, Universal Forwarders are used to collect data from your applications running on your Linux and Windows worker nodes in your cloud, as well as log data from the Docker platform itself and from the infrastructure VMs and servers. Figure 22 shows the Splunk architecture.




Figure 22. Splunk architecture

All the Universal Forwarders run natively on the operating system to allow greater flexibility in terms of configuration options. Each forwarder sends the data it collects to one or more indexers in the central Splunk.

Linux worker nodes: The Universal Forwarders on the Linux worker nodes collect log and metrics data. The log data includes:

- `/var/log/messages` from the Docker host (including the daemon engine logs)
- `/var/log/secure` from the Docker hosts
- container logs via a Splunk technical add-on

The metrics data is collected via a technical add-on and includes:

- `docker stats`
- `docker top`
- `docker events`
- `docker service stats`

Windows worker nodes: The Universal Forwarders running on the Windows worker nodes collect the following data:

- Windows logs
- CPU stats
- Memory stats
- Network Interface stats
- and more

For more information on configuring standalone Splunk for Linux and Windows worker nodes, see the section on [Splunk prerequisites](#).

UCP and ESXi: UCP operational logs and ESXi logs are forwarded to the logger VM via TCP ports 1514 and 514 respectively. Port 1514 is assigned a special `sourcetype` of `ucp` which is then used by the Splunk Docker APP to interpret UCP logs. The Universal Forwarder runs the `rsyslog` daemon which will record the log messages coming from the ESX machines into the `/var/log/messages` file on the VM.

Non-Docker VMs: Other VMs, for example, NFS, use a Splunk `monitor` to collect and forward data from the following files:

- `/var/log/messages`
- `/var/log/secure` (Red Hat)

Note

You can configure the list of files monitored by the Universal Forwarder.

Other syslog senders can be configured to send their data to the logger VM or directly to central Splunk.

Playbooks for installing Splunk

The following playbooks are used to install Splunk.

- `playbooks/splunk_demo.yml` installs a demo of Splunk Enterprise in the cluster (if the `splunk_demo` deployment option is selected. A value of `splunk` is used to configure an external production Splunk deployment.)
- `playbooks/splunk_uf.yml` installs and configures the Splunk Universal Forwarder on each Linux and Windows node in the inventory

Splunk configuration

This solution supports two types of Splunk deployment. Firstly, there is a built-in deployment useful for demos and for getting up to speed with Splunk. Alternatively, the solution can be configured to interact with a standalone, production Splunk deployment that you set up independently. In this case, you must explicitly configure the universal forwarders with external "forward servers" (Splunk indexers), whereas this happens automatically with the built-in option.

In the standalone deployment, you can enable SSL authentication between the universal forwarders and the indexers, by setting the `splunk_ssl` variable to `yes` in the file `group_vars/vars`. The built-in demo deployment does not support SSL and so, in this instance, the value of the `splunk_ssl` variable is ignored. For more information on enabling SSL, see Appendix C.

After the installation is complete, the Splunk UI can be reached at `http://<fqdn>:8000`, where `<fqdn>` is the FQDN of one of your Linux Docker nodes. Mesh routing does not currently work on Windows so you must use a Linux node to access the UI.

Splunk prerequisites

You should select the Splunk deployment type that you require by setting the variable `monitoring_stack` in the `group_vars/vars` file to either `splunk`, to use a standalone Splunk deployment, or `splunk_demo` for the built-in version. If you omit this variable, or if it has an invalid value, no Splunk deployment will be configured.

For both types of deployment, you need to download the Splunk universal forwarder images/packages from https://www.splunk.com/en_us/download/universal-forwarder.html. Packages are available for 64-bit Linux and 64-bit Windows 8.1/Windows 10. Download the RPM package for Linux 64-bit (2.6+ kernel Linux distributions) to `./files/splunk/linux`. If you are deploying Windows



nodes, download the MSI package for Windows 64 bit to `./files/splunk/windows`. For a dual Linux/Windows deployment, the images and packages must have same name and version, along with the appropriate extensions, for example:

- `files/splunk/windows/splunkforwarder-7.1.2.msi`
- `files/splunk/linux/splunkforwarder-7.1.2.rpm`

You need to set the variable `splunk_architecture_universal_forwarder_package` to the name you selected for the package(s), not including the file extension. Depending on the Splunk deployment you have chosen, edit the file `templates/monitoring/splunk/vars.yml` or the file `templates/monitoring/splunk_demo/vars.yml` and set the variable, for example:

```
splunk_architecture_universal_forwarder_package: 'splunkforwarder-7.1.2'
```

As of Splunk version 7.1, the Splunk universal forwarder must be deployed with a password. This password is specified using the variable `splunk_uf_password` which is configured in `group_vars/vault`.

If you are using a standalone Splunk deployment, you must specify the list of indexers using the variable `splunk_architecture_forward_servers` in `group_vars/vars`, for example:

```
splunk_architecture_forward_servers:
- splunk-indexer1.cloudra.local:9997
- splunk-indexer2.cloudra.local:9997
```

By default, the indexers are configured in a single load balancing group. This can be changed by editing the file `outputs.conf.j2` in the folder `template/monitoring/splunk/`. For more information on forwarding using Universal Forwarder, see the Splunk documentation at <http://docs.splunk.com/Documentation/Forwarder/7.0.2/Forwarder/Configureforwardingwithoutoutputs.conf>.

On your standalone Splunk installation, you need to install the following add-ons and apps.

To monitor **Linux worker nodes**, the **Docker app** should be installed on central Splunk. More info on this Docker app can be found at <https://github.com/splunk/docker-itmonitoring> and at <https://hub.docker.com/r/splunk/universalforwarder/>.

To monitor the **Windows worker nodes**, install the **Splunk App for Windows Infrastructure** on central Splunk and its dependencies:

- Splunk App for Windows Infrastructure. The Splunk App for Windows Infrastructure is not compatible with The Splunk Add-on for Windows 5.0 at this time. See <https://splunkbase.splunk.com/app/1680/>
- Splunk Add-on for Microsoft Windows version 4.8.4 - see <https://splunkbase.splunk.com/app/742/>
- Splunk Add-On for Microsoft Active Directory version 1.0.0 - see <https://splunkbase.splunk.com/app/3207/>
- Splunk Add-on for Microsoft Windows DNS version 1.0.1 (if this is not installed on central Splunk, you will see yellow icons on some dashboards with the message `eventtype wineventlog-dns does not exist or is disabled`) - see <https://splunkbase.splunk.com/app/3208/>
- Splunk Supporting Add-on for Active Directory version 2.1.7 (if this is not installed on central Splunk, you will see yellow icons on some dashboards with the message `eventtype wineventlog-ds does not exist or is disabled`) - see <https://splunkbase.splunk.com/app/1151/>

If you want to use your own certificates in your standalone Splunk deployment to secure the communications between the indexers and the universal forwarders, see Appendix D.

You can specify advanced Splunk configuration in the following files:

- `files/splunk/linux/SPLUNK_HOME`
- `files/splunk/linux/DOCKER_TAS`
- `files/splunk/windows/SPLUNK_HOME`



These files will be copied as-is to the systems running the universal forwarder.

Configuring syslog in UCP

In order to see some data in the UCP operational dashboard, you need to have UCP send its logs to the VM configured in the [logger] group. For example, for the following `vm_host` file:

```
[logger]
hpe-logger ip_addr='10.60.59.24/16' esxi_host='esxi-hpe-2.cloudra.local'
```

This will configure UCP to send its logs to `hpe-logger.cloudra.local:1514`. You need to select the TCP protocol as shown in Figure 23.




Figure 23. Configure Remote Syslog Server in UCP



Configuring syslog in ESX

This configuration must be done manually for each ESX server. The syslog server should be the server configured in the [logger] group in your vm_hosts inventory. The protocol should be tcp and the port 514 as shown in Figure 24.




Figure 24. Configure Syslog on ESXi Hosts

For more information, see the VMware documentation at <https://docs.vmware.com/en/VMware-vSphere/6.5/com.vmware.vsphere.security.doc/GUID-9F67DB52-F469-451F-B6C8-DAE8D95976E7.html>.

Limitations

- The Dockerized Splunk App has a number of open issues
 - <https://github.com/splunk/docker-itmonitoring/issues/19>
 - <https://github.com/splunk/docker-itmonitoring/issues/20>
- The Docker events tab is not working

Accessing Splunk UI

After the installation is complete, the Splunk UI can be reached at `http://<fqdn>:8000`, where <fqdn> is the FQDN of one of your Linux Docker nodes. Mesh routing does not currently work on Windows so you must use a Linux node to access the UI. For example:

`http://hpe-ucp01.am2.cloudra.local:8000/`

The default username and password for Splunk is `admin / changeme`.



Use the **Docker App** to view the Docker overview as shown in Figure 25 and the Docker stats as shown in Figure 26.

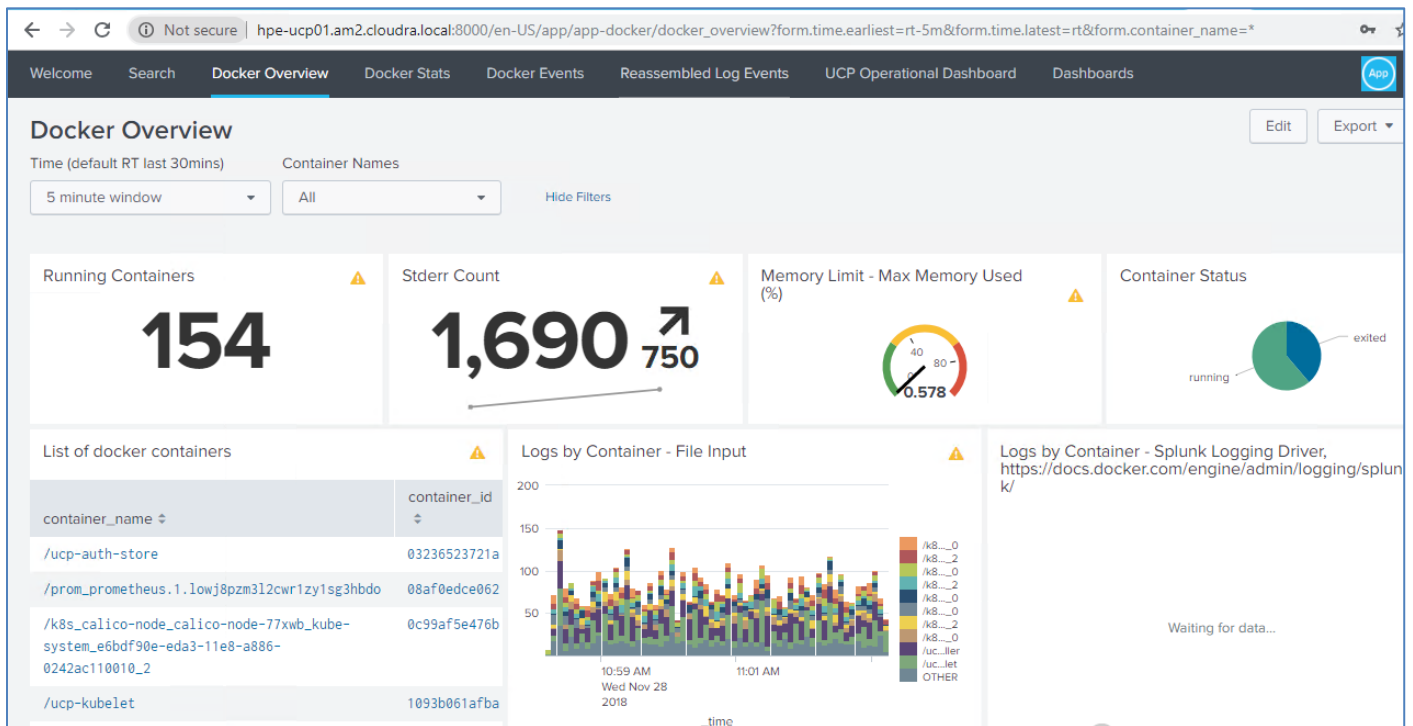


Figure 25. Docker overview

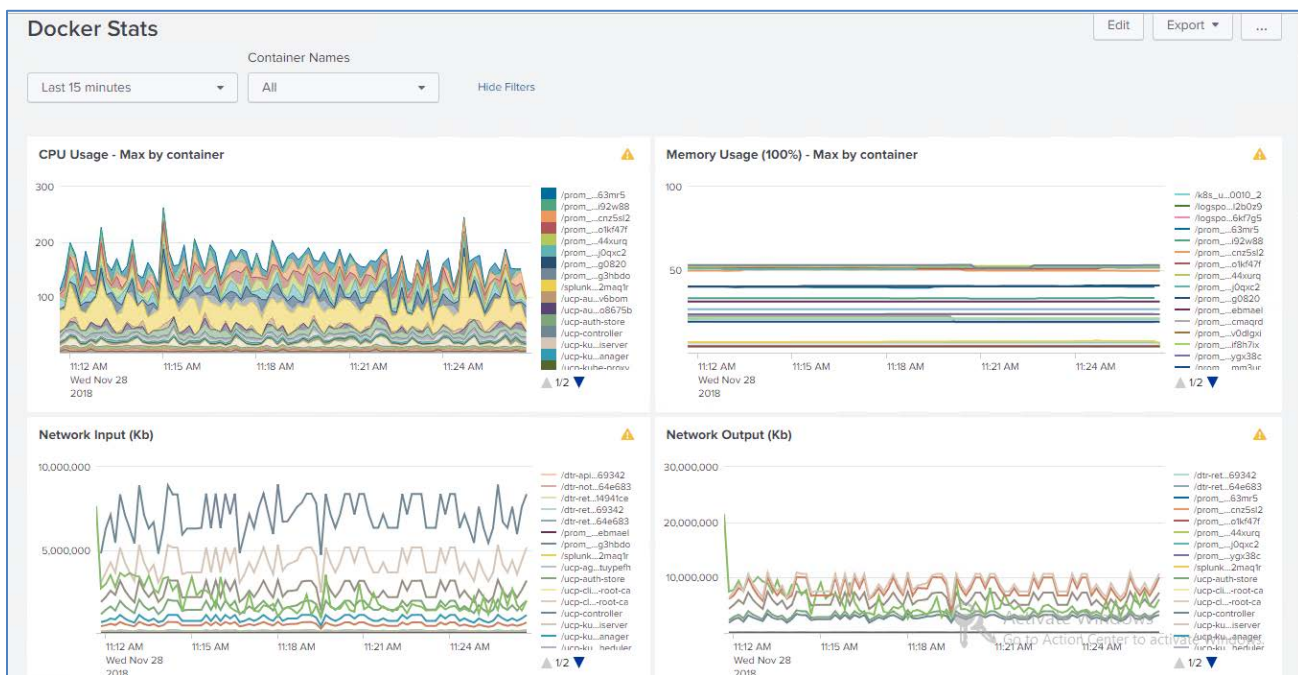


Figure 26. Docker stats

Use the **k8s App** to see the Kubernetes overview as shown in Figure 27 and then access the details for deployments, daemon sets, replica sets, services, etc.




Figure 27. Kubernetes overview

Redeploying Splunk demo

The Splunk demo deployment, whilst fully featured, is **severely** restricted in the amount of data it can process. Once this limit has been reached, often after running for just one or two days, it is necessary to redeploy the application if you want to continue experimenting with the demo.

Before you redeploy, it is necessary to remove the corresponding Docker stack and delete the associated volumes.

```
# ssh hpe-ucp02
```

```
# docker stack rm splunk_demo
Removing service splunk_demo_splunkenterprise
Removing network splunk_demo_default
```

```
# docker volume ls | grep splunk
vsphere:latest      splunk_demo_vsplunk-opt-splunk-etc@Docker_HPE
vsphere:latest      splunk_demo_vsplunk-opt-splunk-var@Docker_HPE
```

```
# docker volume rm splunk_demo_vsplunk-opt-splunk-etc@Docker_HPE
splunk_demo_vsplunk-opt-splunk-etc@Docker_HPE
```

```
# docker volume rm splunk_demo_vsplunk-opt-splunk-var@Docker_HPE
splunk_demo_vsplunk-opt-splunk-var@Docker_HPE
```

Then re-run the playbook on your ansible node.

```
ansible-playbook -i vm_hosts playbooks/splunk_demo.yml --vault-password-file .vault_pass
```



Deploying Prometheus and Grafana on Kubernetes

Monitoring Kubernetes with Prometheus and Grafana

Monitoring a Kubernetes cluster with Prometheus is a natural choice as Kubernetes components themselves are instrumented with Prometheus metrics, therefore those components simply have to be discovered by Prometheus and most of the cluster is monitored.

The solution uses the Prometheus Operator to deploy Prometheus and Grafana. The playbooks install `kube-state-metrics` and `node-exporter` components, as well as supporting `kubelet` and `apiserver` metrics. Sample dashboards for Grafana are installed to help you monitor your Kubernetes infrastructure.

The Prometheus Operator, shown in Figure 28, makes running Prometheus on top of Kubernetes as easy as possible, while preserving Kubernetes-native configuration options. It introduces additional resources in Kubernetes to declare the desired state and configuration of Prometheus. The `Prometheus` resource declaratively describes the desired state of a Prometheus deployment, while a `ServiceMonitor` describes the set of targets to be monitored by Prometheus.




Figure 28. Prometheus Operator

Playbooks for installing Prometheus and Grafana on Kubernetes

Prerequisites

Before you run the playbook to install Prometheus and Grafana on Kubernetes, you need to ensure that you have already downloaded and installed `kubectl` and set up your client bundle. Two convenience playbooks have been provided to make this process easier.

The playbook `playbooks/install-kubectl.yml` installs a specific version of `kubectl` based on the settings in your `group_vars/vars` file.



The playbook `playbooks/kube-prometheus.yml` is used to deploy the Prometheus/Grafana stack on Kubernetes. It is a wrapper for a number of separate playbooks outlined below.

- `playbooks/kube-prometheus/operator.yml`
- `playbooks/kube-prometheus/kube-state-metrics.yml`
- `playbooks/kube-prometheus/node-exporter.yml`
- `playbooks/kube-prometheus/monitors.yml`
- `playbooks/kube-prometheus/prometheus.yml`
- `playbooks/kube-prometheus/grafana.yml`

You can choose not to install certain components by commenting out the appropriate line in the wrapper playbook.

Prometheus Operator

The Prometheus Operator makes running Prometheus on top of Kubernetes as easy as possible, while preserving Kubernetes-native configuration options. For more information on Prometheus Operator, see <https://coreos.com/operators/prometheus/docs/latest/user-guides/getting-started.html>.

The playbook `playbooks/kube-prometheus/operator.yml` installs the operator itself.

Kube state metrics

`kube-state-metrics` is a simple service that listens to the Kubernetes API server and generates metrics about the state of the objects. It is not focused on the health of the individual Kubernetes components, but rather on the health of the various objects inside, such as deployments, nodes and pods. For more information on kube-state-metrics, see <https://github.com/kubernetes/kube-state-metrics>.

The playbook `playbooks/kube-prometheus/kube-state-metrics.yml` installs kube-state-metrics on all UCP, DTR and Kubernetes worker nodes.

Node exporter

The node-exporter provides an overview of cluster node resources including CPU, memory and disk utilization and more. For more information on node-exporter, see https://github.com/prometheus/node_exporter.

The playbook `playbooks/kube-prometheus/node-exporter.yml` installs `node-exporter` as a set of Docker containers on all UCP, DTR and Kubernetes worker nodes. Port 9100 is opened in the firewall on each node where it is installed.

Monitors

While all the other Kubernetes components run on top of Kubernetes itself, `kubelet` and `apiserver` do not, and so they just need service monitors to access these metrics.

The playbook `playbooks/kube-prometheus/monitors.yml` installs Service Monitors for `kubelet` and `apiserver`

cAdvisor

Support for cAdvisor is built-in to Kubernetes, so cAdvisor metrics will automatically be available within Prometheus, without any other configuration required.

Note

Because Docker EE provides a hosted version of Kubernetes, it is not possible to access metrics for `kube-scheduler` and `kube-controller-manager`.

Prometheus

For convenience, the playbook sets up a NodePort so that the Prometheus UI can be accessed on port 33090.

```
# kubectl -n monitoring patch svc prometheus-k8s --type='json' -p
'[{ "op": "replace", "path": "/spec/type", "value": "NodePort" } ]'
```



```
# kubectl -n monitoring patch svc prometheus-k8s --type='json' -p '[{"op": "add",
"path":"/spec/ports/0/nodePort", "value":33090}]'
```

On a production system, it is likely that you will want to remove this NodePort. The following code segment shows how to use the `patch` command to remove the NodePort.

```
# kubectl -n monitoring patch svc prometheus-k8s --type='json' -p '[{"op": "remove",
"path":"/spec/ports/0/nodePort"}]'
```

```
# kubectl -n monitoring patch svc prometheus-k8s --type='json' -p '[{"op": "remove",
"path":"/spec/type"}]'
```

Grafana

For convenience, the playbook sets up a NodePort so that the Grafana UI can be access on the port 33030.

```
# kubectl -n monitoring patch svc grafana --type='json' -p '[{"op":"replace","path":"/spec/type",
"value":"NodePort"}]'
```

```
# kubectl -n monitoring patch svc grafana --type='json' -p '[{"op": "add",
"path":"/spec/ports/0/nodePort", "value":33030}]'
```

On a production system, it is likely that you will want to remove this NodePort. The following code segment shows how to use the `patch` command to remove the NodePort.

```
# kubectl -n monitoring patch svc grafana --type='json' -p '[{"op": "remove",
"path":"/spec/ports/0/nodePort"}]'
```

```
# kubectl -n monitoring patch svc grafana --type='json' -p '[{"op": "remove", "path":"/spec/type"}]'
```

Teardown

The playbook `playbooks/kube-prometheus-teardown.yml` removes the installed Prometheus\Grafana stack.

Prometheus UI

The Prometheus UI is available via your UCP, DTR or Kubernetes worker nodes, using HTTP on port 33090, for example,

```
http://hpe-ucp01.am2.cloudra.local:33090
```

To see what services are being monitored, access the service discovery page, via `Status -> Service Discovery`, or using the `/service-discovery` endpoint:

```
http://hpe2-ucp01.am2.cloudra.local:33090/service-discovery
```

The monitored services are listed as shown in Figure 29.





Figure 29. Prometheus service discovery

To see the status for the monitored services, access the targets page via **Status -> Targets** or using the endpoint `/targets`.
`http://hpe2-ucp01.am2.cloudra.local:33090/targets`

The status of the various monitors are displayed, as shown in Figure 30.




Figure 30. Prometheus targets



To see all the metrics available, click on **Graph** or use the endpoint `/graph`:

`http://hpe2-ucp01.am2.cloudra.local:33090/graph`

Click on the drop-down titled “- insert metric at cursor -” to see all the metrics that are available to Prometheus as shown in Figure 31.




Figure 31. Prometheus metrics

Node Exporter

Metrics specific to the Node Exporter are prefixed with `node_` and include metrics like `node_cpu_seconds_total` and `node_exporter_build_info`. Table 21 below lists some example expressions.

Table 21. Sample Node Exporter metrics

Metric	Meaning
rate(node_cpu_seconds_total{mode="system"}[1m])	The average amount of CPU time spent in system mode, per second, over the last minute (in seconds)
node_filesystem_avail_bytes	The filesystem space available to non-root users (in bytes)
rate(node_network_receive_bytes_total[1m])	The average network traffic received, per second, over the last minute (in bytes)



More information on the use of `node-exporter` metrics is available at https://github.com/prometheus/node_exporter.

cAdvisor

cAdvisor is an open source container resource usage and performance analysis agent. It is purpose-built for containers and supports Docker containers natively. In Kubernetes, cAdvisor is integrated into the Kubelet binary. cAdvisor auto-discovers all containers in the machine and collects CPU, memory, filesystem, and network usage statistics. cAdvisor also provides the overall machine usage by analyzing the 'root' container on the machine.

Kubelet exposes a simple cAdvisor UI for containers on a machine, via the default port **4194**. However, this feature has been marked deprecated in v1.10 and completely removed in v1.12. For more information on how upcoming releases will reduce the set of metrics exposed by the `kubelet`, see the relevant issue page at <https://github.com/kubernetes/kubernetes/issues/68522>.

The Kubelet also starts an internal HTTP server on port 10255 and exposes endpoints including `/metrics` and `/metrics/cadvisor`. As this release of Express Containers uses Kubernetes 1.11, it is able to use this feature. In future releases, it will be necessary to deploy cAdvisor as a DaemonSet for access to the cAdvisor UI.

Table 22 lists some example cAdvisor expressions.

Table 22. Sample cAdvisor metrics

Expression	Description	For
<code>rate(container_cpu_usage_seconds_total{name="redis"}[1m])</code>	The cgroup's CPU usage in the last minute (split up by core)	The <code>redis</code> container
<code>container_memory_usage_bytes{name="redis"}</code>	The cgroup's total memory usage (in bytes)	The <code>redis</code> container
<code>rate(container_network_transmit_bytes_total[1m])</code>	Bytes transmitted over the network by the container per second in the last minute	All containers
<code>rate(container_network_receive_bytes_total[1m])</code>	Bytes received over the network by the container per second in the last minute	All containers

A full listing of cAdvisor-gathered container metrics exposed to Prometheus can be found in the cAdvisor documentation at <https://github.com/google/cadvisor/blob/master/docs/storage/prometheus.md>.

Grafana UI

The Grafana UI is available via your UCP, DTR or Kubernetes worker nodes, using HTTP on port **33030**, for example,

`http://hpe-ucp01.am2.cloudra.local:33030`

The default username and password for Grafana is `admin/admin`. The first time you login, you will be asked to reset the default `admin` password.

A number of dashboards are installed by default. The following figures illustrate some of the dashboard provided.



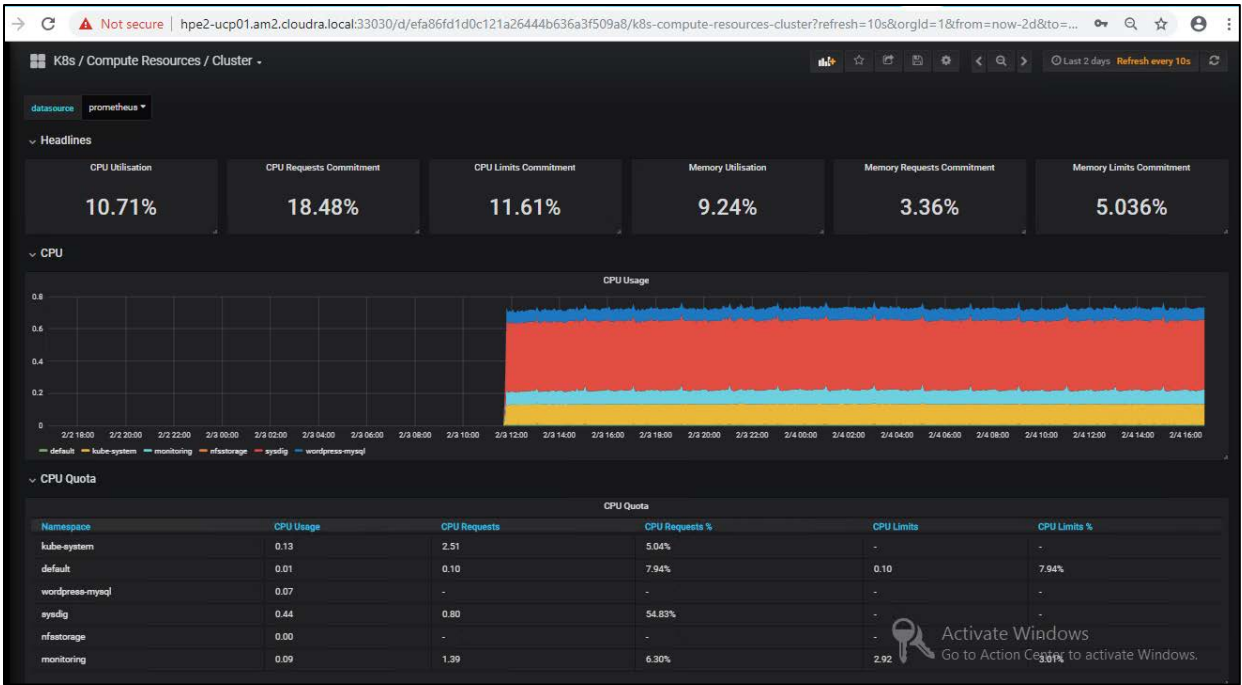


Figure 32. Compute resources dashboard

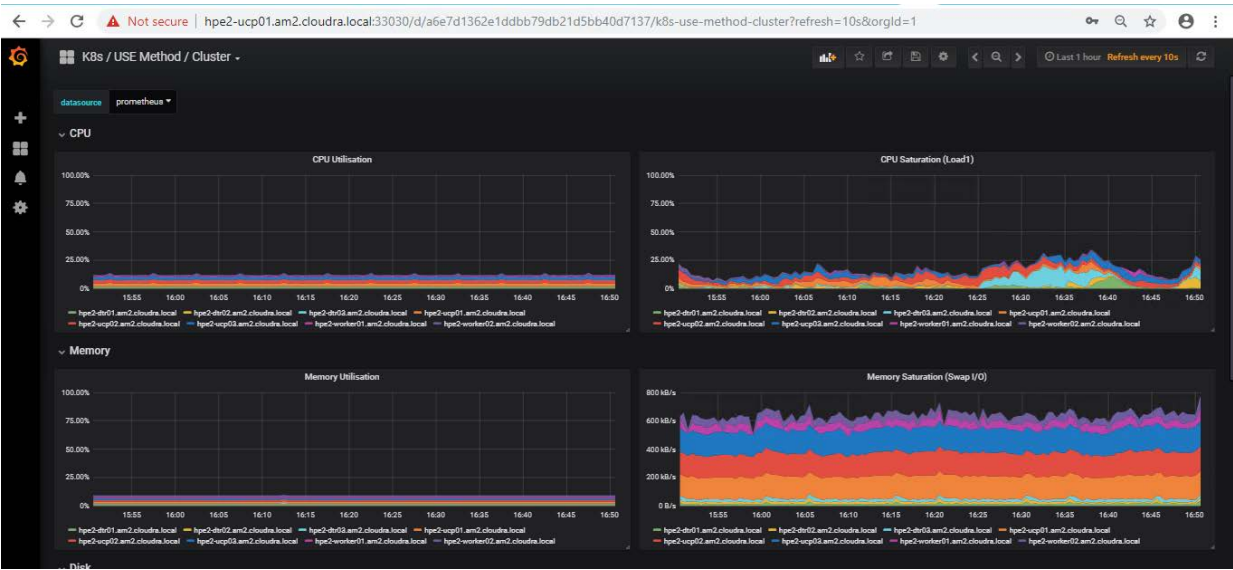


Figure 33. USE method cluster dashboard






Figure 34. USE method node dashboard

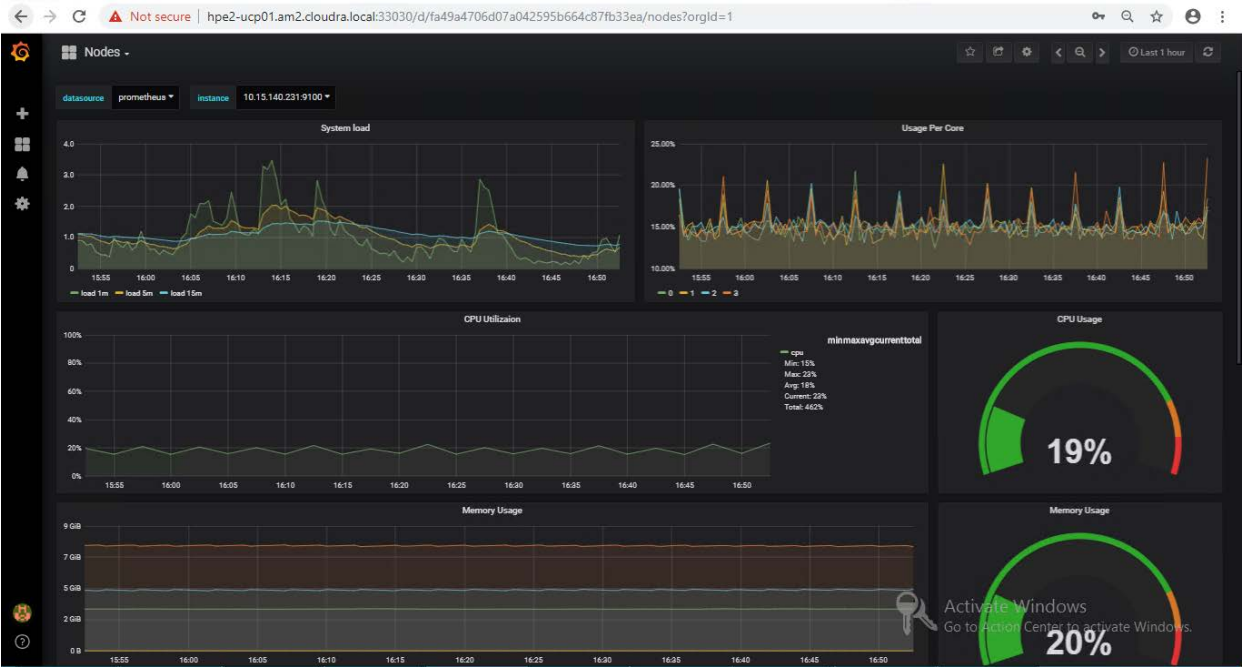


Figure 35. Nodes dashboard



Deploying Prometheus and Grafana on Docker swarm

Monitoring with Prometheus and Grafana

The solution can be configured to enable the use of Prometheus and Grafana for monitoring. In this setup, there is no need for native installs and all the required monitoring software runs in containers, deployed as either services or stacks. The load among the three hosts will be shared as per Figure 36.



Figure 36. Solution architecture: Linux workers with Prometheus and Grafana

The Prometheus and Grafana services are declared in a Docker stack as replicated services with one replica each, so that if they fail, Docker EE will ensure that they are restarted on one of the UCP VMs. `cAdvisor` and `node-exporter` are declared in the same stack as global services, so Docker EE will ensure that there is always one copy of each running on every machine in the cluster.

Note

Prometheus and Grafana functionality is not turned on by default in this solution - see the section on [Prometheus and Grafana configuration](#) for more information on how to enable these tools. Additionally, this functionality will not work for the Windows worker nodes in your environment at present.



Playbooks for installing Prometheus and Grafana

The following playbooks are used to deploy Prometheus and Grafana on Docker RHEL nodes.

- `playbooks/install_logspout.yml` installs and configures **Logspout** on all Docker nodes. Logspout is responsible for sending logs produced by containers running on the Docker nodes to the central logger VM. By default, this playbook is commented out in `site.yml`.
- `playbooks/config_monitoring.yml` configures a monitoring system for the Docker environment based on Grafana, Prometheus, cAdvisor and node-exporter Docker containers. By default, this playbook is commented out in `site.yml`, so if you want to use the solution to automatically deploy a Prometheus/Grafana monitoring system, you must explicitly uncomment both this and the `playbooks/install_logspout.yml` playbook.

Prometheus and Grafana configuration

All monitoring-related variables for Prometheus and Grafana are described in Table 23. The variables determine the versions of various software tools that are used and it is recommended that the values given below are used.

Table 23. Monitoring variables

Variable	Description
<code>cadvisor_version</code>	<code>v0.28.3</code>
<code>node_exporter_version</code>	<code>v1.15.0</code>
<code>prometheus_version</code>	<code>V2.3.2</code>
<code>grafana_version</code>	<code>5.2.3</code>
<code>prom_persistent_vol_name</code>	The name of the volume which will be used to store the monitoring data. The volume is created using the vSphere Docker Volume plugin.
<code>prom_persistent_vol_size</code>	The size of the volume which will hold the monitoring data. The exact syntax is dictated by the vSphere Docker Volume plugin. The default value is 10GB.

Accessing Grafana UI

The Grafana UI is available at the UCP VIP, using HTTP on port 3000, for example,

`http://hpe-ucpvip.am2.cloudra.local:3000`



The default username and password for Grafana is **admin/admin**. The first time you login, you will be asked to reset the default **admin** password. Select the Docker Swarm Monitor dashboard that has already been loaded by the playbooks, as shown in Figure 37 and Figure 38.




Figure 37. Docker Swarm Monitor dashboard




Figure 38. Docker Swarm Monitor dashboard



Backup and restore

This Reference Configuration provides playbooks and scripts to help you back up and restore:

- Docker UCP and DTR
- Docker volumes

Backup and restore UCP and DTR

The playbooks provided in this solution implement the backup and restore procedures as they are described in the Docker documentation at <https://docs.docker.com/enterprise/backup/>. The solution follows the recommendations in the Docker best practices document at <https://success.docker.com/article/backup-restore-best-practices>.

Note

It is important that you make copies of the backed up data and that you store the copies in a separate physical location. You must also recognize that the backed up data contains sensitive information such as private keys and so it is important to restrict access to the generated files. However, the playbooks do not backup the sensitive information in your `group_vars/vault` file so you should make sure to keep track of the credentials for the UCP Administrator.

Warning

The restore procedures do not restore swarm data. You should follow infrastructure as code (IaC) guidelines and maintain your service, stack and network definitions using source code or configuration management tools. You must also ensure that you safely manage the credentials of your administration accounts, as existing UCP Client bundles will not work when you restore UCP on a new swarm.



Backup UCP and DTR

The playbooks support backing up the swarm, UCP, DTR metadata and DTR images.

Backup configuration variables

Table 24 shows the variables related to backing up UCP and DTR. All these variables are defined in the file **group_vars/backups**. All the data that is backed up is streamed over an SSH connection to the backup server. Currently, the playbooks only support the use of the Ansible box as the backup server.

Table 24. Backup variables

Variable	File	Description
backup_server	group_vars/backups	Currently, the playbooks only support the use of the Ansible box as the backup server.
backup_dest	group_vars/backups	This variable should point to an existing folder on your Ansible box where the <code>root</code> user has write access. All the backups will be stored in this folder. For example, <code>/root/backups</code>
backup_passphrase	group_vars/vault	This variable is used to encrypt the tar file with a passphrase that must be at least 12 characters long.
#swarm_offline_backup	group_vars/backups	This variable is commented out by default. More information on this variable is provided below.

Backing up the swarm

When you back up the swarm, your services and stack definitions are backed up together with the networks definitions. However, Docker volumes or their contents will not be backed up. (If Docker volumes are defined in stacks, they will be re-created when you restore the stacks, but their content will be lost). You can back up the swarm using the playbook named **backup-swarm.yml** which is located in the **playbooks** folder on your Ansible server. The playbook is invoked as follows:

```
# ansible-playbook -i vm_hosts playbooks/backup-swarm.yml
```

This playbook creates two archives in the folder specified by the variable **backup_dest** in **group_vars/backups**. By default, the file is named using the following pattern:

```
<backup_dest>/backup-swarm-<vmname>_<timestamp>.tgz
<backup_dest>/backup-swarm-<vmname>_<timestamp>.vars.tgz
```

<vmname> is the name of the host (in the inventory) that was used to take the backup, and **<timestamp>** is the time at which the backup was taken. The file with the extension **.vars.tgz** contains information regarding the system that was backed up.

You can override the generated file name by defining the variable **backup_name** on the command line when running the playbook. In the example below:

```
# ansible-playbook -i vm_hosts playbooks/backup-swarm.yml -e backup_name=my-swarm-backup
```

The generated files won't have **<vmname>** or **<timestamp>** appended:

```
<backup_dest>/my-swarm-backup.tgz
<backup_dest>/my-swarm-backup.vars.tgz
```

Warning

Online versus offline backups: By default, the playbook performs online backups. You can take offline backups by setting the variable **swarm_backup_offline** to **"true"**. The playbook will then stop the Docker daemon on the machine used to take the backup (a manager or UCP node). Before it does so, the playbook will verify that enough managers are running in the cluster to maintain the quorum. If this is not the case, the playbook will exit with an error. For more information, see the Docker documentation at https://docs.docker.com/engine/swarm/admin_guide/#recover-from-disasterv



Backing up the Universal Control Plane (UCP)

When you backup UCP, you save the data/metadata outlined in Table 25.

Table 25. UCP data backed up

Data	Description
Configurations	The UCP cluster configurations, as shown by <code>docker config ls</code> , including Docker EE license and swarm and client CAs
Access control	Permissions for team access to swarm resources, including collections, grants, and roles
Certificates and keys	The certificates, public keys, and private keys that are used for authentication and mutual TLS communication
Metrics data	Monitoring data gathered by UCP
Organizations	Your users, teams, and orgs
Volumes	All <u>UCP named volumes</u> , which include all UCP component certs and data

To make a backup of UCP, use `playbook/backup-ucp.yml` as follows:

```
# ansible-playbook -i vm-host playbooks/backup-ucp.yml
```

This playbook creates two archives in the folder specified by the variable `backup_dest` in `group_vars/backups`. By default, the files are named using the following pattern:

```
<backup_dest>/backup-ucp-<ucpid>_<vmname>_<timestamp>.tgz
<backup_dest>/backup-ucp-<ucpid>_<vmname>_<timestamp>.vars.tgz
```

`<ucpid>` is the ID of the UCP instance, `<vmname>` is the name of the host (in the inventory) that was used to take the backup, and `<timestamp>` is the time at which the backup was taken. The file with the extension `.vars.tgz` contains information regarding the system which was backed up.

You can override the generated file name by defining the variable `backup_name` on the command line when running the playbook. In the example below:

```
# ansible-playbook -i vm-hosts playbooks/backup-ucp.yml -e backup_name=my-ucp-backup
```

The generated files won't have `<vmname>` or `<timestamp>` appended:

```
<backup_dest>/my-ucp-backup.tgz
<backup_dest>/my-ucp-backup.vars.tgz
```

Warning

To create a consistent backup, the backup command **temporarily stops the UCP containers running on the node where the backup is being performed**. User resources, such as services, containers, and stacks are not affected by this operation and will continue to operate as expected. Any long-lasting `docker exec`, `docker logs`, `docker events`, or `docker attach` operations on the affected manager node will be disconnected.

For more information on UCP backup, see the Docker documentation at <https://docs.docker.com/datacenter/ucp/3.0/guides/admin/backups-and-disaster-recovery/>



Backing up the Docker Trusted Registry (DTR)

When you backup DTR, you save the data/metadata outlined in Table 26.

Table 26. DTR data backed up

Data	Backed up?	Description
Configurations	yes	DTR settings
Repository metadata	yes	Metadata like image architecture and size
Access control to repos and images	yes	Data about who has access to which images
Notary data	yes	Signatures and digests for images that are signed
Scan results	yes	Information about vulnerabilities in your images
Certificates and keys	yes	TLS certificates and keys used by DTR
Image content	no	Needs to be backed up separately, depends on DTR configuration
Users, orgs, teams	no	Create a UCP backup to backup this data
Vulnerability database	no	Can be re-downloaded after a restore

To make a backup of DTR metadata, use `playbook/backup_dtr_metadata.yml`

```
# ansible-playbook -i vm_host playbooks/backup_dtr_metadata.yml
```

This playbook creates two archives in the folder specified by the variable `backup_dest` in `group_vars/backups`. By default, the file is named using the following pattern:

```
<backup_dest>/backup_dtr_meta-<replica_id>-<vmname>-<timestamp>.tgz
<backup_dest>/backup_dtr_meta-<replica_id>-<vmname>-<timestamp>.vars.tgz
```

`<replica_id>` is the ID of the DTR replica that was backed up, `<vmname>` is the name of the host (in the inventory) that was used to take the backup, and `<timestamp>` is the time at which the backup was taken. The file with the extension `.vars.tgz` contains information regarding the system that was backed up.

You can override the generated file name by defining the variable `backup_name` on the command line when running the playbook. In the example below:

```
# ansible-playbook -i vm_hosts playbooks/backup_dtr_metadata.yml -e backup_name=my_dtr_metadata_backup
```

The generated files won't have `<vmname>` or `<timestamp>` appended:

```
<backup_dest>/my_dtr_metadata_backup.tgz
<backup_dest>/my_dtr_metadata_backup.vars.tgz
```

For more information on DTR backups, see the Docker documentation at <https://docs.docker.com/datacenter/dtr/2.5/guides/admin/backups-and-disaster-recovery/>

Backing up DTR data (images)

To make a backup of the images that are inventoried in DTR and stored on the NFS server, use `playbooks/backup_dtr_images.yml`

```
# ansible-playbook -i vm_host playbooks/backup_dtr_images.yml
```

This playbook creates two archives in the folder specified by the variable `backup_dest` in `group_vars/backups`. By default, the files are named using the following pattern:

```
<backup_dest>/backup_dtr_data-<replica_id>-<vmname>-<timestamp>.tgz
<backup_dest>/backup_dtr_data-<replica_id>-<vmname>-<timestamp>.vars.tgz
```



<replica_id> is the ID of the DTR replica that was backed up, <vmname> is the name of the host (in the inventory) that was used to take the backup, and <timestamp> is the time at which the backup was taken.

You can override the generated file names by defining the variable **backup_name** on the command line when running the playbook, as shown in the example below:

```
# ansible-playbook -i vm_hosts playbooks/backup_dtr_images.yml -e backup_name=my_dtr_data_backup
```

The generated files won't have <vmname> or <timestamp> appended:

```
<backup_dest>/my_dtr_data_backup.tgz
<backup_dest>/my_dtr_data_backup.vars.tgz
```

For more information on DTR backups, see the Docker documentation at <https://docs.docker.com/datacenter/dtr/2.5/guides/admin/backups-and-disaster-recovery/>

Backing up other metadata, including passwords

The backup playbooks do not backup the sensitive data in your **group_vars/vault** file. The information stored in the **.vars.tgz** files includes backups of the following files:

- **vm_hosts**, a copy of the **vm_hosts** file at the time the backup was taken
- **vars**, a copy of your **group_vars/vars** file at the time the backup was taken
- **meta.yml**, a generated file containing information pertaining to the backup

The **meta.yml** file contains the following information:

```
backup_node="<node that took the backup>"
replica_id="<ID of DTR replica if DTR backup>"
backup_source=""
ucp_version="<UCP version if UCP backup>"
dtr_version="<DTR version of DTR backup>"
```

Backup Utility

The script **backup.sh** can be used to take a backup of the swarm, UCP, DTR metadata and the DTR images in one go. You can pass this script an argument (tag) that will be used to prefix the backup filenames, thereby overriding the default naming. Table 27 shows the file names produced by **backup.sh** based on the argument passed in the command line.

Table 27. Backup utility

Example	Command line	Generated filenames
Default	<code>./backup.sh</code>	<code>backup_swarm_<vmname>_<timestamp>.tgz</code> , <code>backup_ucp_<ucpid>_<vmname>_<timestamp>.tgz</code> , <code>backup_dtr_meta_<replica_id>_<vmname>_<timestamp>.tgz</code> , <code>backup_dtr_data_<replica_id>_<vmname>_<timestamp>.tgz</code> and the corresponding .vars.tgz files
Custom	<code>./backup.sh my_backup</code>	<code>my_backup_swarm.tgz</code> , <code>my_backup_ucp.tgz</code> , <code>my_backup_dtr_meta.tgz</code> , <code>my_backup_dtr_data.tgz</code> , and the corresponding .vars.tgz files
Date	<code>./backup.sh \${date '+%Y_%m_%d_%H%M%S'}</code>	<code><date>_swarm.tgz</code> , <code><date>_ucp.tgz</code> , <code><date>_dtr_meta.tgz</code> , <code><date>_dtr_data.tgz</code> , and the corresponding .vars.tgz files

In addition, the **backup.sh** script accepts an optional switch that will let you specify the location of the password file that will be passed to the **ansible-playbook** commands in the script. This is required if you have encrypted the **group_vars/vault** file. The general syntax for this script is as follows:

```
./backup.sh [ -v <Vault Password File> ] [ tag ]
```



Related playbooks

- `playbooks/backup_swarm.yml` is used to back up the swarm data
- `playbooks/backup_ucp.yml` is used to back up UCP
- `playbooks/backup_dtr_meta.yml` is used to back up DTR metadata
- `playbooks/backup_dtr_images.yml` is used to back up DTR images

Restoring your cluster after a disaster

The playbooks address a disaster recovery scenario where you have lost your entire cluster and all the VMs. Other scenarios and how to handle them are described in the Docker documentation including the following scenarios:

- You have lost one UCP instance but your cluster still has the quorum. The easiest way is to recreate the missing UCP instance from scratch.
- You have lost the quorum in your UCP cluster but there is still one UCP instance running.
- You have lost one instance of DTR but still have a quorum of replicas. The easiest way is to recreate the missing DTR instance from scratch.
- You have lost the quorum of your DTR cluster but still have one DTR instance running.

Before you restore

Step 1. Retrieve the backup files using your chosen backup solution and save them to a folder on your Ansible server. If you have used timestamps in the naming of your backup files, you can use them to determine the chronological order. If you used the `backup.sh` script specifying a date prefix, you can use that to identify the matching set of backup files. You should choose the files in the following reverse chronological order, from the most recent to the oldest file. Make sure you restore both the `*.tgz` and the `*.vars.tgz` files.

1. DTR images backup
2. DTR metadata backup
3. UCP backup
4. Swarm backup

In this example, we will assume a set of backup files stored in `/root/restore` that were created specifying a date prefix. These will have names like `2018_04_17_151734_swarm.tgz`, `2018_04_17_151734_ucp.tgz`, etc and the corresponding `.vars.tgz` files.

Step 2: Retrieve the DTR replica ID, the DTR version and the UCP version

To retrieve the ID of the replica that was backed up, as well as the version of DTR, you need to extract the data from the `.vars.tgz` file associated with the archive of the DTR metadata. You can retrieve this as follows:

```
[root@hpe2-ansible ops]# tar -0xf /root/restore/2018_04_17_151734_dtr_meta.vars.tgz meta.yml
backup_node="hpe-dtr01"
replica_id="ad5204e8a4d0"
backup_source=""
ucp_version=""
dtr_version="2.4.3"
```

```
[root@hpe2-ansible ops]# tar -0xf /root/restore/2018_04_17_151734_ucp.vars.tgz meta.yml
backup_node="hpe-ucp01"
replica_id=""
backup_source=""
ucp_version="3.0.4"
dtr_version=""
```

Take note of the replica ID (`ad5204e8a4d0`), the version of DTR (`2.5.3`) and the version of UCP (`3.0.4`).

Step 3: Populate the `group_vars/backups` file



```

backup_swarm: "/root/restore/2018_04_17_151734_swarm.tgz"
backup_ucp: "/root/restore/2018_04_17_151734_ucp.tgz"
backup_dtr_meta: "/root/restore/2018_04_17_151734_dtr_meta.tgz"
backup_dtr_data: "/root/restore/2018_04_17_151734_dtr_data.tgz"
backup_dtr_id: "ad5204e8a4d0"
backup_dest: "/root/backups"
backup_server: <IP of your ansible box>

```

You should populate your `group_vars/backups` file as above, with the `backup_dtr_id` variable containing the value you retrieved in the preceding step as `replica_id="ad5204e8a4d0"`.

Step 4: Verify that your `group_vars/vars` file specifies the correct versions of DTR and UCP.

The playbooks use the versions of UCP and DTR as specified in your `group_vars/vars` file to restore your backups. You must ensure that the versions specified in your current `group_vars/vars` file correspond to the versions in the backups as determined above.

```

[root@hpe2-ansible ops]# cat group_vars/vars | grep dtr_version
dtr_version: '2.5.3'

```

```

[root@hpe2-ansible ops]# cat group_vars/vars | grep ucp_version
ucp_version: '3.0.4'

```

Step 5: Restore UCP admin credentials if required

You must ensure that the UCP admin credentials in your current `group_vars/vars` file are those that were in effect when you generated the backup files. If they have changed since then, you must restore the original credentials for the duration of the restore procedure.

Step 6: Restore your inventory (`vm_hosts`)

Your inventory must reflect the environment that was present when the backup files were created. You can find a copy of the inventory as it was when the backup was taken in the `*.vars.tgz` files.

Restore UCP and DTR

Warning

This procedure is aimed at restoring a cluster after a disaster. It assumes you have lost all the VMs in your cluster and want to redeploy using data that you backed up earlier. The solution follows Docker best practice, which means the swarm artifacts are not restored. You will need to restore your Docker volumes and your applications (stacks and services) when this procedure is complete.

1. Ensure that you have completed all the preliminary steps as outlined in the section [Before you restore](#).
2. Run the restore playbook


```
ansible-playbook -i vm_hosts restore.yml
```
3. If you are using the image scanning functionality in DTR, you will need to re-download the vulnerability database. For more information, see the Docker documentation [here](#).

You are now ready to restore your Docker volumes and your applications.

Restore DTR metadata and DTR images

Note

This procedure restores DTR metadata and images and assumes you have lost all the DTR VMs in your cluster. It will redeploy using the DTR data that you backed up earlier and will also restore the images if the folder exported by the NFS VM is empty.



1. Ensure that you have completed all the preliminary steps as outlined in the section [Before you restore](#). In this scenario, you need the archives for the DTR metadata and the DTR images.
2. Ensure that all the DTR VMs listed in your inventory are destroyed, using the vSphere Web Client to delete them if required. If you want to restore the DTR images you should also delete the NFS VM.
3. Remove the DTR nodes from the swarm by running the `docker node rm <DTR node>` command on a UCP node for each DTR node in your cluster. The following example shows the sequence of commands to use to remove the DTR nodes:

```
[root@hpe2-ansible ops]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY
aiz...	* hpe-ucp02.cloudra.local	Ready	Active
gvf...	hpe-dtr01.cloudra.local	Down	Active
ir4...	hpe-ucp03.cloudra.local	Ready	Active
mwf...	hpe-dtr02.cloudra.local	Down	Active
oqy...	hpe-ucp01.cloudra.local	Ready	Active
xqe...	hpe-worker01.cloudra.local	Ready	Active
zdu...	hpe-dtr03.cloudra.local	Down	Active

```
[root@hpe2-ansible ops]# docker node rm hpe-dtr01.cloudra.local
hpe-dtr01.cloudra.local
[root@hpe2-ansible ops]# docker node rm hpe-dtr02.cloudra.local
hpe-dtr02.cloudra.local
[root@hpe2-ansible ops]# docker node rm hpe-dtr03.cloudra.local
hpe-dtr03.cloudra.local
```

```
[root@hpe2-ansible ops]# docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY
aiz...	hpe-ucp02.cloudra.local	Ready	Active
ir4...	hpe-ucp03.cloudra.local	Ready	Active
oqy...	* hpe-ucp01.cloudra.local	Ready	Active
xqe...	hpe-worker01.cloudra.local	Ready	Active

4. Run the restore script:

```
./restore_dtr.sh
```

5. If you are using the image scanning functionality in DTR, you will need to re-download the vulnerability database. For more information, see the Docker documentation [here](#).

Related playbooks

- `playbooks/restore_swarm.yml` is used to restore the swarm data
- `playbooks/restore_dtr_meta.yml` is used to restore DTR metadata
- `playbooks/restore_dtr_images.yml` is used to restore DTR images

HPE SimpliVity backups

HPE SimpliVity functionality can be used to backup and restore Docker persistent volumes.

Backup and restore Docker persistent volumes

In order to restore a Docker volume, you need to restore a special VM that has been deployed for the sole purpose of backing up Docker volumes. There is one such VM for each datastore defined in the `datastores` array in the `group_vars/vars` file. By default, a single datastore is specified in the playbooks:

```
datastores: [ 'Docker_HPE' ]
```



Note

The use of a single datastore is recommended. If you have configured multiple datastores, you need to understand and keep track of how your Docker volumes are distributed across the datastores.

The name of the special VM follows the pattern `<prefix>-in-dockervols-<Datastore>` where

- `<prefix>` is the value of the variable `dummy_vm_prefix` from the file `group_vars/vars`
- `<Datastore>` is the name of the datastore

For example, based on the default values in the scripts, the VM name would be `hpe-VM-in-dockervols-Docker_HPE`.

Create a Docker volume

A single Docker volume will have been created for Prometheus using the vSphere driver as part of the initial deployment. To see this volume, use the `docker volume ls` command and limit the results to those volumes created using the vSphere driver.

```
# docker volume ls | grep vsphere
vsphere:latest      prom_hpe-db-data@Docker_HPE
#
```

To create a Docker volume named `test_01`, you can use the `docker volume create` command specifying the vSphere driver:

```
# docker volume create -d vsphere test_01
test_01
```

You can check that the volume exists using the `docker volume ls` command:

```
# docker volume ls | grep vsphere
vsphere:latest      prom_hpe-db-data@Docker_HPE
vsphere:latest      test_01@Docker_HPE
```

You can attach a container to the volume and then add data to it by creating a text file with some arbitrary content:

```
# docker run -it --rm -v test_01:/tmp alpine sh -c "echo some test data here > /tmp/foo.txt"
```

If this is the first time you have used the alpine image, you may see additional output relating to download of image layers:

```
Unable to find image 'alpine:latest' locally
latest: Pulling from library/alpine
88286f41530e: Already exists
Digest: sha256:f006ecbb824d87947d0b51ab8488634bf69fe4094959d935c0c103f4820a417d
Status: Downloaded newer image for alpine:latest
```

The container will exit once the shell command has run and any unnamed volumes will be removed. However, the named volume `test_01:/tmp` will persist. To check that the data is still available after the container exited, spin up a new container and try to retrieve the data:

```
# docker run -it --rm -v test_01:/tmp alpine sh -c "cat /tmp/foo.txt"
some test data here
```

Automated backup

By default, the special VM and any Docker volume in the `dockvols` folder are backed up every hour. This is controlled by the following settings in the `group_vars/vars` file.

```
backup_policies:
- name: 'hpe-gold'
  days: 'All'
  start_time: '00:00'
```




```
frequency: '60'
retention: '43200'
dummy_vm_prefix: 'hpe-VM'
docker_volumes_policy: 'hpe-gold'
```

The backup policy **hpe-gold** is assigned to the special VM that is used to back up the Docker volumes. This policy specifies that a backup is taken every hour (**frequency: '60'** means sixty minutes) while the backup is retained for one month (**retention: '43200'** means 43200 minutes or thirty days).

Manual backup

Rather than waiting for an automated backup to take place, you can create a backup immediately. Right-click on the special VM, in this case **hpe-VM-in-dockervols-Docker_HPE**, select **All HPE SimpliVity Actions** and choose **Backup Virtual Machine** as shown in Figure 39.




Figure 39. Backup virtual machine

You can specify a backup name, in this case `manual_backup_test_01`, as shown in Figure 40.




Figure 40. Backup virtual machine details

Restore

Right-click on the special VM, in this case `hpe-VM-in-dockervols-Docker_HPE`. On the Configure tab, select HPE SimpliVity Search Backups as shown in Figure 41.

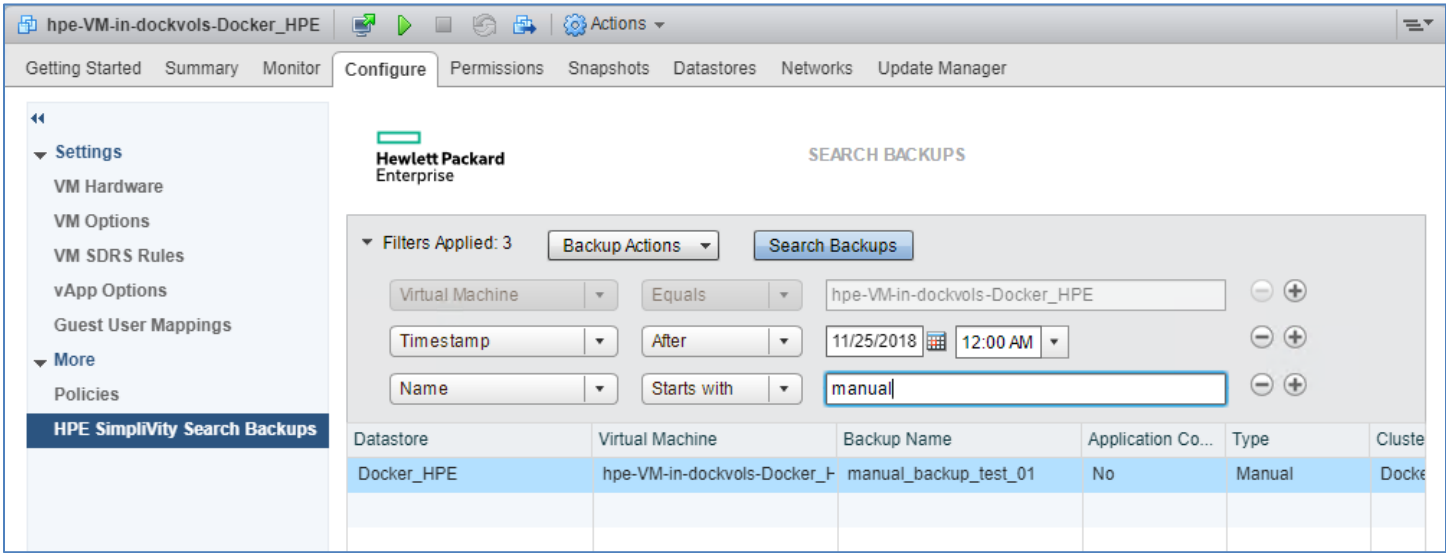


Figure 41. Search backups

You can narrow the search based on the time of the backup. If you are restoring from an automatic backup, the name will be the timestamp of the backup. If you are restoring from a manual backup, the name will be the one you specified earlier when creating the backup, in this case `manual_backup_test_01`.



Right-click on the backup you wish to restore, as shown in Figure 42, and select **Restore Virtual Machine**.

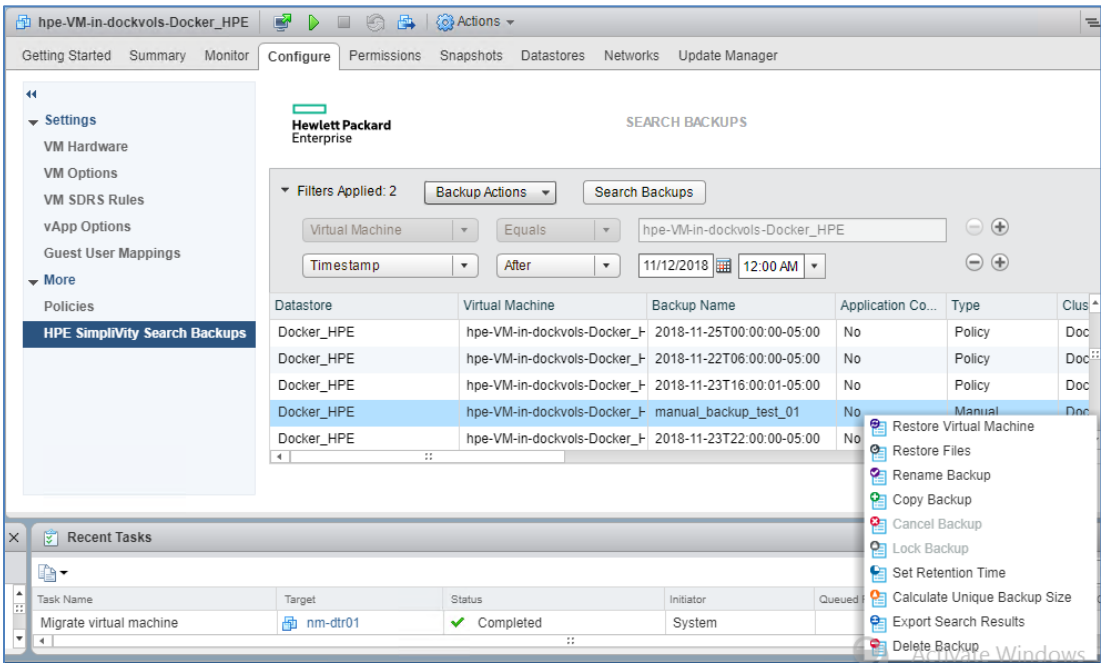


Figure 42. Restore virtual machine

In the details screen, shown in Figure 43, you can choose a name for the new virtual machine and specify the datstore.




Figure 43. Restore virtual machine details



The name of the new virtual machine will default to a combination of the special VM name and a timestamp, in this instance `hpe-VM-in-dockervols-Docker_HPE-2018-11-26-20h47m01s`. The datastore should be the one specified in the datastores array from the `group_vars/vars` file. Click OK to restore the virtual machine.

Once the virtual machine has been restored, navigate to the datastore and locate the new VM in the file browser, as shown in Figure 44.

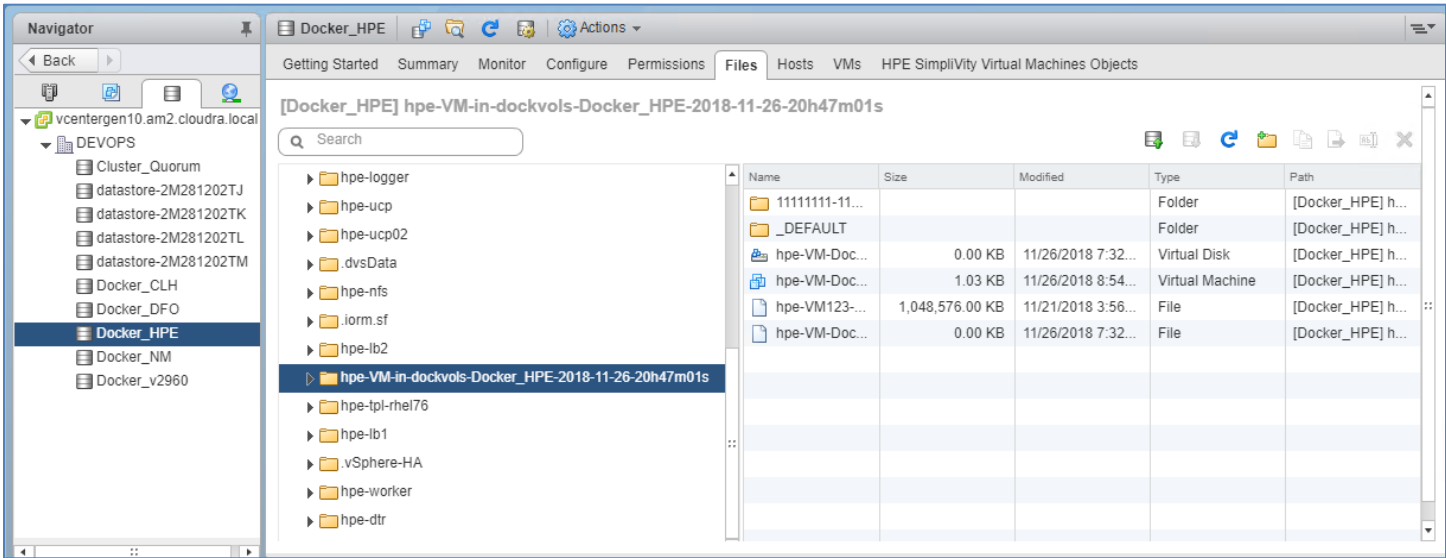


Figure 44. Browse for restored virtual machine

Navigate to the folder named `11111111-1111-1111-1111-...` as shown in Figure 45. You will see files with names based on the Docker volume name that you used at the start, in this instance `test_01.vmdk` and `test_01-478...f1f.vmf`




Figure 45. Locate vmdk and vmf files



You need to move these two files to the dockvols sub-directory named 11111111-1111-1111-1111-... in the same datastore. Right click on the .vmdk file and choose **Move to...** as shown in Figure 46.

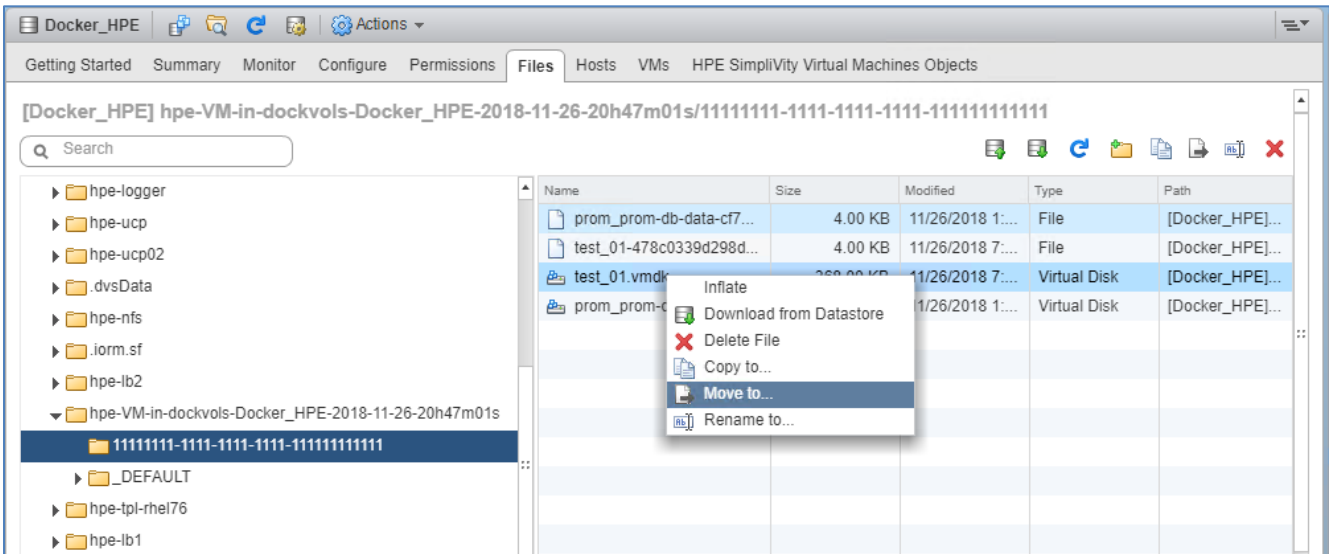


Figure 46. Move files

Set the destination folder to the dockvols sub-directory named 11111111-1111-1111-1111-... as shown in Figure 47.




Figure 47. Move to destination



It is only necessary to move the `.vmdk` file as the `.vmfd` file will automatically follow. The `dockvols` sub-directory named `1111111-1111-1111-1111-...` should now contain both files as shown in Figure 48.

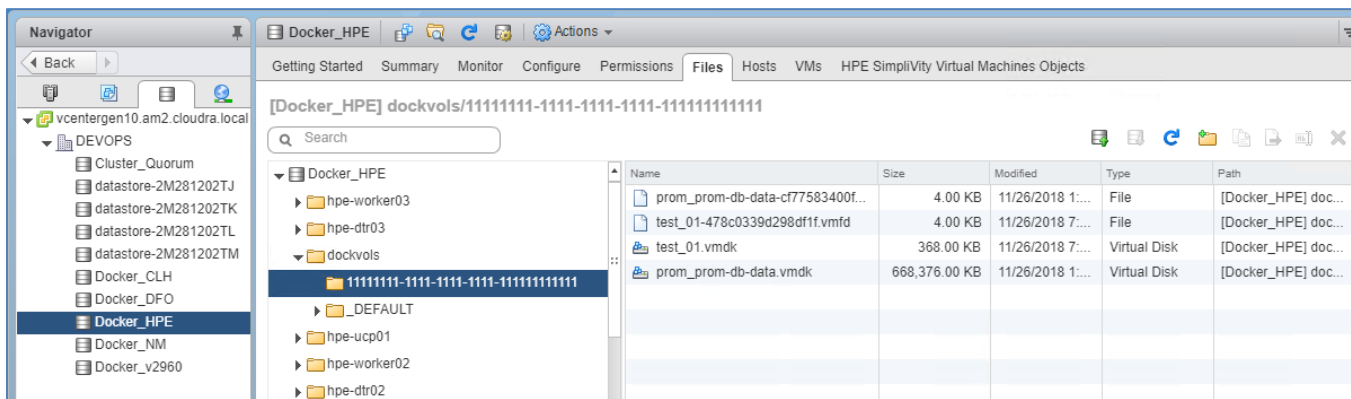


Figure 48. Files moved to destination

Test the restore

You can check that the volume `test_01` has been restored by using the `docker volume ls` command again.

```
# docker volume ls | grep vsphere
vsphere:latest      prom_hpe-db-data@Docker_HPE
vsphere:latest      test_01@Docker_HPE
```

You can verify that the volume contains the correct data by spinning up a container and running a shell command:

```
# docker run -it --rm -v test_01:/tmp alpine sh -c "cat /tmp/foo.txt"
```

some test data here

The data you entered in the text file, before performing the backup and deleting the volume, is available once again after restoring the volume.



Solution lifecycle management

Lifecycle management with respect to this solution refers to the maintenance and management of software and hardware of various components that make up the solution stack. Lifecycle management is required to keep the solution up-to-date and ensure that the latest versions of the software are running to provide optimal performance, security and to fix any existing defects within the product.

In this section, we will cover life cycle management of the different components that are used in this solution. The lifecycle of the following stacks need to be maintained and managed:

- Monitoring Tools (Splunk or Prometheus and Grafana)
- Docker Enterprise Edition Environment
- Virtual Machine Operating systems
- HPE SimpliVity environment

The general practice and recommendation is to follow a bottom-up approach for updating all components of the environment and making sure the dependencies are met. In our solution, we would start with HPE SimpliVity and end with the monitoring environment. If all components are not being updated at the same time, the same approach can be followed – updating only the components that require updates while adhering to the interdependencies of each component that is being updated.

HPE SimpliVity environment

The HPE SimpliVity environment is made up of proprietary HPE SimpliVity software, VMware software and HPE firmware. There are interdependencies between the various components that need to be accounted for and are provided in the table below. The components in Table 28 are part of the HPE SimpliVity environment that require lifecycle management.

In general, ensure that the software bits for the Arbiter and vSphere extension corresponding to an OmniStack release are used.

Table 28. HPE SimpliVity components

Order	Component	Dependency (compatibility)	Download/Documentation
1	HPE SimpliVity Arbiter	1. HPE OmniStack	SimpliVity OmniStack for vSphere Upgrade Guide
2	HPE SimpliVity VMware Plug-in	1. HPE SimpliVity Arbiter 2. HPE OmniStack	Download software bits from HPE's support website. http://www.hpe.com/support
3	HPE Omnistack	1. HPE SimpliVity VMware Plug-in 2. HPE SimpliVity Arbiter	

VMware Components

The solution in this deployment guide is built on VMware vSphere and leverages VMware ESXi and vCenter. For more information on upgrading vSphere, see the VMware documentation, Introduction to vSphere Upgrade, at <https://docs.vmware.com/en/VMware-vSphere/6.5/com.vmware.vsphere.upgrade.doc/GUID-EB29D42E-7174-467C-AB40-DB37236FEAF5.html>.

The VMware ESXi and vCenter versions must be compatible with each other and with the HPE OmniStack version that is running on the HPE SimpliVity systems.

Table 29. VMware components

Order	Component	Dependency (compatibility)	Download/Documentation
1	VMware vCenter	1. HPE OmniStack 2. VMware ESXi	VMware Upgrade for SimpliVity
2	VMware ESXi	1. HPE OmniStack 2. VMware vCenter	



HPE server software

HPE SimpliVity servers are based on HPE server platforms and require a compatible firmware version to function with HPE OmniStack Software, as shown in Table 30.

Table 30. HPE server components

Order	Component	Dependency (compatibility)	Download/Documentation
1	HPE Firmware	1. HPE OmniStack Software	Firmware Upgrade for SimpliVity

vSphere Docker Volume Service Plug-in

vSphere Docker Volume service plug-in is part of an open source project by VMware that enables running stateful containers by providing persistent Docker volumes leveraging existing storage technology from VMware. There are two parts to the plug-in, namely, client software and server software (see Table 31). Every version of the plug-in that is released includes both pieces of software and it is imperative that the version number installed on the client side and server side are the same.

When updating the Docker Volume service plug-in, ensure the ESXi version you are running is supported and that the client software is compatible with the operating system.

Table 31. vSphere Docker Volume service components

Order	Component	Dependency (compatibility)	Download/Documentation
1.	Server Software	1. VMware ESXi 2. Docker EE	vSphere Docker Volume Service on GitHub
2.	Client Software	1. VM Operating System 2. Docker EE	

Red Hat Enterprise Linux operating system

This solution is built using Red Hat Enterprise Linux (see Table 32) as the base operating system. When upgrading the operating system on the VMs, first verify that the OS version is compatible with Docker EE by looking at the Docker OS compatibility matrix.

Table 32. Operating system

Order	Component	Dependency (compatibility)	Download/Documentation
1.	Red Hat Enterprise Linux	1. Docker EE 2. vDVS client software plugin	RHEL



Docker EE Environment

Each release of Docker Enterprise Edition contains three technology components – UCP, DTR and the Docker Daemon or Engine. It is imperative that the components belonging to the same version are deployed or upgraded together – see Table 33.

A banner will be displayed on the UI, as shown in Figure 49, when an update is available for UCP or DTR. You can start the upgrade process by clicking on the banner.




Figure 49. Docker update notification

Table 33. Docker EE components

Order	Component	Dependency (compatibility)	Download/Documentation
1.	Docker Daemon/Engine	1. VM Operating System	Docker Lifecycle Maintenance
2.	Universal Control Plane	2. vDVS plugin	Docker Compatibility Matrix
3.	Docker Trusted Registry	3. Prometheus and Grafana	

Monitoring Tools

To learn more about upgrading Splunk, see the relevant documentation at [How to upgrade Splunk Enterprise](#).

The Sysdig agent runs as a container and the latest version is pulled from the Docker hub on first installation. Re-run the `install_sysdig.yml` playbook to update to the newest version if required.

Prometheus and Grafana monitoring tools (see Table 34) run as containers within the Docker environment. Newer versions of these tools can be deployed by pulling the Docker images from Docker Hub. Verify that the version of Prometheus that is being used is compatible with the version of Docker EE.

Table 34. Monitoring tools: Prometheus and Grafana

Order	Component	Dependency (compatibility)	Download/Documentation
1.	Prometheus	1. Grafana 2. Docker EE	1. Prometheus Images on Docker Hub
2.	Grafana	1. Prometheus 2. Docker EE	2. Upgrading Grafana

Summary

This document has described how to architect and deploy a Docker CaaS platform on HPE SimpliVity, using Ansible playbooks to quickly install and deploy a production-ready container environment. This deployment includes a highly available container cluster with backup services and persistent data support. This solution is ideal for customers looking to run containers on VMs to take advantage of the resource efficient usage of virtual machines for Docker containers, and having the ability to run legacy and new container applications side-by-side. Customers deploying Docker containers on a large scale, on Linux and Microsoft Windows, should consider HPE SimpliVity as the deployment infrastructure.



Appendix A: Software Licenses

Licenses are required for the following software components:

- VMware
- Red Hat Linux
- Docker EE
- Splunk (optional software)
- Sysdig (optional software)

Appendix B: Using customer supplied certificates for UCP and DTR

Table 35 lists the variables used when configuring customer supplied certificates for UCP and DTR.

Table 35. Customer certs variables

Variable	File	Description
ucp_certs_dir	group_vars/vars	<p>If ucp_certs_dir is not defined, UCP is installed with self-signed certificates and DTR is installed with the --ucp-insecure-tls switch</p> <p>If ucp_certs_dir is defined, this is a folder on the Ansible machine that must contain 3 files:</p> <p>ca.pem, the root CA certificate in PEM format</p> <p>cert.pem, the server certificate optionally followed by intermediate CAs</p> <p>key.pem, the private key that comes with the cert.pem certificates</p>
dtr_certs_dir	group_vars/vars	<p>If dtr_certs_dir is not defined, DTR is installed with self-signed certificates</p> <p>If dtr_certs_dir is defined, this is a folder on the Ansible machine that must contain 3 files:</p> <p>ca.pem, the root CA certificate in PEM format</p> <p>cert.pem, the server certificate optionally followed by intermediate CAs</p> <p>key.pem, the private key that comes with the cert.pem certificates</p>

Note

The installation will fail if the **ca.pem**, **cert.pem** and **key.pem** files cannot be found in the folders designated by **dtr_certs_dir** and **ucp_certs_dir** or if they don't constitute valid certificates.

The certificates should specify the names of the FQDNs of the load balancer and the FQDNs of the VMs themselves. This applies to both the UCP server certificate and the DTR server certificate.

Generating and testing certificates

In the example described here we have a root CA named **Example root CA** and an intermediate CA named **Intermediate CA valid 3 years**. The intermediate CA signs the server certificates for UCP and DTR.

Below is the start of the output displayed by running the **openssl x509** utility against the **ca.pem** file (the root CA certificate).

```
[root@ansible ucp_certs]# openssl x509 -text -noout -in ca.pem | head -14
Certificate:
    Data:
        Version: 3 (0x2)
        Serial Number:
            0d:07:ca:ea:00:37:77:6e:25:e0:18:3e:0e:db:80:0f:11:cb:1b:3f
        Signature Algorithm: sha256WithRSAEncryption
```



```

Issuer: CN=Example Root CA
Validity
  Not Before: Apr 24 20:12:01 2018 GMT
  Not After : Apr 21 20:12:30 2028 GMT
Subject: CN=Example Root CA
Subject Public Key Info:
  Public Key Algorithm: rsaEncryption
  Public-Key: [4096 bit]

```

Here is an excerpt from the example `ca.pem` file:

```

-----BEGIN CERTIFICATE-----
MIIFJTCCAww2gAwIBAgIUDDQfK6gA3d2414Bg+DtuADxHLGz8wDQYJKoZIhvcNAQEL
BQAwGjEYMBYGA1UEAxMPRXhhbXBsZSBSb290IENBMB4XDTE4MDQyNDIwMTIwMVox
...
...
uXzYbCtU6Jt9B3fayAeWswQv+jQSuUa3reOM1x838iIZWDx93f4yLJWLJ17xsY
btvKBmqKDCsAqsQLFLNj/JyYq4e9a6Xxcyn9FXNpzuEsFjfNGHn+csY+w3f987T
MNviy376xZbyAc1CV5kgmnZzjU5bDkgT8Q==
-----END CERTIFICATE-----

```

The `cert.pem` file should contain the server certificate itself, followed by your intermediate CA's certificate. The following example shows how to extract the intermediate CA certificate from `cert.pem` and to save it to a file named `intca.pem`. Using the `openssl x509` utility, you can display the content of the `intca.pem` file in human readable form. This certificate was signed by the example CA above (Issuer = 'Example Root CA').

```

[root@ansible ucp-certs]# openssl x509 -text -noout -in intca.pem | head -14
Certificate:
  Data:
    Version: 3 (0x2)
    Serial Number:
      6b:1e:0c:86:20:cf:f0:88:d2:52:0d:5d:b9:56:fa:91:87:a0:49:18
  Signature Algorithm: sha256WithRSAEncryption
  Issuer: CN=Example Root CA
  Validity
    Not Before: Apr 24 20:12:09 2018 GMT
    Not After : Apr 23 20:12:39 2021 GMT
  Subject: CN=Intermediate CA valid 3 years
  Subject Public Key Info:
    Public Key Algorithm: rsaEncryption
    Public-Key: [4096 bit]

```

Here is an excerpt from the `intca.pem` file showing the example Intermediate CA certificate:

```

-----BEGIN CERTIFICATE-----
MIIFCjCCA1qgAwIBAgIUax4MhiDP8IjSUg1duVb6kYegSRgwDQYJKoZIhvcNAQEL
BQAwGjEYMBYGA1UEAxMPRXhhbXBsZSBSb290IENBMB4XDTE4MDQyNDIwMTIwMVox
...
...
o2tL5nwR7R0iAr/kk9MIRzWzLNbc4cYth7jEjSpU9dBqsXgsTozzWlwqI9ybZwvL
Ni1JnZandVlyQdo0aB2M/1DNFfKvwW3JeaKvDA9j95n/BWFTjoZ+Y0z9pYit6T7
1GCGu3be
-----END CERTIFICATE-----

```

The `openssl x509` utility will only decrypt the first certificate found in `cert.pem`, so you don't need to extract the server certificate from `cert.pem`. In this example, the server certificate is signed by the intermediate CA above. Note the **Subject Alternate Names**: `hpe-`



ucp.cloudra.local is the FQDN of the UCP load balancer, and the other names are those of the UCP instances (hpe-ucp01.cloudra.local, hpe-ucp02.cloudra.local, hpe-ucp03.cloudra.local).

```
[root@ansible ucp-certs]# openssl x509 -text -noout -in server.pem
Certificate:
    Data:
        Version: 3 (0x2)
        Serial Number:
            25:d9:f8:1d:9b:1d:23:f1:21:56:54:f2:43:cc:4f:0e:73:22:be:ec
    Signature Algorithm: sha256WithRSAEncryption
        Issuer: CN=Intermediate CA valid 3 years
        Validity
            Not Before: Apr 24 20:17:30 2018 GMT
            Not After : Apr 24 20:18:00 2019 GMT
        Subject: O=HPE, OU=CloudRA Team, CN=hpe-ucp.cloudra.local
        Subject Public Key Info:
            Public Key Algorithm: rsaEncryption
                Public-Key: (2048 bit)
                CA Issuers - URI:http://localhost:8200/v1/intca
            [ portions removed ]

            X509v3 Subject Alternative Name:
                DNS:hpe-ucp.cloudra.local, DNS:hpe-ucp01.cloudra.local, DNS:hpe-ucp02.cloudra.local, DNS:hpe-ucp03.cloudra.local
```

The following excerpts from `cert.pem` show the first certificate which is the server certificate itself and the second certificate which is the intermediate CA's certificate.

```
-----BEGIN CERTIFICATE-----
MIIFGTCCAwGgAwIBAgIUJdn4HZsdI/EhVlTyQ8xPDnMivuwWdQYJKoZIhvcNAQEL
BQAwKDEmMCQA1UEAxMdSW50ZXJtZWRpYXR1IENBIHZhbnB1kIDMgeWVhcnMwHhcN
...
sOR4I3Qnc50oNISng5l7wW1d4RMMwmXQhG1H5QKAUjHfJXH4bNtIzKxw/zGTVr4Z
l1YKbEwJcgAvvfkn+w==
-----END CERTIFICATE-----

-----BEGIN CERTIFICATE-----
MIIFcjcCA1qgAwIBAgIUax4MhiDP8IjSUg1duVb6kYegSRgwDQYJKoZIhvcNAQEL
BQAwGjEYMBYGA1UEAxMPRXhhbnB1kXSbSb290IENBMB4XDTE4MDQyNDIwMTIwOVox
...
Ni1JnZandVlyQdo0aB2M/1DNffKvwW3JeArKvDA9j95n/BWFTjoZ+Y0z9pYit6T7
1GCGu3be
-----END CERTIFICATE-----
```

Finally, here is an excerpt from `key.pem`, the private key that goes with the server certificate.

```
-----BEGIN RSA PRIVATE KEY-----
MIIEpQIBAAKCAQEA5cmmb52ufE80a3cXhY2HSRZNazb7/fipXY1rZ+U5+rJv9BN5
d/X3NTroSE8/PvoS/maGkHCnURGNqbu/G2umKN/tm/eSpDY861YnGWxj+bc0gtiU
...
AOSgidSMk3hFX1Iaftgx4EUGbrzZ07I8M5R064U1aMFNFyj4XghJ2mZTdNelwNBw
```



```
pr/fYulyi5lYPa1QHYH30yvNqQQ3arEbTbZp8hEyY0gxtZRXmmaoqOY=
-----END RSA PRIVATE KEY-----
```

Verify your certificates

The playbooks do not verify the validity of the certificate files you supply so you should verify them manually before you start your deployment.

Verify that the private and the server cert match

On the Ansible box, run the following commands:

```
ckcert=${openssl x509 -noout -modulus -in cert.pem | openssl md5}
ckkey=${openssl rsa -noout -modulus -in key.pem | openssl md5}
if [ "$ckkey" == "$ckcert" ] ; then echo "Private key and Certificate match" ; else echo "STOP! Private Key and Certificate don't match" ; fi
```

Verify that the server certificate was signed by the CA

Extract all but the first certificate from `cert.pem` (i.e. extract the certs for the intermediate CA authorities) into the file `int.pem`

```
sed -e '1,/-----END CERTIFICATE-----/d' cert.pem >intca.pem
```

Combine `intca.pem` and `ca.pem` to form `cachain.pem`:

```
cat intca.pem ca.pem > cachain.pem
```

Finally, verify that `cert.pem` was signed by the CA or by an intermediate CA:

```
openssl verify -verbose -CAfile cachain.pem cert.pem
```

A successful check will generate output similar to:

```
[root@ansible ucp_certs]# cat intca.pem ca.pem > cachain.pem
[root@ansible ucp_certs]# openssl verify -verbose -CAfile cachain.pem cert.pem
cert.pem: OK
```

An unsuccessful check will generate output similar to:

```
[root@ansible ucp_certs]# openssl verify -verbose -CAfile cachain.pem certsignedbyanotherca.pem
certsignedbyanotherca.pem: 0 = HPE, OU = CloudRA Team, CN = hpe-ucp.cloudra.local
error 20 at 0 depth lookup:unable to get local issuer certificate
```

Appendix C: Enabling SSL between the universal forwarders and the Splunk indexers using your certificates

The procedure for enabling SSL between the universal forwarders and the Splunk indexers using your certificates is described below. In summary, the following steps are required:

1. Set the variable `splunk_ssl` to `yes` in `group_vars/vars`
2. Put your root CA certificate and your server certificate files in `/root/Docker-SimpliVity/files/splunk/linux/SPLUNK_HOME/etc/mycerts`
3. Uncomment the `[sslConfig]` stanza in the file `/files/splunk/linux/SPLUNK_HOME/etc/system/local/server.conf`

Limitations

SSL only works with Linux worker nodes. The Universal Forwarders verify that the indexers they connect to have a certificate signed by the configured root CA and that the Common Name in the certificate presented by the indexer matches its FQDN as listed by the variable `splunk_architecture_forward_servers`.



Prerequisites

Configure your indexers to use SSL on port 9998. The following is an example `inputs.conf` file located in `$(SPLUNK_HOME)/etc/system/local` that enables SSL on port 9998 and configures the certificate file for use by the indexer itself, in this instance `/opt/splunk/etc/mycerts/indexer.pem`.

```
[splunktcp-ssl://9998]
disabled=0
connection_host = ip

[SSL]
serverCert=/opt/splunk/etc/mycerts/indexer.pem
#requireClientCert = true
#sslAltNameToCheck = forwarder,forwarder.cloudra.local

[tcp://1514]
connection_host = dns
sourcetype = ucp
```

For more information, see the documentation at

<https://docs.splunk.com/Documentation/Splunk/7.1.2/Security/ConfigureSplunkforwardingtousesignedcertificates>. In addition, you can see how to create your own certificates and the content of the file designated with `serverCert` at

<http://docs.splunk.com/Documentation/Splunk/7.1.2/Security/Howtoself-signcertificates>.

In this instance, the folder `mycerts` was created under `/opt/splunk/etc` and the file `indexer.pem` was copied to this folder.

Indexers are configured with the Root CA cert used to sign all certificates. This can be achieved by editing the file `server.conf` in `$(SPLUNK_HOME)/etc/system/local` on your indexer(s). The following code block shows the relevant portion of this file where `sslRootCaPath` is pointing to the root CA certificate.

```
[sslConfig]
sslRootCaPath = /opt/splunk/etc/mycerts/ca.pem
```

Note

In order to be able to download and install additional applications, you may want to append the file `$(SPLUNK_HOME)/auth/appsCA.pem` to your `ca.pem` file. If you don't do this, the Splunk UI will make this suggestion when you attempt to **Find more apps**.

Splunk should be restarted on the indexers if you had to make these changes (see the Splunk documentation for more information).

Before you deploy

Generate the forwarder certificate and name it `forwarder.pem`. Make sure that you copy the root CA certificate to `ca.pem`

1. Copy both the `ca.pem` and the `forwarder.pem` files to `files/splunk/linux/SPLUNK_HOME/etc/mycerts/` (overwriting any existing files).
2. Edit the file `server.conf` in the folder `files/splunk/linux/SPLUNK_HOME/etc/system/local` and uncomment the last two lines as suggested in the file itself. Your file should look like this:

```
#
# uncomment the section below if you want to enable SSL
#
[sslConfig]
sslRootCaPath = /opt/splunkforwarder/etc/mycerts/ca.pem
```



3. Set `splunk_ssl` to `yes` in the file `group_vars/vars`, uncommenting the line if required. Make sure that the `splunk_architecture_forward_servers` list specifies all your indexers together with the port that was configured to accept SSL:

```
monitoring_stack: splunk
splunk_ssl: yes
splunk_architecture_forward_servers:
- indexer1.cloudra.local:9998
- indexer2.cloudra.local:9998
```

Hybrid environment Linux / Windows

Currently, you cannot deploy your own certificates for use by the Universal Forwarders deployed on Windows machines. If you want to have your Linux machines in a hybrid deployment to use SSL, proceed as follows.

1. Comment out the `splunk_architecture_forward_servers` variable (and its values) from `group_vars/vars`

```
monitoring_stack: splunk
splunk_ssl: yes
#splunk_architecture_forward_servers:
# - hpe2-ansible.cloudra.local:9998
```

2. Create a file named `vms.yml` in the folder `group_vars` and specify the list of forward servers to use by the Linux servers. This list is typically the same as the one used for Windows servers but specifies a TCP port that enables SSL.

```
splunk_architecture_forward_servers:
- hpe2-ansible.cloudra.local:9998
```

3. Edit the `group_vars/win_worker.yml` file and specify the list of forward servers to be used by the Windows servers. This list is typically the same as the one used for Linux servers but specifies a TCP port that does not enable SSL.

```
splunk_architecture_forward_servers:
- hpe2-ansible.cloudra.local:9997
```



Appendix D: How to check that certs were deployed correctly

The following commands should return the CA certificates used by UCP / DTR. This certificates is the same as the one pointed to by the `--cacert` switch.

```
# curl --cacert <ucp_certs_dir>/ca.pem https://<your ucp fqdn>/ca
# curl --cacert <dtr_certs_dir>/ca.pem https://<your dtr fqdn>/ca
```

Output 1: certificates successfully deployed (content will depend on your own CA certificate)

```
-----BEGIN CERTIFICATE-----
MIIDyTCCAaGgAwIBAgIUUeo+H6xGSB7/9gqq9T2SUwJPLggwDQYJKoZIhvcNAQEL
BQAwbDELMAkGA1UEBhMCRR1FTATBgNVBACjDFRoZSBjbhR1cm5ldETMBEGA1UE
ChMKQ2hyaXN0b3BoZTEUMBIGA1UECjMLQ0EgU2VydmljZXMxGzAZBgNVBAMTEkNo
...
XkJ8WcsHocJ08J9J3RaWsM2BQc7wRntJc0kA7ooTH130tQTP1jFcQp5xNdI4J3Mz
j9BAYERjkGqu7v9tf0em99oVGUa120pu4r73eWU1mL948xuw6PgiRSLZrXhn/RS
uvFVnS/vPYJozOXIZA==
-----END CERTIFICATE-----
```

If the deployment was not successful, `curl` will output something like **Output 2**.

Output 2: certificates were not successfully deployed

```
curl: [60] Peer's Certificate issuer is not recognized.
More details here: http://curl.haxx.se/docs/sslcerts.html
...
```

Enable certs for browser (Windows 2016 example)

Choose `Manage computer certificates` in the control panel as shown in Figure 50.




Figure 50. Manage computer certificates

Import the `ca.pem` for UCP into the Trusted Root Certification Authorities, as shown in Figure 51.




Figure 51. Import the `ca.pem`

It should now show up in the list of certificates. You may need to restart your browser to see the green, secure lock symbol as shown in Figure 52.




Figure 52. Secure HTTPS



Resources and additional links

HPE Reference Architectures, hpe.com/info/ra

HPE SimpliVity, hpe.com/info/simplivity

HPE Servers, hpe.com/servers

HPE Storage, hpe.com/storage

HPE Networking, hpe.com/networking


HPE Technology Consulting Services, hpe.com/us/en/services/consulting.html

Docker Reference Architectures, <https://success.docker.com/architectures>

Splunk Validate Architectures, <https://www.splunk.com/pdfs/white-papers/splunk-validated-architectures.pdf>

Sysdig Resources, <https://sysdig.com/resources/>

To help us improve our documents, please provide feedback at hpe.com/contact/feedback.



Sign up for updates

© Copyright 2018 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty. Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Microsoft, Windows, and Windows Server are registered trademarks or trademarks of Microsoft Corporation in the United States and/or other countries. Linux is the registered trademark of Linus Torvalds in the U.S. and other countries. VMware and vSphere are registered trademarks of VMware, Inc. in the United States and/or other jurisdictions. Red Hat and Red Hat Enterprise Linux are registered trademarks of Red Hat, Inc. in the United States and other countries.

